

TOOL SELECTION FOR PRE-SEED AND SEED FOUNDERS

# The Startup Tool Stack

Last updated August 2026

## Contents

### SELLING

01 CRM & Lead Gen

---

02 Email Sequencing

---

03 Newsletter / Marketing Campaigns

---

04 Social Media Scheduling

---

05 Payments & Billing

---

### BUILDING

06 App Building

---

07 Databases & Backend

---

08 Authentication

---

09 Hosting

---

10 Product Analytics & Experimentation

---

### OPERATING

11 Incorporation & Formation

---

12 Bookkeeping & Accounting

---

13 Customer Support & Helpdesk

---

14 Workflow Automation

---

### REFERENCE

Methodology

---

## Why every chapter ends with a build option

Every chapter closes with a seventh option: building the capability yourself instead of buying it. With AI tools like Claude and Codex, a founder can now stand up a working version of almost anything in this guide.

Building yourself covers a wide range. It can be a spreadsheet you keep by hand, a script against an API, a self-hosted open-source app, or a full custom build with an AI agent writing most of the code.

We include it in every comparison because the real question is no longer whether you can build it. It is whether the result is reliable, what it costs you in setup and upkeep, and when that trade beats paying for a product. The build pages weigh exactly that.

# CRM & Lead Gen

Compared: HubSpot CRM · Attio · Pipedrive · Close · Apollo · ZoomInfo · Build

# At a glance

| TOOL               | CHOOSE THIS IF YOU ARE                                                                                                                                              | STARTING PRICE                                                                                                                                       | MAIN BENEFIT                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
|--------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>HubSpot CRM</b> | <b>Pre-seed to Series A B2B founders who want one system to grow into</b><br>non-technical, US/EU, starting free (2 seats, 1,000 contacts).                         | <b>Free</b><br>2 seats, 1,000 contacts; then \$20/user/mo, or \$7 on annual billing.                                                                 | <b>The combination is what stands out: a free CRM to start on, and a full marketing, sales and service suite on the same platform to grow into.</b><br>The free tier itself is modest. New accounts get 1,000 contacts, 2 users and 1 pipeline, and only accounts opened earlier keep the old 1,000,000-contact limit. What teams are buying is the consolidation, one vendor covering CRM, marketing and service as they grow, instead of the size of the free plan. |
| <b>Attio</b>       | <b>Technical, product-led founders who treat CRM as a flexible database</b><br>pre-seed to seed, startups, remote, B2B, wanting custom objects.                     | <b>Free</b><br>3 users, 50,000 records; Plus \$35/user/mo annual (\$44 monthly), Pro \$79 (\$99 monthly).                                            | <b>Records take the shape a team gives them instead of being forced into fixed contact and deal objects, with a developer API behind that flexibility.</b><br>For a company that has already outgrown a rigid CRM, that is usually the deciding factor.                                                                                                                                                                                                               |
| <b>Pipedrive</b>   | <b>Sales-led founders who live inside a deal pipeline</b><br>non-technical, SMB, B2B, global, small AE teams wanting cheap visual pipeline.                         | <b>\$14/user/mo</b><br>Lite, billed annually; Growth \$39, Premium \$49, Ultimate \$79.                                                              | <b>Everything is organized around stages and moving deals toward close, and that is the whole product.</b><br>Setup is quick and the mobile app is well established. A small sales team gets pipeline management and nothing else, which is the point, not a shortcoming.                                                                                                                                                                                             |
| <b>Close</b>       | <b>Founders whose reps live on the phone</b><br>inside-sales, SMB, B2B, US-centric; small high-velocity SDR/AE teams doing call + SMS + email.                      | <b>\$9/user/mo</b><br>Solo; Essentials \$35, Growth \$99, Scale \$139.                                                                               | <b>Calling is native here, not bolted on through a separate tool, because the Power and Predictive Dialer is built in.</b><br>For a team whose selling happens mostly on the phone, having the dialer, the sequences and the coaching inside one CRM is what separates Close from general-purpose options.                                                                                                                                                            |
| <b>Apollo</b>      | <b>Outbound founders who want a contact database AND sequencing in one tool</b><br>seed-stage, B2B, semi-technical, US/global, building an SDR motion from scratch. | <b>Free tier</b><br>then Basic ~\$49, Professional ~\$79, Organization ~\$119/user/mo; cheaper on annual billing.                                    | <b>Combines a 240M+ contact database with enrichment and outreach sequencing in one platform, starting at a \$0 free tier that is usable for solo prospecting.</b><br>For a founder who wants prospecting and sending in one place instead of several tools, that bundle at that entry price is the differentiator.                                                                                                                                                   |
| <b>ZoomInfo</b>    | <b>Well-funded Series A+ teams with a dedicated RevOps/SDR org</b><br>B2B, mid-market/enterprise motion, US-heavy, budget for annual contracts.                     | <b>Quote-only</b><br>reportedly ~\$15,000/yr to start (3 seats, ~5,000 credits), climbing to \$40,000+; median contract ~\$32,000/yr per Vendr data. | <b>Broad coverage of verified direct-dial phone numbers and firmographic depth, plus established intent data and a prospect-to-engage suite.</b><br>When outbound volume justifies the cost, data quality is the reason teams pay.                                                                                                                                                                                                                                    |
| <b>Build</b>       | <b>Technical founders with under ~100 target accounts</b><br>earliest pre-seed, near-zero budget, hands-on, doing founder-led sales personally.                     | <b>No license fee</b><br>cost is time: setup hours plus your AI plan and API/token burn for enrichment.                                              | <b>Full data ownership, no lock-in, near-zero recurring cost, fully customizable.</b><br>Well-suited to a lean MVP where a founder wants to understand and control the whole funnel.                                                                                                                                                                                                                                                                                  |

Full detail for every tool follows on its own page.

# HubSpot CRM

## BUILT FOR

Growing SMBs and mid-market teams with combined marketing and sales staff of mixed technical skill, running the full funnel from lead capture through deal close in one system. Best fit: when you're outgrowing spreadsheets and want CRM, marketing, and support unified instead of stitched together, and can absorb rising per-seat costs.

---

## CHOOSE THIS IF YOU ARE

Pre-seed to Series A B2B founders who want one system to grow into: non-technical, US/EU, starting free (2 seats, 1,000 contacts).

### BEST WHEN

You're building a combined marketing+sales motion and want CRM, forms, email, and reporting under one roof.

### AVOID IF

You're a pure high-volume outbound shop or watching every dollar; Professional jumps to \$90/seat plus a \$1,500 onboarding fee and contact-tiered Marketing Hub balloons fast.

---

## WHAT IT DOES

**Its free tier caps new accounts at 1,000 contacts**, 2 users and 1 pipeline (as of September 2024; accounts opened earlier keep the old 1,000,000-contact ceiling) and is an entry point into paid marketing and sales suites; the first paid Starter step is \$20/user/mo.

**Contacts arrive from web forms**, email sync, bulk imports and marketing-capture flows, and they all land in the same records the marketing, sales and service hubs read from. Marketplace apps add further inbound sources on top of that. The constraint is the free tier. New accounts stop at 1,000 contacts, while accounts opened before September 2024 keep the old 1,000,000 ceiling. A team collecting a lot of inbound form fills will reach that limit far sooner than the old ceiling allowed.

**HubSpot isn't a data vendor.** It stores whatever you and your integrations put in, and built-in enrichment is limited or sold as an add-on.

**Follow-up runs on workflow automation**, with AI assistants layered on top. The free plan includes only the basics. The deeper automation sits behind paid tiers. A single workflow can act across marketing, sales and service records at once. That same reach is what pushes teams up the tiers as usage grows.

**HubSpot holds your contacts, as any CRM does.** Sourcing and consent are on you, and on any enrichment tool you plug in.

---

### COST

**Free for up to 1,000 contacts**, 2 users and 1 pipeline for new accounts (since Sept 2024), but the old 1,000,000-contact free tier survives only on legacy accounts. The first paid step is Starter at \$20/user/mo month-to-month, or \$7/user/mo on annual billing. The real cost is not the entry price but the climb as additional hubs and higher tiers are added across the suite, where capabilities are tier-gated, not bundled in.

---

### ECOSYSTEM & EXIT

**The app marketplace is one of the broader ones in the CRM category**, connecting marketing, sales and service tools directly into the shared contact records. That breadth cuts both ways. Every app a team wires in is one more integration that would have to be rebuilt elsewhere. The cost of leaving climbs alongside the cost of adopting.

**Getting data in is easy.** Getting out is the part worth thinking about, because every hub adopted and every marketplace app connected becomes one more dependency to unwind. Switching cost rises with each addition, not all at once.

---

### TRACK RECORD

**HubSpot has been publicly traded since 2014**, the longest public track record in this group (ZoomInfo, the other public company here, listed in 2020), and ships enterprise-grade security, which puts vendor risk at the low end of this set. Its scale supports extensive documentation, support and education resources. What comes with that scale is the complexity of the suite itself.

---

### WHERE IT STANDS OUT

**The combination is what stands out:** a free CRM to start on, and a full marketing, sales and service suite on the same platform to grow into. The free tier itself is modest. New accounts get 1,000 contacts, 2 users and 1 pipeline, and only accounts opened earlier keep the old 1,000,000-contact limit. What teams are buying is the consolidation, one vendor covering CRM, marketing and service as they grow, instead of the size of the free plan.

---

### LIMITATIONS

**Cost and complexity both climb with scale.** The free entry point is broad, but each hub added and each step up a tier raises the price and the number of moving parts, and capabilities are gated behind tiers, not included. The breadth that consolidates a team's tooling is the same thing that makes leaving expensive.

---

Pricing and limits read from [HubSpot CRM's own site](#) . Verified to the edition date above.

# Attio

## BUILT FOR

Startups and modern go-to-market teams whose product-minded operators are comfortable configuring their own data models and workflows around custom objects. Best fit: early-to-growth companies that find legacy CRMs rigid and want a fast, flexible system they shape themselves instead of adapting to.

---

## CHOOSE THIS IF YOU ARE

Technical, product-led pre-seed to seed founders who treat CRM as a flexible database: startups, remote, B2B, wanting custom objects.

## BEST WHEN

You want to sync product usage data, get real-time bidirectional Gmail/Calendar sync, and build custom workflows. Plus is \$35/seat and Pro \$79/seat (annual).

## AVOID IF

You need turnkey marketing automation, plug-and-play sales templates, or a native dialer out of the box, or you're wary of the AI/enrichment credit metering added in 2026.

---

## WHAT IT DOES

It bills on a hybrid model that combines per-seat pricing (Plus \$35/user/mo, Pro \$79/user/mo on annual billing) with usage-based credits, so cost tracks both seats and consumption.

Rather than push records into a fixed contact and deal structure, **Attio organizes them into a data model the team defines**. Records come in through imports, email and calendar sync, or the API. Anything more automated than those built-in syncs has to be built against the developer API, which puts it out of reach for a team without engineering time.

**Attio holds the data you put in and adds some AI enrichment.** It isn't a data provider, and coverage isn't its job.

Because the data models are defined by the team instead of fixed by Attio, **follow-up rules can reference whatever fields that team created**, and the AI works against the same structure. The flexibility is real, but nothing runs until those models and rules are built. That cost comes up front, before the CRM is doing any work.

**The data in Attio is first-party**, and compliance rests with you.

## COST

**Free for 3 users and 50,000 records**, then Plus at \$35/user/mo annual (\$44 monthly) and Pro at \$79/user/mo annual (\$99 monthly), both raised ~July 2026. The less obvious cost is the hybrid model. Per-seat pricing is layered with usage-based credits, so spend tracks headcount and consumption together instead of seats alone, making bills harder to predict as usage scales.

## ECOSYSTEM & EXIT

**Connections are built against a developer API instead of picked from a prebuilt catalog**, and incoming data maps into the custom data model. For a team with engineering time, that means almost anything can be wired up. For a team without it, the prebuilt set is the practical ceiling.

**Import maps into the custom data model**, so records can be reshaped on the way in, and standard export is available on the way out. The complication is that the structure belongs to the workspace that built it. An export will not drop cleanly into another CRM's fixed contact and deal objects without remapping first.

## TRACK RECORD

**Growth is fast and development is active**, with standard security in place. The operating history is shorter than the incumbents here, and that is the trade a buyer accepts in return for how quickly the product improves.

## WHERE IT STANDS OUT

**Records take the shape a team gives them instead of being forced into fixed contact and deal objects**, with a developer API behind that flexibility. For a company that has already outgrown a rigid CRM, that is usually the deciding factor.

## LIMITATIONS

**Spend is harder to forecast than with flat per-seat pricing**, because the model combines seats with usage credits and the bill moves with consumption as well as headcount. The custom data models also carry a learning curve before the CRM does anything useful. As a newer entrant, Attio's track record is shorter than the incumbents in this group.

---

Pricing and limits read from [Attio's own site](#) . Verified to the edition date above.



# Pipedrive

## BUILT FOR

Small sales teams and solo founders, non-technical reps managing a visual pipeline of deals with straightforward activity tracking. Best fit: when you want an affordable, pipeline-first CRM that's quick to set up and doesn't need marketing or support modules bolted on.

---

## CHOOSE THIS IF YOU ARE

Non-technical, SMB, sales-led founders who live inside a deal pipeline: B2B, global, small AE teams wanting cheap visual pipeline.

## BEST WHEN

You have an inbound/outbound sales motion and just want to drag deals across stages without complexity. Growth \$39, Premium \$49/seat.

## AVOID IF

You need a rich contact database, native enrichment, or real marketing automation; Pipedrive is thin there and add-ons stack up.

---

## WHAT IT DOES

**Core plans run annually from Lite \$14 to Ultimate \$79 per user/mo**, with capabilities like LeadBooster and Smart Docs sold as separate add-ons.

**Every lead that comes in becomes a deal sitting at a stage.** Web forms feed that pipeline through the LeadBooster add-on, billed separately from the plan, and email sync and imports cover the rest. The orientation is toward staged deals instead of broad marketing lists, so holding contacts outside the pipeline is not what this is built for.

**Pipedrive is a pipeline CRM rather than a data vendor**, so you bring the contacts.

## Automation here stays close to the deal:

moving it between stages, creating tasks, triggering follow-ups. An AI assistant covers the same pipeline activity. The scope stops at sales. It does not extend to marketing campaigns. Contacts a team wants to warm up before they enter the pipeline need a separate tool.

**There's no sourcing question here.** Pipedrive holds your first-party data.

## COST

**Billed annually, per-user pricing runs Lite \$14,** Growth \$39, Premium \$49 and Ultimate \$79, the July 2025 rebrand of Essential/Advanced/Professional. The hidden cost is add-ons: LeadBooster and Smart Docs are billed separately on top of the chosen plan. The effective price for a team using those features sits above the headline tier.

## ECOSYSTEM & EXIT

**LeadBooster and Smart Docs are paid add-ons on top of the plan,** not bundled capabilities. The catalog itself covers the common sales tools and stays within that scope and does not reach across a full marketing suite.

**Both directions are standard operations.** The data is simply contacts and deals sitting on stages, so an export carries most of what matters, and switching cost stays lower here than with the broader suites.

## TRACK RECORD

**Pipedrive has been in the sales CRM market a long time and is widely used,** so its feature set and support base are settled instead of fast-changing. Security is standard, and vendor risk sits lower than with the newer entrants here.

## WHERE IT STANDS OUT

**Everything is organized around stages and moving deals toward close,** and that is the whole product. Setup is quick and the mobile app is well established. A small sales team gets pipeline management and nothing else, which is the point, not a shortcoming.

## LIMITATIONS

**The limits follow directly from the focus.**

LeadBooster and Smart Docs are add-ons, not included features, so a team that needs them pays more than the advertised tier suggests. Pipedrive is also deliberately less all-in-one than HubSpot, trading marketing and service breadth for depth in the pipeline.

---

Pricing and limits read from [Pipedrive's own site](#) . Verified to the edition date above.

# Close

## BUILT FOR

High-velocity inside-sales teams at SMBs and startups, reps who live in calls and email, running outbound with built-in calling and SMS. Best fit: small teams whose primary motion is phone-and-email selling and want dialer, inbox, and CRM in one place instead of separate tools.

---

## CHOOSE THIS IF YOU ARE

Inside-sales SMB founders whose reps live on the phone: B2B, US-centric, small high-velocity SDR/AE teams doing call + SMS + email.

### BEST WHEN

You're running high-volume calling and want the Power Dialer, SMS, and email sequencing built directly into the CRM with no Aircall bolt-on. Solo \$9, Essentials \$35, Growth \$99/user/mo on annual billing.

### AVOID IF

You're product-led or marketing-driven with little outbound calling, or the usage-based calling costs on top of seat price scare you; Close bills calling à la carte (outbound ~\$0.02/min, phone numbers from ~\$1/mo, premium IVR/routing numbers charged by usage) rather than a fixed markup.

## WHAT IT DOES

**Per-seat plans run Solo \$9, Essentials \$35, Growth \$99 and Scale \$139/user/mo**, and the Power Dialer sits on the Growth and Scale tiers, with the Predictive Dialer and call coaching on Scale only.

**A dialed or received call becomes a tracked record on its own**, and that is the main way lists grow here. Contacts also arrive through imports, and email captured inside the CRM adds records the same way. Lists build from calling activity more than web-form capture, which suits a team already making calls more than one waiting on inbound.

**You import and enrich contacts yourself.** Close isn't a data provider, and its focus is working those contacts rather than supplying them.

---

## Sequences and the Power and Predictive

**Dialer work together**, with AI in the mix. A follow-up sequence can combine calls and email instead of relying on email alone. Where those features sit is the limit. The Power Dialer requires Growth at \$99 or Scale at \$139 per user per month on annual billing, and the Predictive Dialer and call coaching are Scale-only, which means Solo and Essentials do not include the calling automation most buyers come to Close for.

**This doesn't apply to Close, which holds first-party CRM data.**

---

### COST

**Per-seat plans run Solo \$9, Essentials \$35, Growth \$99 and Scale \$139/user/mo.** The tier-gating is the point to watch. The Power Dialer requires Growth (\$99) or Scale (\$139), and the Predictive Dialer and call coaching are Scale-only (\$139). These are the features Close is chosen for, so the two entry tiers omit its defining capability.

---

### ECOSYSTEM & EXIT

**Integrations cover the sales tools that sit around calling and email**, enough to run call sequences and log activity against a record. There is no general-purpose marketplace behind that. Anything outside the sales stack has to be built and maintained by the team.

Because Close stays within sales and calling instead of spreading across a suite, **there is less connected tooling to unwind when a team leaves**. Import and export themselves work the way you would expect.

---

### TRACK RECORD

**There is a track record in the calling CRM niche**, with standard security behind it. Close is smaller than the publicly traded incumbents in this group and older than the newest entrants, which places it in the middle on both counts.

---

### WHERE IT STANDS OUT

**Calling is native here, not bolted on through a separate tool**, because the Power and Predictive Dialer is built in. For a team whose selling happens mostly on the phone, having the dialer, the sequences and the coaching inside one CRM is what separates Close from general-purpose options.

---

### LIMITATIONS

**The features that define Close are the ones behind the paywall.** The Power Dialer requires Growth at \$99, while the Predictive Dialer and call coaching sit on Scale at \$139 only, so Solo and Essentials leave out the capability most buyers arrive for. Beyond that the scope is narrow by design. This is outbound calling, not a broad CRM suite.

---

Pricing and limits read from **Close's own site** . Verified to the edition date above.

# Apollo

## BUILT FOR

Startup and SMB sales teams whose SDRs blend prospecting with outreach, sourcing B2B contacts and running sequences from a single database. Best fit: budget-conscious teams that want a contact database plus email sequencing combined, without paying separately for data and outreach tools.

---

## CHOOSE THIS IF YOU ARE

Seed-stage outbound B2B founders who want a contact database AND sequencing in one tool: semi-technical, US/global, building an SDR motion from scratch.

## BEST WHEN

The free tier gives 100 email credits/mo.

## AVOID IF

You need pristine, deeply-verified enterprise data or accurate direct-dials at scale.

## WHAT IT DOES

**The free tier is \$0 and usable for solo prospecting**, though mobile numbers and export credits are capped and are where Apollo directs upgrades.

**Bundles a large B2B contact database with enrichment plus CSV import.** You pull contacts from Apollo's own DB into sequences.

**Quality is bounded by Apollo's own 240M+ database**, a single proprietary source, not a multi-provider blend. 2026 practitioner tests (Salesforce's four-week hands-on review, verification-vendor analyses from Prospeo and Tomba, recurring r/coldemail reports) put real-world accuracy at roughly 65-80% depending on segment and geography, with bounce rates on "verified" contacts running from single digits on US mid-market targets to over 30% elsewhere, against Apollo's advertised ~91%. Note that several of these testers sell competing tools. "Inaccurate data" also recurs across hundreds of G2 reviews. Free-tier mobile numbers are capped, limiting verification without paying.

**Built-in outreach sequences and AI writing (email/dialer)**, but it acts as a light CRM/outreach engine.

**Apollo builds its contact data from its own 240M+ database instead of buying from third-party aggregators.** Whatever sourcing risk exists is concentrated in that one dataset, not spread across suppliers. Security is described as standard. No specific certifications appear in the sources used here, so there is not enough to judge the compliance position either way.

---

#### COST

**Four tiers: Free at \$0, Basic ~\$49/user/mo, Professional ~\$79/user/mo and Organization ~\$119/user/mo**, each cheaper on annual billing. The free plan is usable for solo prospecting, but capped mobile numbers and export credits are where Apollo upsells hardest into the paid tiers.

---

#### ECOSYSTEM & EXIT

**CRM sync sits alongside Apollo's own outreach sequencing**, so a record can be enriched and then emailed without leaving the platform. Apollo runs in the browser with an extension. Which CRMs the sync supports is not documented in the sources used here.

**Export is credit-capped.** Moving enriched data out is limited by plan credits.

---

#### TRACK RECORD

**Very popular B2B tool; main knock is data limited to Apollo's own database.**

---

#### WHERE IT STANDS OUT

**Combines a 240M+ contact database with enrichment and outreach sequencing in one platform**, starting at a \$0 free tier that is usable for solo prospecting. For a founder who wants prospecting and sending in one place instead of several tools, that bundle at that entry price is the differentiator.

---

#### LIMITATIONS

**Data is confined to Apollo's own 240M+ database instead of aggregated sources**, so coverage is capped at what Apollo holds. On the free tier, mobile numbers and export credits are limited and are where Apollo upsells hardest; export being credit-capped constrains bulk extraction without upgrading.

---

Pricing and limits read from [Apollo's own site](#) . Verified to the edition date above.

# ZoomInfo

## BUILT FOR

Mid-market and enterprise sales and marketing orgs with dedicated RevOps staff, sourcing high-accuracy B2B data with intent signals and deep integrations.

---

## CHOOSE THIS IF YOU ARE

Well-funded Series A+ B2B teams with a dedicated RevOps/SDR org: mid-market/enterprise motion, US-heavy, budget for annual contracts.

## BEST WHEN

Accurate firmographics, intent data, and verified direct-dials materially change your pipeline and you can absorb ~\$15k+/yr annual commitments.

## AVOID IF

You're pre-seed/seed watching burn; annual-only contracts and five-figure minimums are wrong for early experimentation.

---

## WHAT IT DOES

**Search/filter the database by firmographics, technographics, job title and intent, then export or push to CRM.** Enrichment auto-fills and updates records. WebSights de-anonymizes site visitors, but FormComplete enriches inbound form fills.

**Among the largest and most accurate B2B datasets:** hundreds of millions of contacts with verified direct-dial phone numbers, a key differentiator. Coverage is strongest in North America. International and SMB data is thinner and ages faster.

## Bundles engagement tooling: Chorus

**conversation intelligence**, Engage for sequenced email/dialer outreach, Copilot AI for account prioritization and recommended actions. Depth rivals standalone sales-engagement tools but sits behind higher tiers.

**Data is aggregated from public sources**, contributory community networks, partnerships and web crawling. Offers GDPR/CCPA tooling, DNC scrubbing and notice mechanisms, but its sourcing model draws ongoing privacy scrutiny and opt-out obligations.



---

## COST

**Quote-only, never published.** Tiers Professional/Advanced/Elite reportedly start ~\$15,000/yr (3 seats, ~5,000 credits) and climb to \$40,000+; intent, international and other add-ons cost extra. Median contract ~\$32,000/yr per Vendr data.

---

## ECOSYSTEM & EXIT

**Native two-way sync with Salesforce,** HubSpot, Microsoft Dynamics and major sales-engagement tools (Outreach, Salesloft), plus Slack, marketing automation and a REST API/webhooks for enrichment pipelines.

**Getting in is easy.** CSV import and CRM sync enrich existing records. Exiting is harder. Data is licensed, not owned. Contracts restrict retained use, and auto-renewal clauses with firm opt-out windows are a common lock-in complaint.

---

## TRACK RECORD

**Public company (Nasdaq: GTM, renamed from ZI in May 2025) with a long track record and large customer base.** Enterprise support and CSMs on higher tiers, but reviews cite aggressive sales/renewal tactics and rigid annual contracts.

---

## WHERE IT STANDS OUT

**Broad coverage of verified direct-dial phone numbers and firmographic depth,** plus established intent data and a prospect-to-engage suite. When outbound volume justifies the cost, data quality is the reason teams pay.

---

## LIMITATIONS

**Expensive, quote-gated, 3-seat minimum and add-on-heavy,** a poor fit pre-revenue. Annual lock-in, aggressive renewals, thinner international/SMB data, and privacy/compliance overhead. Cheaper tools (Apollo, Clay) suffice for early prospecting.

---

Pricing and limits read from [ZoomInfo's own site](#) . Verified to the edition date above.



# Build

## BUILT FOR

Not a product but a stack you assemble: Airtable/Notion/Sheets as the DB plus AI enrichment scripts, or a self-hosted EspoCRM/Twenty. Maximum control, zero license fee, all the plumbing is on you. Technical or scrappy pre-seed founders at MVP scale (hundreds, not tens of thousands of leads) who value control and near-zero recurring cost over polish, and can tolerate wiring things together themselves.

---

## CHOOSE THIS IF YOU ARE

Earliest pre-seed technical founders with under ~100 target accounts and near-zero budget: hands-on, doing founder-led sales personally.

### BEST WHEN

Your list is small and tightly curated and you want total control. Airtable or Sheets plus Clay (or an enrichment API / LLM) to fill firmographics costs a fraction of a CRM seat and forces you to know every account cold.

### AVOID IF

More than one person touches deals, or you need pipeline reporting, sequence tracking, or activity logging; you'll drop follow-ups fast.

## WHAT IT DOES

**You own the schema, the data, and every workflow.**

**Manual CSV imports from Apollo/LinkedIn**

**Sales Nav exports**, form/webhook captures, or a scraper you write. No native one-click sync. You build each pipe. Fine at low volume, tedious as inflow grows.

**Only as good as the source exported from and the enrichment prompt written.** No proprietary database, just piggybacking on Apollo/scrapes. Expect stale emails, dedupe headaches, and gaps you clean by hand.

**DIY via scripts, cron jobs, or Make/n8n plus LLM calls for enrichment/scoring/drafting.**

Capable and cheap once built, but every automation is code you must write, debug, and maintain, nothing works out of the box.

**You inherit full liability.** No vendor DPA, opt-out handling, or GDPR/CAN-SPAM guardrails. Scraping LinkedIn breaches its ToS. You must bolt on consent tracking and suppression lists yourself or risk fines.

---

#### COST

**No per-seat fee and no feature gates.** Cost is time-vs-money: mostly one-time setup hours plus your AI plan and API/token burn for enrichment. Only beats a fixed monthly SaaS long-term if you keep using it.

---

#### ECOSYSTEM & EXIT

**Whatever you code.** Open-source CRMs expose REST APIs/webhooks. Airtable/Notion have decent APIs and Zapier/Make connectors. No curated marketplace. Each integration is a build-and-maintain task, not a toggle.

**Easiest exit of the group.** It's your DB, export to CSV/SQL anytime, no contractual lock-in. Import is manual CSV mapping. Moving between Airtable/Notion/EspoCRM is straightforward since you control the schema.

---

#### TRACK RECORD

**No vendor SLA or support desk.** You're the support. Airtable/Notion are mature; EspoCRM/Twenty are solid OSS but you self-patch. Help comes from docs, forums, and community, not a paid line.

---

#### WHERE IT STANDS OUT

**Full data ownership, no lock-in,** near-zero recurring cost, fully customizable. Well-suited to a lean MVP where a founder wants to understand and control the whole funnel.

---

#### LIMITATIONS

**Won't scale past small volume without real eng effort.** You own all security, compliance, and dedupe risk; UI/UX is unpolished unless you have in-house skill. Maintenance is a permanent tax you keep paying.

---

No vendor page: this route is assembled from the components named above.

# Email Sequencing

Compared: Instantly · Smartlead · SalesHandy · QuickMail · Salesforge · Woodpecker · Build

# At a glance

| TOOL              | CHOOSE THIS IF YOU ARE                                                                                                                         | STARTING PRICE                                                                                                                                                                                       | MAIN BENEFIT                                                                                                                                                                                                                                   |
|-------------------|------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Instantly</b>  | <b>Seed-stage outbound founders and agencies scaling cold email volume</b><br>semi-technical, US/global, running many sending inboxes at once. | <b>Outreach: Growth ~\$37.60/mo annual</b><br>(~\$47 monthly), Hypergrowth \$97, Lightspeed \$358                                                                                                    | <b>Sheer scale + ecosystem: unlimited inboxes/warmup, biggest community and playbooks, and add-on modules (leads, inbox placement, DFY) that make it a one-stop outbound stack for volume senders.</b>                                         |
| <b>Smartlead</b>  | <b>Agencies and multi-client seed founders managing many client sending accounts</b><br>semi-technical, global, deliverability-obsessed.       | <b>Base \$39/mo, Pro \$94</b><br>Unlimited Smart \$174, Unlimited Prime \$379<br>(annual saves ~17%)                                                                                                 | <b>Unlimited everything + the deepest API, webhooks, and white-label, ideal for automation and agencies.</b><br>Flat per-account pricing (unlimited seats) beats per-user rivals as teams grow.                                                |
| <b>SalesHandy</b> | <b>Budget-conscious solo and seed founders who want high send volume cheaply</b><br>non-technical, global, straightforward cold email.         | <b>Outreach Starter \$25/mo annual</b><br>(~\$36 monthly), Pro \$69, Scale \$139, Scale Plus \$209, with unlimited email accounts on all                                                             | <b>Price-to-value: Unlimited email accounts from \$25/mo is the cheapest credible way to run real multi-inbox cold email.</b><br>Easy setup and fair pricing are the most-praised traits.                                                      |
| <b>QuickMail</b>  | <b>Experienced outbound founders who already have a data source and prize inbox placement</b><br>semi-technical, US/EU, deliverability-first.  | <b>Single track: Starter \$49/mo</b><br>(5,000 emails/1,000 contacts), Growth \$99/mo (100,000 emails/25,000 contacts), Agency \$299/mo (500,000 emails/100,000 contacts)                            | <b>Deliverability engineering (auto sender-swapping, free warmer) plus no-per-seat pricing, a rare combo that keeps costs flat while protecting inbox placement.</b><br>Clean vendor-conduct record vs rivals.                                 |
| <b>Salesforge</b> | <b>AI-forward seed founders and agencies wanting cheap unlimited mailboxes plus AI personalization</b><br>semi-technical, global, high-volume. | <b>Pro from \$48/mo</b><br>(5,000 emails, 1 user, unlimited premium warmup, 300 validation + 300 personalization credits/mo); Growth from \$96/mo (50k emails, 10k active contacts, unlimited users) | <b>AI-native workflow (write → send → warm → reply) with unlimited mailboxes/LinkedIn at a founder price, plus a clean upgrade to an autonomous AI SDR (Agent Frank) when you're ready to hand off prospecting.</b>                            |
| <b>Woodpecker</b> | <b>Quality-over-quantity B2B founders and agencies who value deliverability</b><br>compliance, and unlimited inboxes/seats                     | <b>Priced per contacted prospect</b><br>(~\$7 per 100/mo monthly; ~33% off annual)                                                                                                                   | <b>Deliverability-first reputation, unlimited inboxes/seats, mature EU/GDPR-friendly compliance stance, and an Agency Panel (\$27/mo per active client, \$5/mo white-label) with per-client databases, bulk actions, and automated audits.</b> |
| <b>Build</b>      | <b>Highly technical founders sending low volume to a tightly curated list who want total control</b><br>near-zero per-email cost.              | <b>Cheap per - email: SES ~\$0.10/1k</b><br>Resend/Postmark have small free tiers                                                                                                                    | <b>Lowest possible send cost, total control of deliverability levers, and AI personalization limited only by your prompt skill.</b><br>Ideal for a founder doing surgical, low-volume outreach who wants to own the pipe.                      |

Full detail for every tool follows on its own page.

# Instantly

## BUILT FOR

Agencies and solo outbound operators, growth marketers running cold email at volume across many inboxes with deliverability and warmup tooling and a lead database. Best fit: high-volume senders who prioritize inbox rotation and low cost over deep CRM integration or per-client reporting.

---

## CHOOSE THIS IF YOU ARE

Seed-stage outbound founders and agencies scaling cold email volume: semi-technical, US/global, running many sending inboxes at once.

## BEST WHEN

You're sending tens of thousands of emails across dozens of domains and want unlimited email accounts plus free warmup on every plan. Pick over per-inbox rivals because Growth is \$47 (1k leads/5k emails) and Hypergrowth \$97 unlocks 25k leads/100k emails plus the Unibox, a ~10x drop in cost per 1k emails.

## AVOID IF

You want tight CRM or multichannel (LinkedIn + calling), or hate that leads and CRM are billed as separate products inflating the invoice.

## WHAT IT DOES

**Sending is one modular product among leads/CRM/enrichment.**

**Unlimited email accounts and unlimited warmup on every plan are the core of the offer.**

Growth handles solo volume;

Hypergrowth/Lightspeed add far higher send caps, dedicated servers, and IP rotation for the heaviest senders.

**Unlimited free warmup network + optional paid Inbox Placement (spam-test) module.**

Benchmarks among the highest (~94% inbox in one 2026 test), but Q1-2026 Reddit threads flag DFY domains arriving reputation-blocked and 'healthy' scores that still hit spam.

**Multi-step conditional sequences, A/Z testing, auto inbox-rotation, unified Unibox reply management (Hypergrowth+), open/click/reply tracking and sub-sequences. Solid but email-first; multichannel is thinner than lemlist/Reply.**

## Spintax, custom variables, and AI copy/optimization tools plus AI reply features.

Competent instead of a differentiator.

Salesforge/lemlist push harder on AI-written, per-prospect personalization at scale.

---

### COST

**Outreach: Growth ~\$37.60/mo annual (~\$47 monthly)**, Hypergrowth \$97, Lightspeed \$358. No permanent free plan. Free trial only. Modular: leads/CRM/inbox-placement stack on top. Real spend is often \$100-150+/mo.

---

### ECOSYSTEM & EXIT

**Native API, webhooks, Zapier/Make**, and its own built-in CRM module. Syncs to HubSpot/Pipedrive via automations. Ecosystem is broad. Deepest native path is staying inside Instantly's own CRM add-on.

**Moderate.** Campaigns/leads export via CSV/API, but warmup reputation and Unibox history don't transfer, and mailbox reconnection + re-warming elsewhere takes weeks. Lock-in is reputational more than contractual.

---

### WHERE IT STANDS OUT

**Sheer scale + ecosystem: unlimited inboxes/warmup**, biggest community and playbooks, and add-on modules (leads, inbox placement, DFY) that make it a one-stop outbound stack for volume senders.

---

### LIMITATIONS

**Modular pricing hides real cost (leads/CRM/placement stack fast).** 2026 complaint spike on DFY-domain reputation, misleading warmup scores, and slow/bot-only support as accounts grow. Overkill for a founder sending a few hundred/mo.

---

Pricing and limits read from [Instantly's own site](#) . Verified to the edition date above.

# Smartlead

## BUILT FOR

Agencies and outbound teams whose technical operators manage many client accounts, scaling cold email with unlimited inboxes and API-driven automation. Best fit: agencies running campaigns for multiple clients that need white-label reporting and heavy programmatic control via API.

---

## CHOOSE THIS IF YOU ARE

Agencies and multi-client seed founders managing many client sending accounts: semi-technical, global, deliverability-obsessed.

## BEST WHEN

You run outbound for multiple brands and need unlimited mailboxes, inbox rotation, and a master inbox with strong API/webhooks. Base \$39 (2k leads/6k emails), Pro \$94 (90,000 sends/mo, 30k leads).

## AVOID IF

You're a solo founder with one domain and modest volume. You'll overpay for agency features.

## WHAT IT DOES

**The whole platform leans heavily on its API and integrations.**

**Unlimited email accounts on every plan with automatic rotation.** Base = 6,000 emails/2,000 leads. Pro = 90,000 emails/30,000 leads. Unlimited tiers add unlimited contact storage and much higher send volume.

**AI-powered unlimited warmup +**

**SmartDelivery spam testing.** Community benchmarks ~73% avg inbox (80%+ on Google OAuth). Watch-out: 2026 reports of mailboxes silently paused mid-warmup with no notification, which can dent reputation.

**Multi-step, multichannel sequences with conditional logic,** A/B testing, auto follow-ups, and webhook-triggered branches. Master inbox centralizes replies. Automation depth is a genuine strength for builders.

**Dynamic variables, spintax, AI copy and AI reply categorization**, plus API to inject custom enrichment per lead. Flexible via API; there is no polished built-in AI writer.

---

Pricing and limits read from [Smartlead's own site](#) . Verified to the edition date above.

## COST

**Base \$39/mo, Pro \$94, Unlimited Smart \$174, Unlimited Prime \$379** (annual saves ~17%). 14-day free trial, no permanent free tier. All plans include unlimited email accounts. Budget ~3-5x plan price for mailboxes/domains.

---

## ECOSYSTEM & EXIT

**Extensive API + webhooks, native HubSpot/Salesforce/Pipedrive/Zapier/Make/Clay connectors**, and white-label portals. Among the more developer-friendly and integration-rich options here.

**Moderate-to-high if you've built API workflows or white-label clients on it**; low for raw data (CSV/API export of campaigns/leads). Warmup reputation and integration wiring are the sticky parts.

---

## WHERE IT STANDS OUT

**Unlimited everything + the deepest API**, webhooks, and white-label, ideal for automation and agencies. Flat per-account pricing (unlimited seats) beats per-user rivals as teams grow.

---

## LIMITATIONS

**Reliability wobbles at scale (campaigns failing to send, warmup pauses, analytics not loading) and inconsistent support are the top 2026 complaints**. One Trustpilot bait-and-switch discount report. Steeper than SalesHandy for a first-timer.



# SalesHandy

## BUILT FOR

Small teams and individual sellers, non-technical users running cold email sequences with basic deliverability and tracking on a tight budget.

---

## CHOOSE THIS IF YOU ARE

Budget-conscious solo and seed founders who want high send volume cheaply: non-technical, global, straightforward cold email.

## BEST WHEN

You want the lowest entry into serious volume (Starter \$25 with 6k emails/2k prospects, then Pro \$69 leaps to 150k emails/30k prospects with unlimited email accounts and no per-seat charge).

## AVOID IF

You need advanced deliverability tooling, multichannel, or polished agency whitelabel. It's more utilitarian.

## WHAT IT DOES

**Unlimited email accounts on every tier.** Starter = 2,000 active prospects. Pro = 150,000 emails/30,000 prospects, but Scale tiers push higher. Priced on active prospects, not raw sends. List churn matters.

**Included TrulyInbox warmup + sender rotation;** ~87% inbox placement in one 2026 benchmark (below Instantly's ~94%). Warmup analytics and monitoring are basic, with no detailed reputation dashboards. Inbox-placement testing exists but as a separate paid product line (free 2 tests/mo, then \$49-\$199/mo) instead of bundled into outreach plans.

**Unlimited multi-stage sequences, auto follow-ups on reply/open conditions,** A/B testing, and email rotation across mailboxes. Covers the founder essentials cleanly; lighter on multichannel and advanced branching.

**Merge tags, spintax, and AI writing assistant for subject/body variants.** Functional and improving, but not a personalization-at-scale engine like Salesforge.

## COST

**Outreach Starter \$25/mo annual (~\$36 monthly)**, Pro \$69, Scale \$139, Scale Plus \$209, with unlimited email accounts on all. 7-day free trial (100 emails, 5 lead credits). Lead Finder is separate, from \$29/mo.

---

Pricing and limits read from [SalesHandy's own site](#) . Verified to the edition date above.

## ECOSYSTEM & EXIT

**Zapier + native HubSpot/Pipedrive/Salesforce sync**, but CRM integrations are gated behind the \$69/mo Pro plan, a catch for Starter users. API available on higher tiers.

**Low-to-moderate.** Prospects/sequences export via CSV/API. Unlimited-account model means no per-inbox unwind. Warmup reputation still needs rebuilding elsewhere, but contractual lock-in is minimal.

## WHERE IT STANDS OUT

**Price-to-value: Unlimited email accounts from \$25/mo is the cheapest credible way to run real multi-inbox cold email.** Easy setup and fair pricing are the most-praised traits.

## LIMITATIONS

**Warmup analytics unreliable, Lead Finder data quality produces high bounce rates**, CRM sync paywalled at Pro, and 2026 G2 reviews flag UI bugs breaking placeholders mid-sequence. Not for strong deliverability or heavy scaling.

# QuickMail

## BUILT FOR

Small B2B teams and founders, hands-on operators running cold outreach with auto-follow-ups and deliverability monitoring. Best fit: established senders who value inbox rotation and reliable warmup for steady, moderate-volume campaigns instead of mass sending.

---

## CHOOSE THIS IF YOU ARE

Experienced outbound founders who already have a data source and prize inbox placement: semi-technical, US/EU, deliverability-first.

## BEST WHEN

Your number-one problem is landing in the primary inbox and you want native MailFlow warmup, inbox rotation, and deliverability tooling baked in from ~\$49/mo.

## AVOID IF

You're purely price-shopping high volume or want built-in lead data and AI personalization.

---

## WHAT IT DOES

### Connect multiple mailboxes with auto-rotation.

Tiers scale on monthly send volume (Starter 5k emails/1k contacts up to Agency 500k/100k; Agency includes 2 workspaces, +\$49 per extra workspace). No per-user charge as the team grows.

**The core strength: Free built-in Auto Warmer + Deliverability AI that analyzes hundreds of signals daily and auto-swaps weak senders via Smart Sender Groups.** 4.7/5 on G2 (116 reviews) with deliverability as the standout.

**Conditional multi-step sequences with A/B testing and reply detection**, but uses a bucket-then-schedule model (add prospects to a bucket, then a schedule feeds the campaign), which is capable but less intuitive than one-click adding.

**Attributes/merge variables, conditional content**, and spintax for variation. AI features are lighter than Salesforce/Instantly; QuickMail bets on deliverability engineering over AI copy.

## COST

**Single track: Starter \$49/mo (5,000 emails/1,000 contacts)**, Growth \$99/mo (100,000 emails/25,000 contacts), Agency \$299/mo (500,000 emails/100,000 contacts). Free AutoWarmer via MailFlow on all plans, with no send limit on warmup. Volume-based caps. No per-seat fees.

---

Pricing and limits read from [QuickMail's own site](#) . Verified to the edition date above.

## ECOSYSTEM & EXIT

**Native HubSpot/Pipedrive sync, Zapier**, and a REST API + webhooks. Solid connectors, though it deliberately has no lead database (doesn't find or verify emails). You bring your own list.

**Low-to-moderate.** Campaigns/contacts export cleanly. No per-seat unwind. The bucket/schedule structure and warmup reputation are what you rebuild elsewhere, not contractual lock-in.

## WHERE IT STANDS OUT

**Deliverability engineering (auto sender-swapping, free warmer) plus no-per-seat pricing**, a rare combo that keeps costs flat while protecting inbox placement. Clean vendor-conduct record vs rivals.

## LIMITATIONS

**Dated UI and a learning curve (bucket/schedule model)**, no built-in lead finding or email verification, and no real-time support during US business hours. Best when you already have verified lists and value placement over polish.

# Salesforge

## BUILT FOR

Outbound teams and agencies whose operators are comfortable with AI tooling, generating and personalizing cold email at scale across many mailboxes. Best fit: senders who want AI-written personalization and mailbox management combined, and are betting on automation-first workflows over manual copywriting.

---

## CHOOSE THIS IF YOU ARE

AI-forward seed founders and agencies wanting cheap unlimited mailboxes plus AI personalization: semi-technical, global, high-volume.

## BEST WHEN

You want the lowest-cost unlimited-inbox plan with AI-written sequences and included Warmforge warmup (Pro \$40 annual-billed with 5k emails and unlimited mailboxes, Growth \$80 with unlimited users and 50k emails).

## AVOID IF

You want a long deliverability track record or mature reply/CRM management.

---

## WHAT IT DOES

**Connect unlimited mailboxes with no per-inbox charge (1 LinkedIn sender on Pro, unlimited on Growth).** Volume scales by plan (Pro ~5k emails/mo, Growth ~50k). Unlimited-mailbox model matches the category leaders at an indie price.

**Free unlimited premium warmup via its own Warmforge network on every plan,** plus built-in validation credits. Newer network with a shorter public track record than Instantly/QuickMail, so deliverability data is less battle-tested.

**Multichannel (email + LinkedIn) sequences with AI-assisted step generation,** conditional follow-ups, and AI reply handling. Automation is AI-forward; depth of manual branching is lighter than Smartlead's API-driven flows.

**The core differentiator is that AI writes and personalizes each email per prospect in 21+ languages,** with personalization credits metered per plan. Strongest AI-writing story of the five for at-scale, per-lead copy.

## COST

**Pro from \$48/mo (5,000 emails, 1 user, unlimited premium warmup, 300 validation + 300 personalization credits/mo).** Growth from \$96/mo (50k emails, 10k active contacts, unlimited users). 14-day trial. Agent Frank AI-SDR is a \$499/mo add-on.

---

Pricing and limits read from [Salesforge's own site](#) . Verified to the edition date above.

## ECOSYSTEM & EXIT

**API, webhooks, and Zapier plus common CRM connectors;** sister tools (Warmforge, Mailforge, Primeforge) cover warmup, domains, and infrastructure. Ecosystem is younger/smaller than Smartlead's but coherent.

**Low-to-moderate.** Data exports via CSV/API and mailboxes are BYO. Leaving means giving up AI-written assets and Warmforge reputation. Being newer, less third-party tooling assumes it, which cuts both ways.

## WHERE IT STANDS OUT

**AI-native workflow (write → send → warm → reply) with unlimited mailboxes/LinkedIn at a founder price,** plus a clean upgrade to an autonomous AI SDR (Agent Frank) when you're ready to hand off prospecting.

## LIMITATIONS

**Newer vendor with a shorter deliverability/reliability track record.** AI copy still needs human review to avoid generic output, but Agent Frank (\$499/mo) is pricey. Smaller community and fewer playbooks than Instantly/Smartlead.

# Woodpecker

## BUILT FOR

Small agencies and B2B SMBs, non-technical sales users running cold email and follow-ups with a focus on deliverability and reply detection.

---

## CHOOSE THIS IF YOU ARE

Quality-over-quantity B2B founders and agencies who value deliverability, compliance, and unlimited inboxes/seats.

## BEST WHEN

You run tightly-targeted, permission-based outreach and want native warmup, inbox rotation, HubSpot/Pipedrive two-way sync, and an Agency Panel out of the box.

## AVOID IF

You blast high volumes across hundreds of inboxes on a budget; the per-contacted-prospect pricing punishes scale.

## WHAT IT DOES

SaaS, EU-based.

**Unlimited email accounts and unlimited team seats included on all plans.** You pay for contacted prospects, not inboxes. Extra mailboxes \$4-6/mo, extra warmups \$5/mo/account, dedicated servers (Infraforge) \$59/mo.

**Free built-in warmup (4 warmups included, more \$5 each),** inbox rotation, adaptive sending, and centralized inbox. Domain audit (SPF/DKIM), spam-word/link checker, and customizable sending limits/intervals ship natively.

**Multi-step conditional sequences with automatic follow-ups that stop on reply,** A/B testing, snippet/condition logic, and multichannel steps (email + LinkedIn, LinkedIn is \$29/mo per connected account). Solid but less "unlimited-scale" than Instantly/Smartlead.

**Deep custom-field/snippet personalization and conditional content;** AI assists copy and an Adaptive Sending algorithm optimizes timing. Not as AI-heavy as Salesforce/Smartlead's fully AI-generated personalization at scale.

## COST

**Priced per contacted prospect (~\$7 per 100/mo monthly; ~33% off annual).** Entry plans ~\$24-35/mo. 14-day free trial or 100 cold emails, no card. API/integrations/webhooks/MCP/CLI are a \$20/mo add-on.

---

Pricing and limits read from [Woodpecker's own site](#) . Verified to the edition date above.

## ECOSYSTEM & EXIT

**Native two-way sync with HubSpot and Pipedrive**, but read/write with Zapier, Google Sheets, Calendly, Clay, Hunter, Zoho. Full API/webhooks/MCP/CLI gated behind the \$20/mo add-on, not included.

**Easy in/out: CSV prospect import/export and standard sequence structures.** No proprietary lock-in on data; main friction is rebuilding warmup reputation and re-wiring native CRM syncs elsewhere.

## WHERE IT STANDS OUT

**Deliverability-first reputation, unlimited inboxes/seats**, mature EU/GDPR-friendly compliance stance, and an Agency Panel (\$27/mo per active client, \$5/mo white-label) with per-client databases, bulk actions, and automated audits.

## LIMITATIONS

**Per-contacted-prospect model gets pricey and unpredictable at high volume vs Instantly/Smartlead's flat unlimited sending.** Add-ons (API \$20, LinkedIn \$29, warmups, mailboxes) stack up, but fewer built-in leads and less aggressive AI than newer rivals.



# Build

## BUILT FOR

Roll your own sender: Amazon SES/Postmark/Resend or a warmed Google Workspace, plus a script that handles cadence, personalization, and reply detection. Raw sending power, none of the cold-outreach guardrails.

Technical founders sending low-volume, highly targeted cold email at MVP stage who want to control deliverability and cost, and can write and babysit the sending script themselves.

---

## CHOOSE THIS IF YOU ARE

Highly technical founders sending low volume to a tightly curated list who want total control and near-zero per-email cost.

### BEST WHEN

You script sending through SES (~\$0.10 per 1k emails) or Postmark and manage your own list.

### AVOID IF

You need warmup, inbox rotation, deliverability tooling, reply detection, or bounce/reputation management; you'll get flagged and burn domains, and Postmark/SES both forbid pure cold spam.

---

## WHAT IT DOES

**No outreach SaaS in the loop.**

**SES scales to huge volume technically**, but cold-emailing at volume torches your domain reputation. Safe cold-outreach means several warmed mailboxes/domains you provision and rotate manually, since there is no built-in inbox rotation.

**The hard part, all on you: SPF/DKIM/DMARC setup**, dedicated IP or subdomain warming, and no automated warmup pool. One misstep and you land in spam or get your main domain blacklisted. Zero safety net.

**Fully DIY: you code the cadence**, delays, follow-up steps, and reply-detection (IMAP/webhook parsing) to stop threads. Flexible and free, but every branch and edge case is yours to build and debug.

**The strong upside is that you can call an LLM per lead for tailored openers/merge-fields at fractions of a cent.** Token efficiency is a skill. Sloppy prompts burn credits fast. No native spintax or spam-check UI to lean on.

---

## COST

**Cheap per-email: SES ~\$0.10/1k,** Resend/Postmark have small free tiers. No per-seat outreach fee. Real cost is setup time plus LLM token burn for personalization. The spend is mostly one-time, so it wins long-term only with steady use.

---

## ECOSYSTEM & EXIT

**None out of the box.** You write the glue to pull leads and push replies/opens back to your CRM or sheet. Doable via APIs/webhooks but it's code to build and maintain, not a prebuilt connector.

**There is nothing to leave.** You own the sending infra, lists, and templates, portable anytime. Trade-off: also nothing to fall back on, and rebuilding warmed domains elsewhere still takes weeks.

---

## WHERE IT STANDS OUT

**Lowest possible send cost, total control of deliverability levers,** and AI personalization limited only by your prompt skill. Ideal for a founder doing surgical, low-volume outreach who wants to own the pipe.

---

## LIMITATIONS

**No warmup pool, spam-testing,** inbox rotation, or unified inbox. You build and own deliverability risk, which can nuke your primary domain. Breaks past small volume, unforgiving without email-infra know-how.

---

No vendor page: this route is assembled from the components named above.

# Newsletter / Marketing Campaigns

Compared: beehiiv · Kit · Ghost · MailerLite · Substack · Mailchimp · Build

# At a glance

| TOOL              | CHOOSE THIS IF YOU ARE                                                                                                                                 | STARTING PRICE                                                                                         | MAIN BENEFIT                                                                                                                                                                                                                                                                                                                 |
|-------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>beehiiv</b>    | <b>Founders building a media/newsletter brand who want growth and monetization tooling</b><br>non-technical, US/global, scaling a list fast.           | <b>Launch is free up to 2,500 subscribers</b>                                                          | <b>Newsletter growth paired with native monetization.</b><br>The Boosts ad network and paid subscriptions with a 0% beehiiv take on subscription revenue (only Stripe's 2.9% + \$0.30 applies). For a writer whose goal is to grow a list and earn from it directly, that combination on one platform is the differentiator. |
| <b>Kit</b>        | <b>Creators and solo founders monetizing an audience via automations and digital products</b><br>non-technical, US/global, list-building focused.      | <b>Free up to 10,000 subscribers</b><br>(single automation), one of the larger free ceilings here      | <b>The largest free subscriber ceiling here, up to 10,000 subscribers free, combined with automation-first sequences.</b><br>For a creator who wants to build and nurture a sizable audience before paying, and to run branching automations once they do, that free-tier depth plus workflow tooling is the standout.       |
| <b>Ghost</b>      | <b>Founders who want an owned publication plus membership site with 0% revenue share</b><br>semi-technical, global, building a blog/newsletter brand.  | <b>Self-host free</b><br>(realistic self-hosted infra \$39-55/mo)                                      | <b>A creator-owned full publishing stack (site + newsletter + memberships) that takes 0% of member revenue, self-hostable for full control.</b>                                                                                                                                                                              |
| <b>MailerLite</b> | <b>Budget-conscious SMB and founder marketers who want clean, cheap email plus automation</b><br>non-technical, global, ecommerce or simple campaigns. | <b>Free: 250 subscribers</b><br>2,500 emails/mo (cut from 500/12k on 1 July 2026), MailerLite branding | <b>A drag-and-drop editor plus website, landing-page and digital-product tools bundled in, EU data hosting, and pricing below Mailchimp and Kit for comparable features.</b>                                                                                                                                                 |
| <b>Substack</b>   | <b>Writers and solo founders who want zero-setup paid newsletters plus built-in discovery</b><br>non-technical, US/global, audience-first.             | <b>Free to start with no monthly platform fee and unlimited subscribers/sends</b>                      | <b>Zero fixed cost, dead-simple setup, and a built-in discovery network (Recommendations, Notes, the app) that actively grows your subscriber base, something no traditional ESP offers.</b><br>Best for audience-first, content-led founders.                                                                               |
| <b>Mailchimp</b>  | <b>SMBs and non-technical founders wanting all-in-one marketing (email + ecommerce + ads) under a familiar brand</b><br>global, B2C/ecommerce.         | <b>Free tier cut in Jan 2026 to just 250 contacts and 500 sends/month with no automation</b>           | <b>Breadth: mature automation, deep segmentation, a wide integration ecosystem, and an all-in-one marketing toolkit.</b><br>A versatile choice when a founder needs real campaign automation and e-commerce ties, not just a newsletter.                                                                                     |
| <b>Build</b>      | <b>Highly technical founders wanting full ownership and near-zero send cost at scale</b><br>self-managed infra, any geography.                         | <b>Software is free</b>                                                                                | <b>Own your list outright, escape per-subscriber pricing, and (with Ghost) get a strong editor plus native paid subscriptions, a Substack alternative for a founder willing to self-host.</b>                                                                                                                                |

Full detail for every tool follows on its own page.

# beehiiv

## BUILT FOR

Newsletter creators and media startups, growth-focused writers publishing and monetizing email newsletters with built-in ads, referrals, and audience analytics. Best fit: creators treating the newsletter as a business who want growth and monetization tools without stitching together separate services.

---

## CHOOSE THIS IF YOU ARE

Founders building a media/newsletter brand who want growth and monetization tooling: non-technical, US/global, scaling a list fast.

### BEST WHEN

You want built-in referral programs, the Boosts ad network, paid subscriptions, and cross-recommendations to grow, and you'd rather pay a flat fee than share revenue. Launch is free to 2,500 subscribers. Scale (from \$43/mo annual) unlocks monetization, automations and Boosts (custom domains come on every plan). Max (from \$96/mo annual) adds team and multi-publication features. Enterprise is custom for 100,000+ subscribers.

### AVOID IF

You're a solo writer who just wants to publish and charge readers with minimal overhead, or you'll balk at the 2,500-sub free cliff.

## WHAT IT DOES

**Alongside the newsletter itself it runs the Boosts ad network and paid subscriptions**, and takes a 0% beehiiv cut of paid-subscription revenue (only Stripe's 2.9% + \$0.30 applies). Monetization and automations unlock on the Scale plan.

**Built-in monetization is a core part of the product:** the Boosts ad network and paid subscriptions, with a 0% beehiiv take on paid-subscription revenue (only Stripe's 2.9% + \$0.30 processing applies). Both levers unlock on the Scale plan (\$43/mo, annual) instead of the free tier.

**Growth centers on the Boosts ad network,**

which connects newsletters to acquire subscribers, alongside the newsletter site (custom domains come on every plan, including free Launch). Boosts sits with the monetization tooling on the Scale plan (\$43/mo, annual), so the same tier that unlocks revenue also unlocks the paid growth lever.

**Email automations are part of the platform but gated to the Scale plan (\$43/mo, annual) and above.** The free Launch tier (up to 2,500 subscribers) does not include them. Automation is therefore tied to the same paid tier as monetization, so upgrading unlocks both the revenue and the workflow tooling together.

**Deliverability is managed by beehiiv as part of its hosted newsletter sending.** The sources used here discuss it qualitatively, without a delivery-rate figure. As a hosted, web-based platform, sending reputation and infrastructure sit with beehiiv, not the writer.

**A web-based newsletter editor tied to its newsletter site and branding customization.**

The sources used here describe it as newsletter-focused and easy to set up. Composition is oriented to writing and sending issues instead of product or transactional templates.

---

**COST**

**Launch is free up to 2,500 subscribers.** Scale is \$43/mo (annual) and is the plan that unlocks monetization (Boosts ads, paid subscriptions) and automations. Custom domains are included on every plan, and branding removal comes only with Max or Enterprise. Max is \$96/mo, and Enterprise is custom for 100,000+ subscribers.

---

**ECOSYSTEM & EXIT**

**The documented integrations are the Boosts ad network and Stripe.** Stripe underpins paid subscriptions (with beehiiv's 0% take on that revenue) and Boosts connects newsletters for growth. Integration is oriented to monetization and audience growth, not a broad SaaS-data ecosystem.

**Subscriber export is supported,** so the list itself is portable. As a hosted newsletter platform, moving away means recreating the newsletter site, automations and monetization setup elsewhere, but the core asset, the subscriber list, can be exported out.

---

**WHERE IT STANDS OUT**

**Newsletter growth paired with native monetization.** The Boosts ad network and paid subscriptions with a 0% beehiiv take on subscription revenue (only Stripe's 2.9% + \$0.30 applies). For a writer whose goal is to grow a list and earn from it directly, that combination on one platform is the differentiator.

## LIMITATIONS

**Monetization and automations are gated to the Scale plan (\$43/mo, annual).** The free Launch tier (up to 2,500 subscribers) has neither. It is also newsletter-first, not a product- or transactional-email tool, so a SaaS team needing lifecycle or system mail would look elsewhere.

---

Pricing and limits read from [beehiiv's own site](#) . Verified to the edition date above.

# Kit

## BUILT FOR

Professional creators and course sellers, non-technical solo operators running email marketing with automations, subscriber tagging, and simple commerce. Best fit: creators selling digital products who need behavioral automations and paid recommendations beyond basic broadcasting.

---

## CHOOSE THIS IF YOU ARE

Creators and solo founders monetizing an audience via automations and digital products: non-technical, US/global, list-building focused.

## BEST WHEN

You want visual automations, tagging, a free tier up to 10k subscribers, and creator commerce (products, referral system).

## AVOID IF

You're highly price-sensitive at small list sizes (the 2025 hike pushed Creator to a ~\$33+ minimum) or want an owned blog/site.

## WHAT IT DOES

Formerly ConvertKit, it is a mature creator platform aimed at audience-building over publishing.

Kit centers on audience-building (a free tier scaling to 10,000 subscribers, visual automations, forms and landing pages) over native creator monetization. The sources used here do not document built-in payment or paid-subscription tooling.

### Audience-growth tooling, per the sources

**used here**, is forms and landing pages, backed by a large creator-tool ecosystem for integrations. These support list capture instead of viral referral loops. Combined with the 10,000-subscriber free tier, the emphasis is on collecting and nurturing an audience over paid acquisition.

### Automation is Kit's core strength per the

**sources used here**: visual automations and sequences, described as automation-first. This is the primary reason to choose it over a plain broadcast tool: building branching sequences and audience workflows. The sources used here emphasize automation more than detailed segmentation controls.



**Deliverability rests on Kit's managed sending as a mature creator platform.** The sources used here treat it qualitatively, without a delivery-rate metric. Sending infrastructure and reputation are handled by the hosted, web-based service, not configured by the creator.

**The customization detail in the sources used here covers forms and landing pages,** not the compose editor. Kit is characterized as automation-first. The writing experience is framed around broadcasts and sequences. It is a web tool aimed at creator email, not long-form publishing.

---

## COST

**Free up to 10,000 subscribers (single automation),** one of the larger free ceilings here. Paid Creator pricing is subscriber-based: \$39/mo at 1,000 contacts, \$59 at 3,000, \$89 at 5,000, and ~\$119 at 10,000 (monthly rates; annual ~16% less). Cost climbs with list size once you pass the free tier.

---

## ECOSYSTEM & EXIT

**A large creator-tool ecosystem is the documented integration strength,** alongside forms and landing pages for capture. The breadth is oriented to the creator stack, connecting Kit to the other tools an audience-builder uses, more than to product or SaaS data pipelines.

**Export is supported, so the subscriber list is portable.** Because Kit's value is concentrated in its visual automations and sequences, the practical switching cost is rebuilding those workflows in another tool instead of moving the audience itself, which can be exported out.

---

## WHERE IT STANDS OUT

**The largest free subscriber ceiling here,** up to 10,000 subscribers free, combined with automation-first sequences. For a creator who wants to build and nurture a sizable audience before paying, and to run branching automations once they do, that free-tier depth plus workflow tooling is the standout.

---

## LIMITATIONS

**Creator pricing climbs with list size (\$39/mo at 1,000 contacts, \$89 at 5,000, ~\$119 at 10,000),** so costs rise as the audience grows past the free tier. It is also a marketing-email tool. The documented feature set does not cover transactional or product email, limiting it to audience marketing.

---

Pricing and limits read from **Kit's own site** . Verified to the edition date above.

# Ghost

## BUILT FOR

Independent publishers and developers, technically comfortable writers running a website plus membership newsletter with paid subscriptions on open-source software. Best fit: creators wanting to own their full publishing stack and site, and willing to self-host or pay for managed hosting.

---

## CHOOSE THIS IF YOU ARE

Founders who want an owned publication plus membership site with 0% revenue share: semi-technical, global, building a blog/newsletter brand.

## BEST WHEN

You want a full website + newsletter + paid memberships you fully own and to keep 100% of subscription revenue (only Stripe fees). Starter \$18 (free members only), Publisher \$29 (paid memberships), Business \$199.

## AVOID IF

You want a built-in discovery/growth network, or you're non-technical and Starter's default-theme and free-only limits frustrate you.

## WHAT IT DOES

**Native paid memberships/subscriptions with a 0% Ghost cut (only Stripe's processing fees apply)**, but paid memberships require the Publisher plan.

**A full SEO-friendly website, membership signups**, and portal, with no viral referral loops.

**Newsletter and membership automations with member segmentation (tiers/labels)**. Lighter than dedicated marketing-automation tools.

**Managed sending reputation on Ghost(Pro) (email via bundled infrastructure)**. Self-hosters configure their own mail provider.

**A polished, distraction-free publishing editor**, one of Ghost's strengths.

---

## COST

**Self-host free (realistic self-hosted infra \$39-55/mo).** Ghost(Pro): Starter \$18/mo (annual; 1 staff user, no paid memberships), Publisher \$29/mo scaling with members (\$63 at 5,000, \$88 at 10,000, \$274 at 100,000), Business \$199/mo; 14-day free trial, no card.

---

Pricing and limits read from [Ghost's own site](#) . Verified to the edition date above.

---

## ECOSYSTEM & EXIT

**Native integrations plus Zapier**, and Stripe for memberships; open API.

**Open-source and highly portable.** You own your site, content and member list, with standard export.

---

## WHERE IT STANDS OUT

**A creator-owned full publishing stack (site + newsletter + memberships) that takes 0% of member revenue**, self-hostable for full control.

---

## LIMITATIONS

**Paid memberships require Publisher (\$29/mo+, scaling with members).** Lighter marketing automation than Kit, and it is not a transactional-email tool.

# MailerLite

## BUILT FOR

Small businesses and solo creators, non-technical marketers sending newsletters and automations with landing pages on a generous free tier.

---

## CHOOSE THIS IF YOU ARE

Budget-conscious SMB and founder marketers who want clean, cheap email plus automation: non-technical, global, ecommerce or simple campaigns.

## BEST WHEN

You want a low price-per-subscriber with a drag-and-drop builder, automations, and a free plan (now 250 subs / 2,500 sends).

## AVOID IF

You need deep creator commerce, a media-growth network, or an owned publishing site; Note the 2026 plan renames (now Comfort and Power) and 10-30% price bump.

## WHAT IT DOES

**Sell digital products, subscriptions,** and paid newsletters natively via Stripe; e-commerce integrations (Shopify, WooCommerce) for product-based automations. No transaction fee beyond Stripe's.

**Signup forms/pop-ups, landing pages,** and a basic website builder. The free tier caps these (1 landing page, 1 site, 3 forms), Comfort caps them at 10, and only Power removes the limits. Built-in A/B testing and referral-style features on paid tiers.

**Visual automation builder with triggers,** conditions, and multi-step flows (Free: 3 flows/5 steps; Comfort: 50; Power: unlimited/100 steps). Unlimited tags and dynamic segments on all tiers, though logic is simpler than enterprise ESPs.

**Strong reputation aided by a strict manual account-approval/anti-spam review that keeps sender pools clean,** but EU data storage, ISO 27001, GDPR. Optional dedicated IP add-on on Power/Enterprise.

**Widely praised drag-and-drop editor plus a distraction-free rich-text "Simple" editor and full HTML;** AI writing assistant on Advanced/Power. Clean, modern, low learning curve.

---

## COST

**Free: 250 subscribers, 2,500 emails/mo (cut from 500/12k on 1 July 2026)**, MailerLite branding. Paid Comfort from ~\$12/mo, Power from ~\$25/mo (unlimited emails); ~10% off annual. Scales with subscriber count.

Pricing and limits read from **MailerLite's own site** . Verified to the edition date above.

---

## ECOSYSTEM & EXIT

**100+ native integrations (Stripe, Shopify, WooCommerce, Zapier, plus ChatGPT/Claude) and full REST API**, webhooks, and MCP access on paid plans. Free plan has limited/no API sending.

**Contacts export cleanly to CSV**, but segments, automation flows, and template designs don't export tidily; capture them manually first. Over-limit free accounts enter a restricted state that locks editing until you reduce usage or upgrade.

---

## WHERE IT STANDS OUT

**A drag-and-drop editor plus website**, landing-page and digital-product tools bundled in, EU data hosting, and pricing below Mailchimp and Kit for comparable features.

---

## LIMITATIONS

**Repeated free-tier cuts (now 250 subs) erode its old value edge**. Segmentation/reporting shallower than enterprise tools. Strict approval can delay/reject some accounts, branding on free tier.

# Substack

## BUILT FOR

Individual writers and journalists, non-technical creators publishing free and paid newsletters with near-zero setup and built-in discovery. Best fit: writers who want the fastest path to a paid newsletter and audience network, accepting a revenue cut and limited customization.

---

## CHOOSE THIS IF YOU ARE

Writers and solo founders who want zero-setup paid newsletters plus built-in discovery: non-technical, US/global, audience-first.

## BEST WHEN

You want to start free, publish instantly, and tap network effects (recommendations, Notes, discovery) with no monthly fee.

## AVOID IF

Your list is large and paid (10% of revenue dwarfs a flat SaaS fee) or you want to own your platform, data, and design.

## WHAT IT DOES

**Writers publish posts that go out by email and live on a public archive**, with optional podcast and community/chat features.

**Native paid subscriptions (monthly/annual)**, plus founding-member tiers, group/gift subscriptions and paywalled posts. Substack takes a flat 10% of subscription revenue. Stripe adds ~2.9%+\$0.30, so real take is ~13% of gross earnings.

**Built-in network effects: Recommendations (other writers refer readers)**, the Substack app/Notes feed, Explore discovery, and easy referral sharing. This organic discovery engine is a genuine differentiator versus standalone email tools.

**Weak.** Only a basic automated welcome email, but no drip sequences, no behavioral triggers, minimal list segmentation. You broadcast to everyone or to the free/paid split. It is not a marketing-automation platform.

**Generally solid on a managed shared infrastructure with no setup required**; handles sending reputation. But you cannot use your own sending domain/IP or deeply tune deliverability the way dedicated ESPs allow.

**A clean, distraction-free WYSIWYG editor praised for simplicity**: embeds, images, buttons, paywall markers and podcast/audio. Great for writing, but limited design/branding control and no true drag-and-drop template builder.

## COST

**Free to start with no monthly platform fee and unlimited subscribers/sends.** Substack makes money only via its revenue cut when you charge subscribers, so a free newsletter of any size costs nothing to run.

Pricing and limits read from [Substack's own site](#) . Verified to the edition date above.

## ECOSYSTEM & EXIT

**Minimal by design.** Stripe (payments) is built in; otherwise few native integrations, no deep CRM/e-commerce connectors and only a limited API. Zapier and CSV export cover basic data movement.

**Easy to leave.** You can export subscriber emails (including paid-status) and post content as CSV/HTML at any time. Substack doesn't lock your list. Paid subscriptions migrate via Stripe, though re-subscribing friction exists.

## WHERE IT STANDS OUT

**Zero fixed cost, dead-simple setup,** and a built-in discovery network (Recommendations, Notes, the app) that actively grows your subscriber base, something no traditional ESP offers. Best for audience-first, content-led founders.

## LIMITATIONS

**The flat 10% cut gets expensive at scale,** and it lacks automation, segmentation, custom domains for sending, deep analytics and integrations. Not a fit for transactional email or sophisticated multi-step marketing campaigns.

# Mailchimp

## BUILT FOR

Small-to-mid businesses and e-commerce sellers, general marketers running email campaigns, automations, and audience management with broad integrations. Best fit: established SMBs wanting a general-purpose marketing platform with commerce integrations, instead of a creator- or newsletter-first tool.

---

## CHOOSE THIS IF YOU ARE

SMBs and non-technical founders wanting all-in-one marketing (email + ecommerce + ads) under a familiar brand: global, B2C/ecommerce.

## BEST WHEN

You need broad integrations, landing pages, and multichannel marketing on a single well-known platform and want CRM, ads, and ecommerce under one roof instead of a pure email tool. Essentials from \$13, Standard from \$20 (500 contacts).

## AVOID IF

You're cost-sensitive or a creator; the Jan 2026 free-tier gutting (250 contacts/500 sends) and contact-tiered pricing (Standard ~\$100 at 5k) make it pricey.

---

## WHAT IT DOES

**Not a creator-monetization tool.** Mailchimp is a paid SaaS you pay by contacts/features. It has no native paid-subscription/paywall model. Revenue comes from your own campaigns (e-commerce, promotions), not subscriber billing through the platform.

**Signup forms, hosted landing pages,** pop-ups, social-post scheduling and basic ads, plus a customer-journey builder. Solid acquisition toolkit, but no discovery/referral network like Substack, but you drive traffic yourself.



**Strong and a core strength: multi-step customer journeys**, behavioral triggers, and detailed/predictive segmentation with AI. But automation is tiered: the free plan has no flows, Essentials (\$13/mo) caps multi-step automation at 4 flow steps, and full journeys (up to 200 flows) need Standard+ (\$20/mo).

**Established sender infrastructure with a generally good reputation**, plus DKIM/SPF authentication and a dedicated-IP option on Premium. Reliable, though some large senders report deliverability slipping versus specialist ESPs.

**A mature drag-and-drop email builder with many templates**, a content studio and AI copy/creative assistant. More design/branding control than Substack, but the interface is heavier and less focused on long-form writing.

---

## COST

**Free tier cut in Jan 2026 to just 250 contacts and 500 sends/month with no automation.**

Essentials from \$13/mo, Standard from \$20/mo, Premium from \$350/mo, all priced by contact count, so cost scales steeply as lists grow.

---

## ECOSYSTEM & EXIT

**Extensive: 300+ integrations including Shopify**, WooCommerce, Salesforce, Canva, and social platforms, plus a full-featured API and Zapier. One of the widest integration ecosystems in email marketing, an advantage for connected stacks.

**Contacts and campaign data export easily via CSV/API**, meaning your list isn't locked. But rebuilding automations, templates, segments and integrations elsewhere is real work, creating practical switching friction as your setup deepens.

---

## WHERE IT STANDS OUT

**Breadth: mature automation, deep segmentation**, a wide integration ecosystem, and an all-in-one marketing toolkit. A versatile choice when a founder needs real campaign automation and e-commerce ties, not just a newsletter.

---

## LIMITATIONS

**Costly as lists grow.** It bills on total contacts including unsubscribed, and gates key features (automation, dedicated IP) behind higher tiers. The gutted 2026 free tier, SMS/transactional add-ons and complexity push real costs 20-40% above list price.

---

Pricing and limits read from **Mailchimp's own site** . Verified to the edition date above.

# Build

## BUILT FOR

Self-hosted broadcasting. You own the list and the sending. You also own the ops, deliverability, and uptime. Technical founders and small teams at MVP scale who want to own their audience list and dodge per-subscriber pricing, and are comfortable running a server and DNS/email config themselves.

---

## CHOOSE THIS IF YOU ARE

Highly technical founders wanting full ownership and near-zero send cost at scale: self-managed infra, any geography.

## AVOID IF

You don't want to own deliverability, IP warmup, bounce/complaint handling, uptime, and DNS auth (SPF/DKIM/DMARC) yourself; one misconfig tanks inbox rates.

---

## WHAT IT DOES

**Your server, your subscriber DB.**

**Ghost ships native paid memberships/subscriptions (Stripe).** Listmonk has none. You'd bolt on your own paywall. Neither offers ad networks or referral marketplaces the way hosted platforms bundle.

**Thin.** Ghost has signup forms and basic SEO. Listmonk is bare-bones. No built-in recommendation network, boosts, or discovery feed. The growth flywheels hosted platforms give you are absent unless you engineer them.

**Listmonk offers SQL-query segments and**

**basic automations.** Ghost has simpler tiers/labels. Capable for MVP but not a full journey builder. Complex drip logic means scripting against the API, not clicking a flow UI.

**Rides on your SES setup and sender**

**reputation, all on you.** Correct SPF/DKIM/DMARC and list hygiene get it there, but there's no shared-IP reputation or deliverability team as backstop. Misconfig = spam folder.

**Ghost's editor is excellent: clean**, modern, near-Substack quality. Listmonk's is basic HTML/rich-text, functional but plain. Neither matches the polished template galleries of hosted marketing suites.

---

## COST

**Software is free.** You pay only for a VPS (~\$5-20/mo) and SES (~\$0.10/1k emails). No per-subscriber tax, which is huge at scale. Cost is setup time plus hosting. Wins over per-contact SaaS only if the list and usage persist.

---

## ECOSYSTEM & EXIT

**API/webhook-driven; Ghost has an integrations directory (Zapier, Slack, analytics).** Listmonk is API-only. No sprawling marketplace. Most connections are ones you wire up and maintain yourself.

**Excellent: export subscribers to CSV and content to Markdown/JSON anytime.** It's your DB. Getting in is easy too. The snag is you must stand up and maintain the infra, not just flip an import switch.

---

## WHERE IT STANDS OUT

**Own your list outright, escape per-subscriber pricing**, and (with Ghost) get a strong editor plus native paid subscriptions, a Substack alternative for a founder willing to self-host.

---

## LIMITATIONS

**You own uptime, security patches, and deliverability.** No growth/discovery network. Segmentation and automation are limited vs hosted suites. Fine for MVP, but scaling and ops overhead grow with the list.

---

No vendor page: this route is assembled from the components named above.

# Social Media Scheduling

Compared: Buffer · Typefully · Later · Publer · Metricool · Hypefury · Build

# At a glance

| TOOL      | CHOOSE THIS IF YOU ARE                                                                                                                                     | STARTING PRICE                                                                                                                                | MAIN BENEFIT                                                                                                                                                                                                                                                                              |
|-----------|------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Buffer    | Solo founders and small teams wanting dead-simple per-channel cross-platform scheduling                                                                    | Free covers 3 channels with 10 scheduled posts each<br>refillable as posts publish, and includes API access and basic analytics on every plan | Its free plan is the most generous here, with 3 channels, 10 scheduled posts each (refillable), and API access plus basic analytics on every plan, combined with the simplest setup.<br>Buffer's edge is low-friction scheduling at minimal or no cost instead of depth.                  |
| Typefully | X and LinkedIn-first founders building in public through threads<br>long-form writing                                                                      | Free covers one social set and 15 posts a month<br>enough to trial the editor, not run at volume                                              | A writing-first workflow for X and LinkedIn, plus Threads, Bluesky and Mastodon, built around a distraction-free thread editor with AI.<br>It stands out for composing text content, not visual planning or bulk operations.                                                              |
| Later     | Visual-first brands and creators centered on Instagram<br>TikTok, and Pinterest                                                                            | No free plan (Hypefury also lacks one)<br>and it sits at the higher end at entry among these schedulers                                       | Visual planning built on a grid preview and link-in-bio tools, centered on Instagram and TikTok.<br>Its differentiator is feed-level visual curation, seeing how posts sit together before publishing, instead of volume or analytics breadth.                                            |
| Publer    | Budget-conscious founders wanting the most features per dollar across many networks                                                                        | A free plan exists                                                                                                                            | Bulk and batch scheduling of up to 500 posts via CSV, with content recycling, at low per-account pricing from \$5/mo. Its standout is high-volume queue-building across multiple accounts, not composing or analytics depth.                                                              |
| Metricool | Founders and small agencies who treat analytics and paid ads as first-class<br>not an afterthought                                                         | Free covers one brand and 20 published posts a month but excludes LinkedIn and X                                                              | Analytics that several competitors charge for, including engagement, optimal times, demographics and hashtags, included even on the free tier; competitor tracking reaches 100 profiles on paid plans (5 on free).<br>Reporting depth bundled with scheduling is what sets it apart here. |
| Hypefury  | Solo creators on Threads, Bluesky<br>LinkedIn, Instagram or TikTok optimizing hard for audience and monetization growth (X/Twitter is no longer supported) | No free plan; 7-day trial<br>(no card)                                                                                                        | A purpose-built founder growth loop: auto-plugs on viral posts, auto-DM lead capture (100-300/day by tier), an engagement builder for targeted replies, and evergreen recycling.<br>Posting feeds audience growth.                                                                        |
| Build     | Technical founders who want zero subscription cost<br>full data control                                                                                    | No SaaS fee, but not free<br>~\$0-6/mo hosting plus API costs                                                                                 | Total control, no per-seat fees, and unlimited custom logic (conditional posts, cross-repost, AI-drafted queues).<br>If posting is part of a bigger automation you already run, folding it in is nearly free.                                                                             |

Full detail for every tool follows on its own page.

# Buffer

## BUILT FOR

Solo creators, small businesses, and lean marketing teams who want a simple, reliable queue across Instagram, LinkedIn, X, Facebook, and TikTok without a steep learning curve. Best fit: When you need dependable set-and-forget scheduling and clean basic analytics, not deep engagement automation or heavy multi-client agency workflows.

---

## CHOOSE THIS IF YOU ARE

Solo founders and small teams wanting dead-simple, per-channel cross-platform scheduling.

## BEST WHEN

You post to a mix of LinkedIn, Instagram, TikTok and X and want predictable per-channel pricing. Free covers 3 channels/10 posts each, Essentials runs ~\$5-6/channel/mo with volume discounts down to \$1/channel past 51.

## AVOID IF

Your whole game is X threads and build-in-public growth, or you need agency-grade analytics and a unified inbox; Buffer's reporting is thin and per-channel costs balloon at high channel counts.

## WHAT IT DOES

**Its free tier covers 3 channels with 10 scheduled posts on each**, refillable as posts publish, and includes API access plus basic analytics on every plan, features some tools reserve for paid tiers. Paid plans scale per channel from Essentials upward instead of per seat.

**The major social networks.** The free plan connects 3 channels, with paid plans adding channels one at a time. A channel is each connected profile. The per-channel model applies directly to how many networks you can publish to at once.

**Auto-publishes to the supported major**

**networks on schedule**, not sending manual reminders to post. Scheduling works against the queue. The free tier holds 10 scheduled posts per channel, refillable as they go out. Publishing is hands-off within those limits.

**A simple, fast composer aimed at quick**

**scheduling instead of long-form drafting**, paired with an AI assistant and best-time posting suggestions. It matches Buffer's positioning as the least complex tool here, favoring speed over depth of editing features.

**Queue-based scheduling instead of heavy bulk tooling;** posts flow through a per-channel queue that holds 10 scheduled at a time on the free tier, refillable as they publish. There is no large-batch CSV workflow here. It is built for steady queueing, not high-volume imports.

---

Pricing and limits read from [Buffer's own site](#) . Verified to the edition date above.

**Basic engagement metrics plus best-time-to-post suggestions,** included on every plan, with basic analytics shipping even on the free tier. It covers fundamentals, not deep reporting and is lighter than Metricool's analytics, a known tradeoff for Buffer's simplicity.

---

## COST

**Free covers 3 channels with 10 scheduled posts each,** refillable as posts publish, and includes API access and basic analytics on every plan. Essentials runs \$5/channel/mo (\$6/channel at 1-10 channels on the volume rate). Team is \$10/channel/mo. Costs scale per channel, not per seat.

---

## WHERE IT STANDS OUT

**Its free plan is the most generous here,** with 3 channels, 10 scheduled posts each (refillable), and API access plus basic analytics on every plan, combined with the simplest setup. Buffer's edge is low-friction scheduling at minimal or no cost instead of depth.

---

## LIMITATIONS

**Analytics stay basic next to Metricool's,** and the per-channel model means cost climbs as you add channels: \$5 each on Essentials, \$10 each on Team. A multi-network setup adds up. The free tier's 10-post-per-channel cap also limits volume.

# Typefully

## BUILT FOR

Writers, founders, and X/Twitter-first creators, often solo, who draft threads and long-form posts and value a distraction-free composer. Best fit: When your growth runs through X and LinkedIn text content and you want thread scheduling plus analytics, not a broad multi-network visual planner.

---

## CHOOSE THIS IF YOU ARE

X and LinkedIn-first founders building in public through threads and long-form writing.

## BEST WHEN

Your distribution IS writing: thread composing, tweet scheduling, AI ghostwriting drafts, and X-native analytics. Free gives 15 posts/mo and paid runs from ~\$8/mo (Pro, annual) to ~\$16/mo (Business, annual), bundled into "social sets".

## AVOID IF

You're Instagram/TikTok/visual-led, managing many client brands with reporting needs, or you want monetization-oriented growth automation.

## WHAT IT DOES

Rather than a visual calendar, **it centers a thread and post editor**. The free tier covers one social set and 15 posts a month, positioning it for individual writers before per-social-set paid upgrades.

**Text networks only: X, LinkedIn, Threads, Bluesky and Mastodon.** There is no support for image-led networks such as Instagram or TikTok, which keeps it aligned with written threads and posts but rules it out for anyone whose plan centers on visual content.

**Auto-publishes to its supported text networks, X, LinkedIn, Threads, Bluesky and Mastodon,** so scheduled threads and posts go out without manual intervention. The free tier caps this at 15 posts a month. Higher volume needs a paid social set.

**A distraction-free thread and post editor with AI writing support,** built specifically for composing on text networks. Composition is the tool's center of gravity, not an afterthought to scheduling, which distinguishes it from the visual- and calendar-first tools here.



**Oriented to scheduling individual posts and threads**, without large-batch CSV imports. With the free tier capped at 15 posts a month, the workflow suits a writer publishing deliberately over one loading a big backlog at once.

---

Pricing and limits read from [Typefully's own site](#) . Verified to the edition date above.

**Writing-focused post analytics, covering the performance of individual threads and posts on text networks**, not deep cross-network social analytics. Reporting is scoped to the written content it publishes, consistent with the tool's editor-first focus.

---

## COST

**Free covers one social set and 15 posts a month**, enough to trial the editor, not run at volume. Pro is ~\$8/mo per social set billed annually (~\$10 monthly). Business is ~\$16/mo annual (~\$20 monthly). Enterprise is custom. Pricing scales per social set, not per user.

---

## WHERE IT STANDS OUT

**A writing-first workflow for X and LinkedIn**, plus Threads, Bluesky and Mastodon, built around a distraction-free thread editor with AI. It stands out for composing text content, not visual planning or bulk operations.

---

## LIMITATIONS

**Text networks only (X, LinkedIn, Threads, Bluesky, Mastodon)**, with no Instagram/TikTok or other visual-first tooling, ruling it out for image-led plans. The free tier's 15-posts-a-month cap also limits volume before a paid social set is needed.

# Later

## BUILT FOR

Instagram- and visually-driven brands, e-commerce sellers, and small social teams who plan feeds around images, reels, and link-in-bio. Best fit: When visual grid planning, media library, and link-in-bio matter more than text-heavy scheduling or advanced cross-network reporting.

---

## CHOOSE THIS IF YOU ARE

Visual-first brands and creators centered on Instagram, TikTok, and Pinterest.

### BEST WHEN

You need a drag-and-drop visual content calendar, grid preview, Link-in-bio, best-time-to-post, and a UGC/media library.

### AVOID IF

You're a text-first build-in-public founder, price-sensitive and multi-platform, or you need unified cross-network analytics plus paid-ads reporting.

---

## WHAT IT DOES

**It has no free plan and prices per social set,** sitting at the higher end of these tools at entry.

**Visual networks, centered on Instagram and TikTok.** The grid preview and link-in-bio features are built around these image- and video-led feeds, so its platform coverage is narrower and more specialized than the general-purpose schedulers in this group.

**Auto-publishes to its supported visual networks,** with the grid preview and visual calendar acting as planning aids ahead of publishing. Because coverage centers on Instagram and TikTok, its auto-publishing behavior is tuned to those image- and video-led feeds.

**A visual composer organized around the Instagram grid.** Drafting is tied to how a post will sit in the feed, not to long-form text. Writing tools are secondary to visual placement, matching Later's image-first focus and grid preview.

**Visual calendar planning instead of CSV bulk imports;** scheduling is done by arranging posts on the grid and calendar, not by loading spreadsheets. The emphasis is on how the sequence looks in the feed, not on raw batch volume.

---

Pricing and limits read from [Later's own site](#) . Verified to the edition date above.

**Analytics focused on visual performance for Instagram-first accounts,** tied to the feed and grid it is built around. Reporting is specialized to image- and video-led networks and does not attempt broad multi-network analytics across every platform.

---

## COST

**No free plan (Hypefury also lacks one),** and it sits at the higher end at entry among these schedulers. Pricing is per social set: Starter \$25/mo, Growth \$50/mo, Scale \$110/mo (monthly; ~25% off annual).

---

## WHERE IT STANDS OUT

**Visual planning built on a grid preview and link-in-bio tools,** centered on Instagram and TikTok. Its differentiator is feed-level visual curation, seeing how posts sit together before publishing, instead of volume or analytics breadth.

---

## LIMITATIONS

**No free plan (Hypefury also lacks one),** and it sits at the higher end at entry.

# Publer

## BUILT FOR

Budget-conscious freelancers, agencies, and SMBs, low-to-mid technical, managing many accounts across most major networks with bulk scheduling and workspaces. Best fit: When you want the widest network coverage and power features per dollar, and can trade a busier interface for cheaper multi-account management.

---

## CHOOSE THIS IF YOU ARE

Budget-conscious founders wanting the most features per dollar across many networks.

## AVOID IF

You specifically want creator-growth automation like auto-DMs and evergreen reposts, a polished visual IG planner, or an agency reporting suite with paid-ads tracking; Publer's analytics are functional but unremarkable.

---

## WHAT IT DOES

**It is priced per social account**, not per user, starts low, and supports bulk scheduling of up to 500 posts on its paid plans, a workflow the other tools here handle more lightly.

**The major social networks, with CSV import feeding posts across them in bulk.** Per-account pricing means each connected profile is billed separately, from \$5/mo with +\$4 per extra account on Professional, so coverage scales with how many accounts you attach.

**Auto-publishes on schedule, including large batches loaded via CSV.** Up to 500 posts on paid plans go out without manual reposting. This suits setting up long queues in one pass instead of confirming each post at publish time.

**A composer geared to batch creation**, working alongside CSV import so many posts can be prepared and queued together. It favors throughput, preparing up to 500 scheduled posts on paid plans, over a focused single-post writing environment.

**The clearest fit here for bulk work.** CSV import handles bulk scheduling of up to 500 posts on all paid plans, plus content recycling to repeat evergreen posts. This batch-and-reuse workflow is the core reason to pick Publer over the queue-first tools.

---

Pricing and limits read from [Publer's own site](#) . Verified to the edition date above.

**Scheduling and post analytics that cover the essentials for tracking published output**, but lighter than Metricool's reporting. Sufficient for monitoring bulk-scheduled posts, without the demographic, competitor and hashtag depth Metricool includes.

---

### COST

**A free plan exists.** Paid tiers price per social account. Professional starts at \$5/mo for one account (+\$4 per extra account, +\$2 per member). Business starts at \$10/mo for one account (+\$7 per account, +\$3 per member). Bulk scheduling of up to 500 posts via CSV is on all paid plans.

---

### WHERE IT STANDS OUT

**Bulk and batch scheduling of up to 500 posts via CSV**, with content recycling, at low per-account pricing from \$5/mo. Its standout is high-volume queue-building across multiple accounts, not composing or analytics depth.

---

### LIMITATIONS

**Analytics are lighter than Metricool's**, lacking the same demographic, competitor and hashtag depth. Per-account add-on pricing (+\$4 per account, +\$2 per member on Professional) can also accumulate as you attach more accounts and teammates.

# Metricool

## BUILT FOR

Data-minded marketers, small agencies, and consultants who need scheduling plus genuine cross-channel analytics, including ads and competitor tracking. Best fit: When reporting and measurement drive decisions and you want publishing and analytics in one tool, not a posting-first scheduler with shallow metrics.

---

## CHOOSE THIS IF YOU ARE

Founders and small agencies who treat analytics and paid ads as first-class, not an afterthought.

## BEST WHEN

You want scheduling PLUS unified analytics, competitor benchmarking, and Google/Meta/TikTok ads management in one dashboard. Free covers 1 brand, Starter ~\$20-36/mo (up to 10 brands), Advanced from ~\$53/mo (up to 50 brands).

## AVOID IF

You just need a simple posting queue, you're purely X-growth focused, or you're a solo creator who'll never open the analytics; you'd overpay for unused reporting.

---

## WHAT IT DOES

**It bundles reporting that several competitors gate behind paid tiers**, including engagement, optimal posting times, demographics and hashtags, available even on its free plan, which covers one brand and 20 published posts a month.

### **The major networks, but with caveats.**

LinkedIn and X are excluded on the free tier, and X/Twitter is a paid add-on at \$10/mo per account since 13 July 2026 (accounts added before then keep a grandfathered \$5/mo). Full coverage of those two networks therefore requires a paid plan plus, for X, the extra add-on.

### **Auto-publishes to connected networks on**

**schedule.** Note the connection caveats: the free tier excludes LinkedIn and X, and X publishing needs the \$10/mo add-on (grandfathered \$5/mo before 13 July 2026), so auto-publishing to those two networks depends on a paid plan and, for X, the extra.

### **A composer that sits next to Metricool's**

**reporting.** Posts can be drafted with analytics context such as optimal posting times at hand. Composition is functional and paired with the tool's broader analytics, but it is not a standalone writing focus.

**Bulk planning organized per brand.** The free tier caps publishing at 20 posts a month for one brand; paid tiers raise that limit and add brands: Starter covers up to 10 brands (\$20-36/mo) and Advanced up to 50 (from \$53/mo). There is no dedicated large-CSV import, so it suits multi-brand planning over single high-volume loads.

**The deepest analytics here, available even on the free tier:** engagement, optimal posting times, audience demographics and hashtag performance. Competitor tracking runs to 5 profiles on the free tier (30 days of analytics) and up to 100 on paid plans. This reporting, which several tools gate behind paid plans, is Metricool's main draw.

---

## COST

**Free covers one brand and 20 published posts a month but excludes LinkedIn and X.** Starter runs ~\$20-36/mo for up to 10 brands, but Advanced from ~\$53/mo for up to 50 brands (annual billing ~24% less). X/Twitter is a \$10/mo add-on per account (grandfathered \$5/mo for accounts added before 13 July 2026). Pricing scales per brand.

---

## WHERE IT STANDS OUT

**Analytics that several competitors charge for, including engagement,** optimal times, demographics and hashtags, included even on the free tier; competitor tracking reaches 100 profiles on paid plans (5 on free). Reporting depth bundled with scheduling is what sets it apart here.

---

## LIMITATIONS

**The free tier excludes LinkedIn and X entirely,** and X/Twitter requires a \$10/mo add-on per account even on paid plans (grandfathered \$5/mo before 13 July 2026), so covering those two networks costs extra. Full analytics and multi-brand use also require moving off the free tier, from ~\$20/mo (annual) and up.

---

Pricing and limits read from **Metricool's own site** . Verified to the edition date above.

# Hypefury

## BUILT FOR

Creators on Threads, Bluesky, LinkedIn, Instagram and TikTok (X/Twitter support was dropped), typically solo builders and audience-growth hackers, comfortable with automation who want to monetize a following. Best fit: When you want aggressive growth automation (auto-reposts, evergreen recycling, DM funnels) over a calm multi-network scheduler for a brand.

---

## CHOOSE THIS IF YOU ARE

Solo creators on Threads, Bluesky, LinkedIn, Instagram or TikTok optimizing hard for audience and monetization growth (X/Twitter is no longer supported).

### BEST WHEN

You want aggressive growth automation no other tool here bundles: auto-reposts of top posts, evergreen recycling, auto-DMs to engagers, post-to-lead funnels, inspiration posts. Flexible \$6/mo per channel or Full \$19/mo for all channels (no free plan, 7-day trial).

### AVOID IF

You're X/Twitter-first (Hypefury dropped X support), visual-led, you want minimalist distraction-free writing, or automated DMs and reposts feel off-brand; the automation that's its strength can read as spammy.

---

## WHAT IT DOES

**Formerly X-first, it now covers Bluesky, Threads, LinkedIn, Instagram and TikTok after dropping X support.**

**Bluesky, Threads, LinkedIn, Instagram and TikTok (X/Twitter is no longer supported).**

Cross-post one piece to many, plus thread→LinkedIn-carousel and post→short-form-video repurposing.

**Fully automated publishing across all supported networks (no manual-reminder step),** plus scheduled auto-reposts, evergreen re-posting and rule-based/optimal-time queue management.

**Strong thread composer with an inspiration/post library,** quote-image generator and live preview. Note there is no AI content generation. You write everything, and Hypefury schedules and automates it.



**A core strength: an evergreen queue auto-recycles top-performing posts**, plus bulk CSV imports, categories, and auto-repost/auto-plug to keep older content compounding reach across time zones.

---

Pricing and limits read from [Hypefury's own site](#) . Verified to the edition date above.

**Basic engagement stats (impressions, likes, follower growth)**. 2026 reviews flag it as shallow. It shows what happened but offers no pattern analysis, content recommendations or AI-driven insights.

---

## COST

**No free plan; 7-day trial (no card)**. Current plans: Flexible \$6/mo for 1 channel (+\$6/mo per additional channel) and Full \$19/mo with all channels included; annual billing available. Full unlocks the advanced analytics and engagement tools.

---

## WHERE IT STANDS OUT

**A purpose-built founder growth loop**: auto-plugs on viral posts, auto-DM lead capture (100-300/day by tier), an engagement builder for targeted replies, and evergreen recycling. Posting feeds audience growth.

---

## LIMITATIONS

**No longer supports X/Twitter at all**. Current platforms are Bluesky, Threads, LinkedIn, Instagram and TikTok. X-first creators must look elsewhere. No free plan, and its automation/auto-DM tactics can read as spammy if overused, which some audiences penalize.

# Build

## BUILT FOR

The build-it-yourself route: raw platform APIs plus a cron job on a \$5 VPS. Beats paid schedulers on cost only if you already ship code and post enough to amortize the one-time setup. Solo technical founders at MVP stage posting to 1-2 channels, who value control and hate recurring fees more than they value a polished UI. Breaks the moment a non-coder needs to touch it.

---

## CHOOSE THIS IF YOU ARE

Technical founders who want zero subscription cost and full data control.

### BEST WHEN

You post to 1-2 platforms with stable APIs (LinkedIn, Bluesky/Mastodon) and can write a cron job or GitHub Action that pushes from a markdown/JSON queue.

### AVOID IF

You need Instagram/TikTok scheduling (their content-publishing APIs are painful and gated), visual previews, analytics, or team approval flows; note X's API is now pay-per-use (~\$0.015/post, \$0.20 with a link; the flat Basic tier was retired in 2026), and token refreshes plus API changes become your unpaid maintenance job.

## WHAT IT DOES

**No dashboard, no vendor.** You are the scheduler, the queue, and the on-call fix-it person.

**Whatever exposes an API: X, LinkedIn, Reddit, Mastodon, Bluesky, Telegram.** Instagram/TikTok/Threads are painful: gated, review-heavy, or reels-only via Graph API. You build each integration separately.

**Auto-publish is the point.** Cron fires the API call unattended. But there is no fallback. If a token expires or the API 500s at 2am, the post silently dies. You get reminders only if you code alerting yourself.

**Nonexistent unless you build it.** You draft in a text file, Notion, or markdown and pass strings to the API. No preview, no character counter, no thread builder. You eyeball formatting per platform.

**Trivial once coded.** A CSV or JSON queue plus a loop handles bulk and evergreen recycling better than most SaaS. This is where DIY actually wins. It's just data you control. Building the queue logic is the upfront cost.

---

No vendor page: this route is assembled from the components named above.

**You pull raw metrics from each platform's API and store them yourself:** full ownership, zero dashboard. Realistically most DIY founders skip this and check native app insights, because building charts isn't worth the time.

---

## COST

**No SaaS fee, but not free: ~\$0-6/mo hosting plus API costs.** X retired its flat \$200/mo Basic tier and made pay-per-use the default in Feb 2026 (legacy Basic migrated off by 1 June 2026): ~\$0.015 per post created (\$0.20 if it contains a link), no monthly minimum, so low-volume DIY posting is cheap, but reads and high volume add up.

---

## WHERE IT STANDS OUT

**Total control, no per-seat fees,** and unlimited custom logic (conditional posts, cross-repost, AI-drafted queues). If posting is part of a bigger automation you already run, folding it in is nearly free.

---

## LIMITATIONS

**Won't scale past a few channels or a teammate.** API terms shift and break you overnight, but token/security hygiene is on you. UI is whatever you tolerate. X's API is now pay-per-use (~\$0.015/post, \$0.20 with a link). Heavy posting or any read access can still outrun a flat-fee scheduler.

# Payments & Billing

Compared: [Stripe](#) · [Paddle](#) · [Polar](#) · [Lemon Squeezy](#) · [RevenueCat](#) · [Square](#) · [Build](#)

# At a glance

| TOOL                 | CHOOSE THIS IF YOU ARE                                                                                           | STARTING PRICE                                                                                                                                                 | MAIN BENEFIT                                                                                                                                                                                                                                                                                                                                                                          |
|----------------------|------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Stripe</b>        | <b>Technical founders who want the most flexible</b><br>best-documented payments API and will own tax compliance | <b>2.9% + \$0.30 per successful US online card</b><br>with add-ons that stack: +1.5% international cards, +1% currency conversion, +0.5% manually-entered      | <b>Its standout is breadth: the widest API and product suite, the lowest headline processing rate here, and the deepest third-party integration ecosystem (Shopify, most SaaS billing and accounting tools).</b><br>Add strong docs, 100+ payment methods and a full native Billing engine, and it is one of the most flexible options in the group.                                  |
| <b>Paddle</b>        | <b>SaaS founders selling globally who never want to think about sales tax or VAT</b>                             | <b>5% + \$0.50 per checkout transaction</b><br>all-inclusive of processing and global tax, with no monthly or migration fees                                   | <b>The 5% + \$0.50 fee is all-inclusive, bundling payments, global tax across ~200 countries, fraud, chargeback coverage and support into one rate as Merchant of Record.</b><br>The distinguishing point is consolidating billing and tax liability under a single legal-seller platform.                                                                                            |
| <b>Polar</b>         | <b>Indie hackers and developers wanting a modern</b><br>open-source Merchant of Record with the best DX          | <b>Starter free at 5% + \$0.50</b><br>Pro \$20/mo at 3.8% + 40c; Growth \$100/mo at 3.6% + 35c; Scale \$400/mo at 3.4% + 30c (\$240-\$4,800/yr for paid tiers) | <b>Open-source and developer-first as a Merchant of Record, with support for usage-based and credit-style models beyond flat subscriptions.</b><br>The combination of an open codebase, seller-of-record tax handling and consumption pricing is what separates it from the others here.                                                                                              |
| <b>Lemon Squeezy</b> | <b>Solo makers and creators selling digital downloads</b><br>courses, licenses, and simple SaaS                  | <b>5% + \$0.50 base</b><br>with no monthly fee and no paid tier                                                                                                | <b>The simplest all-in-one MoR for indie software sellers: a hosted storefront, license-key delivery, discounts and an affiliate system, all no-code and on Stripe-grade infrastructure, with global tax handled.</b><br>For a solo founder who wants to ship and start charging quickly without assembling billing, storefront and compliance separately, it is the fastest on-ramp. |
| <b>RevenueCat</b>    | <b>Mobile app founders monetizing via App Store</b><br>Google Play in-app subscriptions                          | <b>Free up to \$10,000/month MTR on the Basic plan</b><br>or up to \$2,500 MTR on Pro, after which Pro charges 1% of MTR                                       | <b>It is oriented specifically to mobile in-app subscriptions across both the App Store and Play Store, wrapping store billing with SDKs, analytics and payroll tooling.</b><br>The distinguishing point is unifying store subscription state and paywalls across the two platforms.                                                                                                  |
| <b>Square</b>        | <b>Founders with a physical retail</b><br>hospitality, or in-person services component alongside online sales    | <b>In-person 2.6% + 15c</b><br>(2.4% + 15c on Premium \$149/mo); online 3.3% + 30c on Free, 2.9% + 30c on Plus \$49/mo; manually keyed 3.5% + 15c              | <b>Unified in-person + online: one free POS, first-party hardware, one ledger, instant onboarding, no chargeback fee.</b><br>Stripe has no comparable card-present + POS + hardware stack.                                                                                                                                                                                            |
| <b>Build</b>         | <b>Engineering-led teams with genuinely unusual billing needs</b><br>appetite to own the stack                   | <b>Just the processor's cut</b><br>~2.9% + \$0.30 per charge on Stripe, no MoR markup (Paddle/LemonSqueezy add ~2-3% on top)                                   | <b>Lowest fees, total control, and no ceiling on custom logic.</b><br>If billing is a core competency or your margins are thin, owning the Stripe integration beats renting a MoR, as long as you can staff the tax and edge cases.                                                                                                                                                   |

Full detail for every tool follows on its own page.

# Stripe

## BUILT FOR

Engineering-led startups through large enterprises with developers who want programmable payments, subscriptions, and financial infrastructure via its APIs. Best fit: When you have technical resources and need maximum flexibility and global coverage, and are willing to handle sales-tax and merchant-of-record concerns yourself.

---

## CHOOSE THIS IF YOU ARE

Technical founders who want the most flexible, best-documented payments API and will own tax compliance.

### BEST WHEN

You're building custom billing (usage-based, metered, complex subscriptions, or marketplaces via Connect) and want the deepest ecosystem (Stripe Tax, Billing, Radar, Invoicing) at a 2.9% + 30¢ base.

### AVOID IF

You sell globally and don't want to become a tax filer in dozens of jurisdictions. A Merchant of Record remits tax for you; or you're mobile in-app subscription-led or in-person retail.

## WHAT IT DOES

**It is NOT a Merchant of Record.** You remain the legal seller and own sales-tax/VAT registration and remittance. A separate MoR option, Stripe Managed Payments, entered public preview in Feb 2026.

**No: YOU are the Merchant of Record and must register,** calculate, collect, file and remit sales tax/VAT in each jurisdiction. Stripe Tax (0.5%/txn no-code, or \$0.50/txn via API) calculates and collects but does not file or remit. The newer Stripe Managed Payments is a separate MoR option, in public preview since Feb 2026.

### Flat-rate and per-seat subscriptions plus

**usage/metered billing,** with metering covered up to 100 million events per month, and multi-phase schedules for pricing that changes over the contract term. All of these sit within the single 0.7% billing-volume layer on top of Stripe.

**No separate processor is needed,** but that is because Stripe itself is the processor. Billing runs on Stripe's payment rails at roughly 2.9% + \$0.30, and using Stripe Billing effectively commits the account to Stripe for money movement.

**Covers the full lifecycle within Stripe:** trials, proration on plan changes, upgrades and downgrades, and multi-phase schedules where pricing shifts across defined periods. Because it is native to Stripe, these lifecycle states are managed through the same API as payments.

**Smart Retries and automatic dunning reminders are built in and included within the 0.7% billing fee at no extra charge.** Failed-payment recovery does not require a separate tool or add-on. Retry logic and reminders are handled inside the same Stripe Billing layer.

**Native invoicing and quotes are part of the billing layer,** both covered within the 0.7% fee alongside recurring and usage charges. Invoices are generated inside Stripe instead of through a separate tool, keeping billing documents in the same system as payments.

**The widest set here: cards, Apple/Google Pay, ACH, SEPA, Link and 100+ local payment methods.** Stripe Billing adds a full native subscription engine with metered/usage-based billing, automatic card retries (Smart Retries), dunning emails and Revenue Recovery, while Radar provides ML fraud scoring.

**Rolling ~2 business days on a standard,** configurable schedule for established US accounts. The first payout is delayed ~7-14 days, up to 30 in higher-risk industries or countries. Instant Payouts settle in ~30 minutes for a 1.5% fee (50c minimum) once an account has ~60 days of processing history, in eligible countries.

**Business accounts can be created in ~46 countries,** while you can charge customers globally in 135+ presentment currencies plus a wide range of local payment methods. That makes it broad for accepting payments worldwide, though the set of countries you can actually operate an account from is narrower than its currency reach.

**Instant and self-serve.** You sign up and can accept payments within minutes, with risk review running in the background instead of as an upfront gate. Because that review happens after the fact, accounts can later be restricted or placed under reserve, and several high-risk business categories are restricted outright.

**Documented pattern of sudden freezes and rolling reserves on flagged or high-risk accounts.** Independent reports cite ~10-25% of volume held ~90 days, often triggered by risk flags or sudden sales spikes. Payouts can be paused during a review, and after a closure residual balances are typically held 90-180 days. A restricted-business list applies.

**The most engineering-heavy option here.** You wire up the API/SDK or embed hosted Checkout, and configure your own tax setup, webhooks and subscription logic. The docs and Elements/API are strong, and Payment Links offer a no-code path that can be live in minutes, but a full integration is real build-and-maintain work best handled by a developer.

## COST

**2.9% + \$0.30 per successful US online card**, with add-ons that stack: +1.5% international cards, +1% currency conversion, +0.5% manually-entered. Stripe Tax adds 0.5%/txn (or \$0.50 via API). Stripe Billing 0.7% of billing volume, or from \$620/mo on contract. ACH is 0.8% capped at \$5. Instant Payouts 1.5% (50c min). Effective all-in ~3-4%+.

**\$15 per dispute, charged even when you submit evidence and win**, because as your own Merchant of Record you carry the dispute liability and the operational work of contesting it directly. That cost and time burden grows with volume, and it is a direct consequence of Stripe not acting as the seller of record.

## ECOSYSTEM & EXIT

**Native to the Stripe ecosystem and its APIs**, so integrations run through Stripe's own platform instead of external connectors. This keeps billing, payments and related Stripe products in one system, but ties the integration surface to Stripe.

**Tied to Stripe.** Billing configuration moves within the Stripe ecosystem but not easily out of it, since the billing layer and processing are the same vendor. Leaving Stripe for payments effectively means rebuilding the billing setup elsewhere.

## WHERE IT STANDS OUT

**Its standout is breadth: the widest API and product suite**, the lowest headline processing rate here, and the deepest third-party integration ecosystem (Shopify, most SaaS billing and accounting tools). Add strong docs, 100+ payment methods and a full native Billing engine, and it is one of the most flexible options in the group.

## LIMITATIONS

**You own global tax registration, filing and remittance.** You carry chargeback liability directly (\$15/dispute even if you win), and fees stack once add-ons (int'l +1.5%, Tax, Billing) pile on. There is a documented freeze/rolling-reserve risk on flagged accounts, no storefront out of the box, and full billing/tax must be built and maintained yourself.

---

Pricing and limits read from [Stripe's own site](#) . Verified to the edition date above.



# Paddle

## BUILT FOR

SaaS and software companies, often small teams without a finance department, who want a merchant of record handling global tax, compliance, and billing. Best fit: When selling software worldwide and you want VAT/sales-tax and fraud offloaded, accepting less low-level control than a raw payments API gives.

---

## CHOOSE THIS IF YOU ARE

SaaS founders selling globally who never want to think about sales tax or VAT.

### BEST WHEN

You're a B2B/B2C software company wanting a Merchant of Record that becomes the legal seller (handling global tax registration and remittance, fraud, chargebacks, and localized checkout) for a flat 5% + 50¢, no monthly fee.

### AVOID IF

You want the lowest fees and will handle tax, you're an indie wanting a cheaper dev-friendly MoR, or you sell mobile in-app or in person.

## WHAT IT DOES

**In that role it bundles payment processing, subscription billing and global tax handling into one platform instead of leaving them to separate tools.**

**Yes: Paddle is the legal seller of record and registers,** files and remits VAT/sales tax across ~200 countries on your behalf. It issues compliant invoices, absorbs the tax liability, and produces one consolidated payout to reconcile, which removes the entire registration-and-filing burden that a non-MoR processor leaves to you.

### Recurring subscriptions and one-time

**purchases for digital goods,** billed and taxed by Paddle as Merchant of Record. Because it acts as the legal seller, global tax calculation and remittance are folded into the same pricing flow, not handled as a separate step.

**No: Paddle is both the payment processor and the Merchant of Record,** so it moves the money and acts as legal seller in one platform. There is no external gateway to connect or maintain, at the cost of not choosing your own processor.

**Subscription lifecycle is managed by Paddle in its role as seller of record.** Renewals, changes and cancellations run through the platform that also owns the tax and payment relationship. The customer relationship sits with Paddle instead of the merchant.

**Churn recovery is included as part of the Merchant-of-Record service,** not billed separately, so failed-payment retries fall under the single 5% + \$0.50 checkout fee alongside payments, tax and support.

**Tax-compliant invoices are generated by Paddle as the seller of record.** The invoice reflects Paddle as legal seller and carries the tax it collects and remits. Invoicing is therefore tied to the Merchant-of-Record role, but the merchant does not issue it.

**Cards, PayPal, Apple Pay and wallets,** with mature subscription management, a RevenueCat integration for apps, and a developer API/webhooks. Failed-payment recovery and dunning (Retain/ProfitWell), churn analytics and free ProfitWell Metrics are built in, with fraud protection included as MoR. Fewer third-party integrations than Stripe.

**Slow: Paddle creates a payout on the 1st and sends it by the 15th,** taking up to 3 working days to arrive, subject to a \$100 (£100/€100) minimum and paid by wire transfer or Payoneer. Cadence is account-configurable (some report ~weekly/bi-weekly), but the monthly cycle plus a holding period is the slowest cash access here and strains early cashflow.

**Sells to almost 200 countries and prices in 30+ currencies as the MoR,** with balances held in USD/EUR/GBP and payouts in 13 currencies. Russia, Belarus, Iran and North Korea are excluded due to sanctions. Seller onboarding is global but gated by a manual approval step, so where you can operate from depends partly on clearing that review.

**Manual underwriting, typically ~2-7 business days,** and Paddle can reject new or unusual product types. Reported triggers include an undisclosed ~3 months of prior payment-processing history, conditional/qualified refund-policy language, and a legal-entity name that doesn't match your site. Multiple rejection cycles are commonly reported by indie founders.

**Documented sudden account closures with terse notices and category reclassifications (e.g. to 'human services'),** and its risk team reportedly uses AI and can close accounts after years. As MoR it enforces strict product eligibility and can decline or offboard accounts. Support is reported as slow (2-3 days). The gated, opaque approval is the main structural risk here.

**Low-code via a hosted checkout,** so tax and billing infrastructure do not have to be built. The main friction is onboarding. Going live requires manual approval by Paddle before the account can transact, which adds a review step ahead of launch.

## COST

**5% + \$0.50 per checkout transaction**, all-inclusive of processing and global tax, with no monthly or migration fees. A currency-conversion margin of up to 1.5% applies, so the effective rate climbs toward ~6-7% on international sales, and independent sellers report ~12-13% on small subscriptions. Products under \$10 and invoicing need custom pricing, but the fee is NOT refunded when you refund a customer.

**\$15 (£15/€15) per card chargeback and \$20 for PayPal**, passed to the seller in addition to the original transaction amount. Paddle fights disputes on your behalf as MoR, and both the fee and the original amount are returned if it wins. Fraud and chargeback protection are bundled into the base 5% fee, not billed separately.

## ECOSYSTEM & EXIT

**A hosted or overlay checkout plus an API and webhooks**, so integration centers on embedding Paddle's checkout and reacting to its events. The connection points are oriented around the Merchant-of-Record checkout, without a broad third-party connector library.

**As Merchant of Record, Paddle owns the customer and payment relationship**, so the customers are legally Paddle's. Leaving means re-collecting payment credentials and re-establishing those relationships under a new seller.

## WHERE IT STANDS OUT

**The 5% + \$0.50 fee is all-inclusive**, bundling payments, global tax across ~200 countries, fraud, chargeback coverage and support into one rate as Merchant of Record. The distinguishing point is consolidating billing and tax liability under a single legal-seller platform.

## LIMITATIONS

**At 5% + \$0.50 the all-inclusive rate is higher than assembling Stripe Billing plus separate processing**. Going live also depends on manual approval, and category-based account closures have been reported, adding review and continuity risk.

---

Pricing and limits read from **Paddle's own site** . Verified to the edition date above.

# Polar

## BUILT FOR

Indie developers and open-source maintainers, highly technical and small, monetizing digital products, subscriptions, and usage with a developer-first merchant of record. Best fit: When you're a solo dev or tiny team wanting modern DX and MoR tax handling, over a heavier enterprise billing platform.

---

## CHOOSE THIS IF YOU ARE

Indie hackers and developers wanting a modern, open-source Merchant of Record with the best DX.

---

## WHAT IT DOES

**Alongside standard subscriptions it supports usage-based and credit-style pricing models,** and its billing runs on Stripe payment rails underneath.

**Yes: Polar is the legal seller of record and remits VAT/sales tax across 60+ jurisdictions,** running on underlying Stripe, PCI-DSS-Level-1 infrastructure. As MoR it also handles dispute and fraud prevention. Coverage is solid and growing but, as a younger platform, has a shorter track record than Paddle's, and Stripe's holds and reviews apply underneath.

**Subscriptions alongside usage-based billing and credit-style models,** so pricing can be charged by consumption or against prepaid credits, not just on a fixed recurring plan. This range beyond basic subscriptions is Polar's stated focus.

**No: Polar acts as the Merchant of Record,** so it takes the legal-seller role instead of requiring a separate gateway. Payments run on Stripe rails underneath, but that is arranged through Polar, not a processor you connect yourself.

**Subscriptions with upgrades and downgrades plus usage tracking.** Consumption is metered over the billing period alongside standard plan changes. Lifecycle events run through Polar as the Merchant of Record, with the underlying payments settled on Stripe rails.

**Automated dunning for failed payments is provided.** Retries on declined charges are handled by the platform instead of requiring the merchant to build recovery logic. It runs as part of the Merchant-of-Record service.

**Invoices are generated by Polar in its Merchant-of-Record role.** The document reflects Polar as the legal seller, not produced by the merchant directly. Invoicing is bundled into the same seller-of-record service that handles tax and payments on Stripe rails.

**Card payments only, via Stripe,** with subscriptions, one-time purchases and strong usage-based/metered billing, a strength for developer and AI tools billed by consumption. It also offers license keys, benefits, a customer portal and webhook events, though it carries lighter enterprise-contract tooling than Paddle.

**A 7-day settlement delay applies by default to organizations created on or after 2026-05-12;** organizations created before that keep instant payouts. Withdrawals are initiated manually via Stripe Connect Express, and Stripe's payout fees pass through at cost. User reports flag payout delays (funds not received after ~25 days) and occasionally blocked payouts.

**Sells globally as MoR, with seller onboarding relying on Stripe-supported regions.** Official docs list 160+ payout countries via Stripe Connect Express. Cuba, Russia, Iran, North Korea and Syria are excluded. Card payments run via Stripe, and the integration set is developer-oriented (SDKs, API, webhooks), not broad local-method coverage.

**Self-serve onboarding with KYC handled via Stripe,** and developer-friendly by design. API-first with SDKs, so an engineer can wire up a working integration fast. There is no heavy manual approval gate. The main expectation is developer comfort, since it offers less no-code hand-holding than Lemon Squeezy or Gumroad.

**As MoR, dispute and fraud handling shift to Polar,** but it is a smaller platform with a shorter dispute-handling track record, and Stripe's underlying holds and reviews still apply. User reports describe payout delays (~25 days), blocked payouts, sudden account closures with poor communication, checkout outages, and support that can take a week.

**Fast for developers via its API,** aimed at technical founders integrating through REST endpoints and webhooks. It assumes comfort with code instead of offering a no-code path, but for that audience the integration is quick.

---

## COST

**Starter free at 5% + \$0.50; Pro \$20/mo at 3.8% + 40c;** Growth \$100/mo at 3.6% + 35c; Scale \$400/mo at 3.4% + 30c (\$240-\$4,800/yr for paid tiers). Add +1.5% on international (non-US) cards and \$15/dispute. Stripe payout fees pass through at cost. Orgs created before 2026-05-27 keep a grandfathered 4% + 40c Early Member rate (+0.5% subs) until they upgrade.

**\$15 per dispute, flat, on every plan.** As an MoR built on Stripe, Polar handles the dispute and fraud-prevention process on your behalf. The operational side is largely off you, though the per-dispute fee still applies. The exact chargeback fee passed through from Stripe is not separately published by Polar.

---

#### ECOSYSTEM & EXIT

**A REST API and webhooks, with an open-source codebase.** Integration is code-first through documented endpoints and event callbacks. The open-source foundation means the implementation itself can be inspected and extended.

**The Merchant-of-Record relationship means migrating off re-onboards customers to a new seller,** since Polar was the legal seller of record. The payment relationship does not transfer cleanly. Customers effectively have to be re-established elsewhere.

---

#### WHERE IT STANDS OUT

**Open-source and developer-first as a Merchant of Record,** with support for usage-based and credit-style models beyond flat subscriptions. The combination of an open codebase, seller-of-record tax handling and consumption pricing is what separates it from the others here.

---

#### LIMITATIONS

**As the youngest Merchant of Record in this set it carries more payout and support maturity risk than longer-established platforms.** Its payments run on Stripe rails and are limited to card payments. The track record and rail breadth are what you give up.

---

Pricing and limits read from [Polar's own site](#) . Verified to the edition date above.

# Lemon Squeezy

## BUILT FOR

Solo founders and small teams selling digital products, SaaS, and licenses who want an all-in-one merchant of record with minimal setup. Best fit: When you want fast, low-code global selling with tax handled, and don't need deep API customization or enterprise-scale billing logic.

---

## CHOOSE THIS IF YOU ARE

Solo makers and creators selling digital downloads, courses, licenses, and simple SaaS.

## BEST WHEN

You want the fastest no-code path to a global storefront with MoR tax handling, license keys, and built-in affiliate and email tools: 5% + 50¢ plus surcharges (1.5% international cards, 1.5% PayPal, 0.5% subscriptions).

## WHAT IT DOES

**Popular with indie hackers, it runs on Stripe-grade infrastructure and the 13-person team was acquired by Stripe in July 2024.** It is now being folded into Stripe Managed Payments, in public preview since Feb 2026.

**Yes: Lemon Squeezy is the legal seller of record and takes on global VAT/sales-tax collection and remittance**, processing refunds and producing a single consolidated payout. It runs on Stripe-grade, PCI-DSS-Level-1 infrastructure and lists AI fraud protection, so the same core compliance burden a non-MoR processor leaves to you is handled here.

**One-time, subscriptions, usage-based billing (new)**, license keys and discounts/affiliates, all as Merchant of Record.



**No: Lemon Squeezy is itself the processor and Merchant of Record (cards, PayPal, wallets).** You integrate only LS.

**Subscriptions with a customer portal and failed-payment recovery**, managed via hosted checkout as MoR.

**Built-in failed-payment recovery/dunning to reduce involuntary churn.**

**Issues receipts and tax-compliant invoices per order as Merchant of Record.**

**Up to 21 payment methods including PayPal**, with subscriptions, newer usage-based billing, a customer portal, failed-payment recovery/dunning and fraud prevention. Automated software license-key issuance is a standout for selling code. Subscription depth is narrower than Stripe's, and the toolset is aimed at simple digital-product billing instead of enterprise contracts.

**Roughly a 13-day hold on net sales**, with payouts processed twice monthly on the 1st and 15th and funds arriving around the 14th and 28th (then 1-5 days to reach the bank). Payouts are free to US banks and cost 1% to international banks, faster than Paddle's monthly cycle but slower and fee-bearing compared with a direct processor.

**Sells globally as the Merchant of Record and pays out across a broad set of geographies:** bank transfer to ~79 countries and PayPal to 200+ countries/regions, roughly 279+ in total. Seller signup is self-serve with KYC and quick to go live, making it practical for indie founders in most places to start selling internationally.

**Self-serve signup with KYC and no manual approval gate**, so it is relatively quick to go live. This low-friction onboarding, combined with MoR tax handling out of the box, is a core reason indie hackers reach for it. You can set up a hosted checkout and start charging without clearing a review.

**As MoR it absorbs chargeback risk but can hold or review accounts.** Day-to-day holds are less reported than Paddle, yet transition risk is high. Support has reportedly slowed since the Stripe acquisition (users cite waits of 'weeks'), and in May 2026 a support agent accidentally mass-cancelled one customer's (Screen Studio) subscriptions. Community sentiment has cooled.

**Among the easiest here.** A hosted checkout, storefront and checkout overlays need no backend code, with an API and webhooks for deeper builds. You can embed a link or button and be live quickly. The whole flow is designed so a developer (or a non-technical seller) can start selling a digital product with minimal setup.



## COST

**5% + \$0.50 base, with no monthly fee and no paid tier.** Surcharges stack: +1.5% international cards, +1.5% PayPal, +0.5% on subscriptions, 5% on recovered carts and 3% on affiliate sales, plus a 1% international payout fee. Independent estimates put the real effective rate anywhere from 5.5% to 11% depending on geography, methods and features used.

**A \$15 dispute fee that is not returned whether or not the dispute is contested.** As the Merchant of Record, Lemon Squeezy manages disputes directly with the card networks. The fraud/chargeback liability is largely off you, but the flat per-dispute fee still lands on your account when a chargeback occurs.

## ECOSYSTEM & EXIT

**Cards, PayPal and wallets.** API and webhooks. No-code storefront. Integrations narrower than Stripe.

**Central risk: users are being migrated toward Stripe Managed Payments.** Some features won't carry over and new development has slowed, but effectively a forced migration.

## WHERE IT STANDS OUT

**The simplest all-in-one MoR for indie software sellers:** a hosted storefront, license-key delivery, discounts and an affiliate system, all no-code and on Stripe-grade infrastructure, with global tax handled. For a solo founder who wants to ship and start charging quickly without assembling billing, storefront and compliance separately, it is the fastest on-ramp.

## LIMITATIONS

**Stacked surcharges push the effective rate to 5.5%-11%,** and payouts carry a ~13-day hold plus a 1% international fee. The bigger risk is strategic: acquired by Stripe and now being folded into Stripe Managed Payments (3.5% on top of standard Stripe processing, ~6.4% + \$0.30 all-in on a domestic US card), with some features (affiliates, storefront, downloads) reportedly not carrying over, slowed development, and no firm sunset date.

Pricing and limits read from [Lemon Squeezy's own site](#) . Verified to the edition date above.

# RevenueCat

## BUILT FOR

Mobile app developers and product teams, technical, running iOS/Android subscriptions who need cross-platform in-app purchase infrastructure and subscription analytics. Best fit: When your revenue is app-store IAP and you need to manage entitlements and metrics across platforms, not general web card payments.

---

## CHOOSE THIS IF YOU ARE

Mobile app founders monetizing via App Store and Google Play in-app subscriptions.

### BEST WHEN

You ship iOS/Android and need to manage subscriptions, entitlements, receipt validation, paywalls, and cross-platform subscriber analytics without building StoreKit/Billing plumbing (free up to \$2,500 monthly tracked revenue, then 1% of MTR). The stores are your processor and RevenueCat sits on top.

### AVOID IF

You sell web/desktop SaaS or digital goods, you process payments outside the app stores, or you're pre-revenue and want to avoid another SDK, though the free tier makes early adoption low-risk.

---

## WHAT IT DOES

**It layers purchase SDKs, subscription analytics and paywall tooling over native store billing**, so the stores still process payment while RevenueCat manages entitlement state and reporting.

**Apple/Google are the Merchant of Record for in-app purchases and handle store-region tax**, but RevenueCat wraps App Store/Play billing and is not the seller of record.

**Mobile in-app subscriptions and one-time purchases**, transacted through the App Store and Play Store, plus Web Billing (via Stripe) for selling subscriptions on the web. Store purchases stay bounded by what the underlying store billing supports, with RevenueCat managing the entitlements on top.

**No separate processor is used.** RevenueCat wraps the App Store and Play Store billing systems, which handle the actual charge. The stores act as the payment mechanism, so there is no external gateway but also no way around store billing.

**Lifecycle is driven by the underlying stores.**

Renewals, grace periods and related states originate in App Store and Play Store billing, and RevenueCat surfaces that store-driven state consistently across both platforms; it does not control the billing cycle itself.

**Recovery is largely handled by the underlying app stores' billing retry logic, not by**

RevenueCat, since the stores process the payment. Failed-charge handling therefore depends on App Store and Play Store retry behavior, not a dunning system the merchant configures.

**Transactions produce store receipts through the App Store and Play Store, not traditional B2B invoices.** Since the stores are the seller to the end user, RevenueCat reports on those receipts instead of issuing merchant invoices of its own.

**Recurring runs through App Store/Play in-app purchases (end-user methods = whatever the stores accept, incl. Apple/Google Pay).**

RevenueCat manages subscription state, entitlements and paywalls on top.

**RevenueCat never touches funds;** Apple and Google pay out on their own schedules (typically monthly, with a lag).

**Available wherever the App Store and Google Play operate,** but the RevenueCat SDK is global.

**Needs Apple Developer and/or Google Play accounts (and their review);** RevenueCat itself is self-serve SDK setup.

**Funds and account standing sit with**

**Apple/Google, not RevenueCat.** Store account suspension is the risk to watch.

**Setup is a mobile SDK integration on iOS and**

**Android,** wiring the app to the stores' billing through RevenueCat's SDKs. The effort is concentrated in the mobile client, not a server-side billing build, but it requires native app integration.

---

**COST**

**Free up to \$10,000/month MTR on the Basic plan,** or up to \$2,500 MTR on Pro, after which Pro charges 1% of MTR. Because the App Store and Play Store process the payments, this fee is charged on the tracked revenue that flows through that store billing.

**Refunds/disputes are handled by**

**Apple/Google under store policy;** no separate RevenueCat dispute fee.

---

**ECOSYSTEM & EXIT**

**iOS and Android SDKs for the client,** plus analytics and attribution integrations; subscription data can flow into product and marketing tooling. The integration surface centers on mobile clients and the App Store/Play Store billing it wraps.

**Mobile-store subscription state complicates moving off,** because active subscriptions and entitlements live in the App Store and Play Store systems that RevenueCat wraps. Migrating means carrying that store-held subscription state to another setup.

---

## WHERE IT STANDS OUT

**It is oriented specifically to mobile in-app subscriptions across both the App Store and Play Store**, wrapping store billing with SDKs, analytics and paywall tooling. The distinguishing point is unifying store subscription state and paywalls across the two platforms.

---

Pricing and limits read from [RevenueCat's own site](#) . Verified to the edition date above.

---

## LIMITATIONS

**It wraps App Store and Play Store billing**, plus newer Web Billing via Stripe. It is no longer strictly mobile-only. On the Pro plan the fee is 1% of MTR beyond the free threshold, so cost scales with tracked revenue once past \$2,500 MTR.

# Square

## BUILT FOR

Brick-and-mortar retailers, restaurants, and service businesses, non-technical owners, needing point-of-sale hardware plus integrated online payments.

Best fit: When you sell in person and want POS, inventory, and payments bundled, instead of a developer-first API for a software product.

---

## CHOOSE THIS IF YOU ARE

Founders with a physical retail, hospitality, or in-person services component alongside online sales.

## BEST WHEN

You take in-person card payments and want a free POS app, hardware, and unified online+offline inventory (2.6% + 15¢ tapped/dipped in person, ~2.9% + 30¢ online, free POS software).

## AVOID IF

You're pure SaaS/digital with no in-person sales, you sell internationally and need tax remittance, or you're mobile-IAP.

---

## WHAT IT DOES

**You accept card, tap, and online payments on one Square account and dashboard**, no separate acquirer.

**No: Square is a payment processor, not a Merchant of Record.** You are the seller of record and remain responsible for calculating, collecting, and remitting your own sales tax and VAT.

**One-time, in-person, invoices, and recurring via the Subscriptions API (weekly/monthly/yearly cadences, trials, discount phases).** No native usage-based metering. You meter externally and bill through the API.

**No. Square is the processor and acquirer.** It settles funds directly to your bank. It replaces both the gateway and the merchant account in a single stack.

**Subscriptions API creates plans**, plan variations, and per-customer subscriptions. Supports trials, discount phases, pauses, and cancellation. Card-on-file customers are auto-charged, but others get an emailed invoice with a pay link.

**Basic retries and automatic reminder emails on failed or unpaid invoices/subscriptions.** No advanced smart-retry or configurable dunning schedules like Stripe Billing. You build escalation logic yourself.

**Yes: native Square Invoices, free to create.** You pay standard processing only when paid. Supports one-off, recurring, milestone, and estimate-to-invoice, with card-on-file, tips, and payment links.

**Cards, Apple Pay, Google Pay**, tap-to-pay, Cash App Pay, gift cards, and Afterpay (BNPL). Recurring via card-on-file + Subscriptions API. No native ACH for standard checkout the way some rivals offer.

**Free next-business-day deposits (payments before 5pm PT arrive next business day).**

Instant or same-day transfers cost 1.95% per transfer, 24/7, with a \$25 minimum balance after fees.

**Payments available in 8 markets:** US, Canada, UK, Ireland, Australia, Japan, France, and Spain. Narrower global reach than Stripe or Adyen, mostly single-country selling per account.

**Instant: aggregator model, sign up and take payments in minutes with no separate underwriting wait.** In exchange there is risk-based review that can trigger later holds once volume grows.

**A real risk.** Square can freeze funds 90-180 days or deactivate accounts algorithmically, often with limited appeal, especially on sudden volume spikes or high-risk categories. Reserves can be debited without prior notice.

**Low for point-and-click (POS, invoices, hosted online store, no code).** Moderate for custom builds via Web Payments SDK and REST APIs, well-documented but less deep than Stripe's developer tooling.

---

## COST

**In-person 2.6% + 15c (2.4% + 15c on Premium \$149/mo).** Online 3.3% + 30c on Free, 2.9% + 30c on Plus \$49/mo. Manually keyed 3.5% + 15c. No monthly fee on the Free plan.

**No per-dispute fee (vs Stripe's \$15),** but you only lose the disputed amount if you lose. Eligible sellers get free Chargeback Protection covering up to \$250/month in qualifying disputes.

---

## ECOSYSTEM & EXIT

**Large App Marketplace plus WooCommerce,** BigCommerce, Wix, QuickBooks, and payroll/HR tools. Native Square Online store and hardware. Fewer third-party dev libraries and plugins than Stripe's ecosystem.

**Card-on-file and customer data can be migrated in via Square's import and Customers/Cards APIs.** Moving hardware and POS workflows is heavier. PAN migration requires a compliant transfer request between processors.

---

Pricing and limits read from [Square's own site](#) . Verified to the edition date above.

---

#### WHERE IT STANDS OUT

**Unified in-person + online: one free POS,** first-party hardware, one ledger, instant onboarding, no chargeback fee. Stripe has no comparable card-present + POS + hardware stack.

---

#### LIMITATIONS

**Only 8 countries.** Higher online rates than Stripe on the free plan. Weaker dunning and no usage-based billing. Account-freeze/hold risk. Thinner developer tooling for complex custom platforms.

# Build

## BUILT FOR

Maximum control and lowest fees, but you inherit every edge case, including sales tax, that a MoR would have absorbed. Technical teams with a backend and someone who understands billing edge cases. Great for a lean MVP charging a few plans; dangerous once tax jurisdictions, proration, and dunning pile up.

---

## CHOOSE THIS IF YOU ARE

Engineering-led teams with unusual billing needs and appetite to own the stack.

## AVOID IF

You become liable for global sales tax/VAT yourself, plus PCI scope, webhook reliability, and refund edge cases; most founders should start managed and drop to raw APIs only where forced.

---

## WHAT IT DOES

**You're assembling a billing system, not buying one.** Your engineers replace the SaaS layer.

**You are the merchant.** No one shields you. You register, collect, and remit sales tax/VAT yourself. Stripe Tax calculates rates but does NOT file. This is the single biggest hidden cost of going DIY.

**Anything you're willing to code:** flat, tiered, per-seat, usage/metered, one-time. Stripe Billing primitives cover most. Exotic hybrid models mean custom logic. Flexibility is unlimited, effort scales with weirdness.

**No: Stripe IS the processor.** That's the whole model. You talk to it directly instead of through a reseller. You do need to add tax, analytics, and email layers yourself.



**Stripe Billing handles renewals**, upgrades, proration, and cancels via API/webhooks, but you must wire and test every webhook. Miss one and a canceled user keeps access, or a card update silently fails.

**Stripe's Smart Retries and email reminders exist but must be enabled and configured.** Deeper recovery flows you build. Weaker out-of-the-box than dedicated dunning tools. Involuntary churn creeps up if you ignore it.

**Stripe Invoicing generates hosted invoices and PDFs natively**, solid for most. Custom branding, consolidated invoices, or region-specific compliance formats are on you to configure or template.

**Cards, ACH, SEPA, wallets, and 40+ local methods via Stripe**, recurring is first-class. You enable and test each method. Every added method is more UI states and failure paths you own.

**Stripe's schedule: ~2 business days in the US after the initial ~7-day rolling hold**, faster/instant options for a fee. No MoR sitting on your money, but no MoR advancing it either. You carry refund risk.

**Stripe operates in ~46 countries for accepting.** You can sell globally but must handle each region's tax/compliance yourself. A MoR would cover global sales for you. DIY means you gate or lawyer-up per market.

**You open your own Stripe account**, fast for standard businesses, minutes to live. Higher-risk verticals face review or rejection with no reseller to hide under. Your legal entity is exposed directly to the processor.

**Risk is fully yours.** Stripe can freeze funds or offboard you for spikes, disputes, or policy hits, and you have no MoR buffer. Mitigate by staying low-risk and diversifying, but you're one flag from a frozen balance.

**The real price.** Days to a basic checkout, weeks-to-months for full-featured subscriptions, tax, dunning, and webhook reliability. This one-time build only pays off versus a MoR's % fee at meaningful, sustained volume.

---

## COST

**Just the processor's cut: ~2.9% + \$0.30 per charge on Stripe**, no MoR markup (Paddle/LemonSqueezy add ~2-3% on top). You save that spread, and spend it back in engineering time and tax tooling.

**You eat them directly, at ~\$15 per dispute plus the lost amount**, and assemble the evidence yourself. Stripe Radar helps but needs tuning. No MoR to absorb fraud, so chargebacks hit your revenue and your account health.

---

## ECOSYSTEM & EXIT

**Stripe has the deepest ecosystem and SDKs of any processor.** You can connect to almost anything, but you write the glue. Accounting, tax filing, and analytics tools integrate, yet wiring and maintaining them is your job.

**Onboarding is easy.** Leaving is the hard part. You own the customer/card data (PCI scope) and Stripe helps export tokens to another processor. Migrating off your own custom billing logic later is a full re-engineering project.

---

#### WHERE IT STANDS OUT

**Lowest fees, total control, and no ceiling on custom logic.** If billing is a core competency or your margins are thin, owning the Stripe integration beats renting a MoR, as long as you can staff the tax and edge cases.

---

#### LIMITATIONS

**You become the tax authority,** the fraud team, and the compliance department. Slower to launch, risky if under-built, and legal exposure is direct. For global consumer sales a MoR is usually worth its cut over pure DIY.

---

No vendor page: this route is assembled from the components named above.

# App Building

Compared: Lovable · Bolt.new · Replit · FlutterFlow · v0 · Bubble · Build

# At a glance

| TOOL        | CHOOSE THIS IF YOU ARE                                                                                   | STARTING PRICE                                                                                                                                                                   | MAIN BENEFIT                                                                                                                                                                                                                                                                          |
|-------------|----------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Lovable     | Non-technical and semi-technical founders shipping a polished full-stack web MVP fast                    | Free tier: 5 daily build credits (up to 30/mo)                                                                                                                                   | <b>A bundled Lovable Cloud backend plus Stripe and a clear 'you own your code' stance make it a common default for shipping a full web MVP without stitching backend services together.</b><br>You describe the app, get visible React you own, and export to GitHub on paid plans.   |
| Bolt.new    | Technical-ish founders wanting in-browser full-stack prototyping with direct control over the code stack | Free \$0 with no card (1M tokens/mo, 300K daily cap); Pro \$25/mo (10M+ tokens with rollover); Teams \$30/member/mo; Enterprise custom                                           | <b>Zero lock-in paired with an instant, install-free in-browser full-stack runtime (WebContainers).</b><br>The app builds and runs client-side as you prompt, output is standard React/Vite, and GitHub auto-sync plus Netlify deploy export and host anywhere.                       |
| Replit      | Founders wanting a full cloud IDE plus an AI Agent that can build run, host, and deploy in one place     | Core \$20/mo bundles \$25 in usage credits up to 5 collaborators and unlimited workspaces; Pro \$100/mo adds teams (up to 15 builders), pooled credits and rollover              | <b>One place to build, run and host (a full IDE, built-in database and one-click deploys) combined with an AI Agent whose effort-based billing keeps small changes typically under \$0.25 per checkpoint.</b><br>Build-plus-host in a single environment is the differentiator.       |
| FlutterFlow | Founders who need real native iOS AND Android mobile apps beyond web alone                               | Basic \$39/mo (annual \$351/yr)<br>Growth \$80/mo for the first seat (annual \$720/yr), Business (annual \$1,350/yr), with annual billing running about 25% cheaper than monthly | <b>The only pick here that produces genuine native mobile apps (real iOS/Android and web from one project) while still handing you clean, fully-owned Flutter/Dart you can export from the Basic plan and develop anywhere.</b><br>Native output plus code ownership is the standout. |
| v0          | Founders and developers who already own a backend just need beautiful frontend UI fast                   | Permanent free tier worth ~\$5 of tokens/mo (roughly 3-8 prompts/day)                                                                                                            | <b>The tightest Next.js/Vercel/GitHub workflow of the group: native GitHub branches and PRs, direct Vercel deploy, a permanent free tier, and Vercel's SOC 2 Type II posture behind it.</b><br>For a front-end team already on Vercel, the integration is the draw.                   |
| Bubble      | Non-technical founders wanting deep no-code control over complex web app logic data                      | Free plan (50k Workload Units, dev-only, no live deploy)                                                                                                                         | <b>The deepest no-code backend of the bunch: a database, complex multi-step workflows, and a massive plugin and agency ecosystem.</b><br>Best when a non-coder needs actual application logic, not just a polished front end.                                                         |
| Build       | Technical founders who can ship code and want zero platform lock-in                                      | A time -vs- money trade<br>not a flat fee: ~\$20-200/mo for the AI plan plus API/token burn that spikes with sloppy prompting                                                    | <b>High flexibility and ownership at near-no marginal cost per feature.</b><br>If you can code even a little, it collapses build time while leaving you a real, scalable, portable product, with no platform tax and no walls.                                                        |

Full detail for every tool follows on its own page.

# Lovable

## BUILT FOR

Non-technical founders and product people prototyping full-stack web apps through natural-language prompts, generating React front ends with backend integration. Best fit: When you want to describe an app and get a working, editable codebase fast, over hand-coding or a purely visual no-code builder.

---

## CHOOSE THIS IF YOU ARE

Non-technical and semi-technical founders shipping a polished full-stack web MVP fast.

### BEST WHEN

You want to describe an app in chat and get a React + Supabase full-stack product with auth, database, and clean UI, plus GitHub sync and Supabase integration for when you outgrow the prompt box. Free gives 5 daily/30 monthly credits, Pro \$25/mo (100 credits), Business \$50/mo.

### AVOID IF

You need native mobile apps, pixel-level no-code control, or you're a strong engineer who'd move faster hand-coding, and watch credit burn on heavy debugging iteration.

## WHAT IT DOES

**You edit by chat, all code stays visible**, and Stripe plus GitHub export (paid) are wired in, a prompt-to-product path for web MVPs, internal tools and landing pages.

**A working React/Tailwind web app plus a bundled Lovable Cloud (Supabase-style) backend**, with Stripe integration available. All code is visible and editable by chat, and per Lovable's own terms 'you own your code' and any AI output. GitHub export is unlocked on paid plans.

**Minutes to a working web prototype from a single prompt.** The no-code entry means you can go from description to a running React/Tailwind app on the free tier (5 daily build credits) without any setup beyond a browser and account.

**Low to build: the prompt-first**, no-code entry gets a non-technical founder to a running app. Fixing is where skill matters. All code is visible and edits happen by chat or in GitHub, so some coding and GitHub comfort helps when the generated React or the Lovable Cloud backend needs correcting.

**Responsive web apps only, with no native mobile build.** Output targets the browser (React/Tailwind). If you need a true iOS/Android app this isn't the tool, but it's web-first by design.

**Good for web MVPs, internal tools and landing pages;** because all code is visible and exportable to GitHub (paid), complex or custom logic can eventually be hand-coded in the exported React repo, though doing so means working around the Lovable Cloud backend it bundles.

---

### COST

**Free tier: 5 daily build credits (up to 30/mo).** Paid Pro \$25/mo and Business \$50/mo, each including 100 monthly credits (Pro ~\$21/mo billed annually). Each chat message costs ~0.5 credits (simple) up to ~1.5 (complex), so cost scales directly with how much you iterate; heavy back-and-forth burns the allotment fast.

---

### ECOSYSTEM & EXIT

**Per Lovable's own terms you own the generated code and any AI output.** It's standard React and exportable to GitHub on paid plans. The limit is Lovable Cloud. Because it also hosts your backend/database, leaving means re-homing that backend, so the front end is portable but the bundled stack creates soft lock-in.

---

### WHERE IT STANDS OUT

**A bundled Lovable Cloud backend plus Stripe and a clear 'you own your code' stance make it a common default for shipping a full web MVP without stitching backend services together.** You describe the app, get visible React you own, and export to GitHub on paid plans.

---

### LIMITATIONS

**Soft lock-in via Lovable Cloud,** which hosts the bundled backend even though the React front end is exportable. Credit burn that scales with iteration (~0.5-1.5 credits per message), but web only, with no native mobile. Generated apps also still need your own security review for AI-code risks.

---

Pricing and limits read from [Lovable's own site](#) . Verified to the edition date above.

# Bolt.new

## BUILT FOR

Developers and technical tinkerers who want AI to scaffold and run full-stack apps in an in-browser environment with live preview and deploy. Best fit: When you want prompt-driven generation you can immediately edit and ship in-browser, not a managed no-code platform hiding the code.

---

## CHOOSE THIS IF YOU ARE

Technical-ish founders wanting in-browser full-stack prototyping with direct control over the code and stack.

### BEST WHEN

You want to generate AND run a full app entirely in the browser (StackBlitz WebContainers), pick your framework, and edit files directly (Free 1M tokens/mo with a 300K daily cap, Pro \$25/mo with ~10M tokens, Teams \$30/member).

### AVOID IF

You want maximum design polish and hand-holding, native mobile, pure UI components, or you'd burn tokens fast on debugging loops; Bolt's token model gets expensive on complex apps.

## WHAT IT DOES

**It generates and runs full-stack JavaScript apps entirely in the browser:** full Node runtime, package manager and build all client-side, with no remote sandbox. Output is standard React/Vite, downloadable, and GitHub auto-sync plus Netlify deploy are built in.

**Standard React/Vite project code that runs live in-browser via WebContainers and that you can download,** modify and host anywhere, with no proprietary format. GitHub auto-sync keeps a repo in step, and you can deploy out to Netlify or other hosts.

**Immediate: there is no install and no remote sandbox to spin up.** The app builds and runs in the browser via WebContainers as soon as you prompt. The free tier (\$0, no card) starts instantly within the 1M-token monthly / 300K daily allowance.

**Moderate: it's aimed at developers,** and output is standard React/Vite, so familiarity with that stack helps a lot when generations break or need manual fixing. The in-browser runtime lowers setup friction, but debugging still assumes you can read and edit JavaScript.

**Web apps only, no native mobile.** Everything runs and ships as browser-based JavaScript (React/Vite) via WebContainers, so 'mobile' here means a responsive web app, not a native binary.

**Bounded by the in-browser WebContainers runtime and by token budget**, which can climb steeply on complex apps. The escape hatch is portability. Download the standard React/Vite code, or push via GitHub auto-sync into a dev environment to go past what the browser build handles.

---

## COST

**Free \$0 with no card (1M tokens/mo, 300K daily cap).** Pro \$25/mo (10M+ tokens with rollover). Teams \$30/member/mo, but Enterprise custom. Billing is token-based and hard to predict (complex apps have reportedly run to \$1,000+), and the caps themselves changed twice in early 2026.

---

## ECOSYSTEM & EXIT

**Low lock-in by design.** Output is standard React/Vite with no proprietary format, GitHub auto-sync mirrors it to a repo, and you can download and self-host anywhere or deploy to Netlify. Nothing ties the code to Bolt, so the exit ramp is essentially free.

---

## WHERE IT STANDS OUT

**Zero lock-in paired with an instant**, install-free in-browser full-stack runtime (WebContainers). The app builds and runs client-side as you prompt, output is standard React/Vite, and GitHub auto-sync plus Netlify deploy export and host anywhere.

---

## LIMITATIONS

**Unpredictable token-based billing that has reportedly reached \$1,000+ on complex apps**, with caps that changed twice in early 2026. Bounded by the in-browser WebContainers runtime. Web only, no native mobile. Generated code still needs review before shipping.

---

Pricing and limits read from **Bolt.new's own site** . Verified to the edition date above.



# Replit

## BUILT FOR

Developers, students, and teams wanting a full cloud IDE with AI agent, hosting, and collaboration across many languages, well beyond web alone.

Best fit: When you need a coding environment plus AI assistance and deployment, not a single-purpose app generator or visual builder.

---

## CHOOSE THIS IF YOU ARE

Founders wanting a full cloud IDE plus an AI Agent that can build, run, host, and deploy in one place.

### BEST WHEN

You want Agent 3 to build autonomously WHILE you retain a dev environment (terminal, packages, databases, deployments, collaboration) so you can graduate from prompting to actual coding without leaving; Core \$20/mo includes ~\$25 usage credits, Agent 3 runs up to 200 min autonomously.

### AVOID IF

You want the slickest no-code full-stack UX, native mobile, or you're non-technical and the IDE surface overwhelms. The power comes with more complexity than pure prompt builders.

---

## WHAT IT DOES

**The Agent works in billable 'checkpoints',** and there is real code access, a built-in database, and one-click deploys (autoscale or reserved VM).

**App code inside a full cloud IDE with a built-in database and one-click deployments (autoscale, which bills for requests/compute, or a dedicated reserved VM).** The AI Agent produces work in discrete checkpoints, each a billable unit, and you keep full code access throughout.

**Fast agentic building via the AI Agent,** though the full IDE surface adds a little overhead versus a pure prompt box. Each Agent task lands as a checkpoint, and one-click deploys mean you can move from build to hosted quickly once the prototype is running.

**Low-to-moderate: The AI Agent does much of the building,** but the full IDE deliberately exposes real code, a database and deploy infra (autoscale vs reserved VM). That surface helps you fix and extend, but it also means more real concepts than a pure prompt builder.

**Primarily web apps, APIs and static sites,**

which is the core of what it builds and hosts. A separate mobile app product exists in Replit's lineup, but the mainstream path here is web, not native mobile.

**A higher ceiling than pure prompt builders because there is a full IDE with real code**

**access**, a built-in database and real deploy infra (autoscale or reserved VM). You take on more hands-on work, and going bigger leans on managing that infrastructure yourself.

---

**COST**

**Core \$20/mo bundles \$25 in usage credits**, up to 5 collaborators and unlimited workspaces; Pro \$100/mo adds teams (up to 15 builders), pooled credits and rollover. Agent billing is effort-based: each completed task is a checkpoint, with simple changes typically under \$0.25 and complex ones bundled higher.

---

**ECOSYSTEM & EXIT**

**The code itself is standard and yours to take**, but the tie is hosting. The integrated experience (built-in database and one-click autoscale/reserved-VM deploys) lives on Replit. Leaving means moving off that hosting and rewiring deployment elsewhere. Code portable, infra sticky.

---

**WHERE IT STANDS OUT**

**One place to build, run and host (a full IDE, built-in database and one-click deploys) combined with an AI Agent whose effort-based billing keeps small changes typically under \$0.25 per checkpoint.** Build-plus-host in a single environment is the differentiator.

---

**LIMITATIONS**

**Effort-based usage billing can be unpredictable as tasks grow (each checkpoint is a billable unit) and the model changed in 2026.** You're tied to Replit hosting for the integrated experience. The IDE surface also means more real concepts to manage than a pure prompt builder.

---

Pricing and limits read from [Replit's own site](#) . Verified to the edition date above.

# FlutterFlow

## BUILT FOR

Product teams and semi-technical builders creating cross-platform mobile and web apps visually on Flutter, with custom code when needed. Best fit: When you specifically want native iOS/Android apps from a visual builder with real Flutter export, over web-only generators or heavy hand-coding.

---

## CHOOSE THIS IF YOU ARE

Founders who need real native iOS AND Android mobile apps, beyond web alone.

### BEST WHEN

You want a visual builder that outputs actual Flutter code, with Firebase/Supabase integration, custom functions, and App Store/Play Store deployment, the only tool here purpose-built for cross-platform native mobile; ~\$39/mo Basic up through Growth/Business tiers.

### AVOID IF

You're building a web-first SaaS, you want pure prompt-to-app with no visual-builder learning curve, or you need complex backend logic better hand-coded; FlutterFlow's strength is the mobile UI layer, and deep logic still means custom code.

## WHAT IT DOES

**It outputs real Dart source with full code export**, and connects to Firebase/Supabase, external APIs and GitHub.

**Clean, real Flutter/Dart source for native iOS**, Android and web from one project, generated via visual building plus AI. Code export is available from the Basic plan, so you can download the Dart and continue developing outside the platform, with Firebase/Supabase and API hookups.

**Slower to a first prototype than prompt-only tools because of the visual-development model (you assemble screens and logic instead of describing them once)**, but the payoff is a genuine native app for iOS, Android and web, not a web page.

**Moderate: the visual-development model and underlying Flutter/Dart concepts carry a learning curve**, steeper than pure prompt builders. To fix or extend meaningfully you work with the exported Dart and integrations like Firebase/Supabase, which rewards some engineering comfort.

## Native iOS and Android plus web from a single

**Flutter project**, the native-mobile pick of this group. It generates real Dart for the app stores (store deploy needs your own Apple/Google developer accounts) instead of just a responsive web page.

**High for mobile.** Full Flutter/Dart export means you can extend indefinitely in a mainstream framework outside the platform, wiring in Firebase/Supabase and APIs. The ceiling is set by your Flutter engineering instead of by the builder itself.

---

## COST

**Basic \$39/mo (annual \$351/yr)**, Growth \$80/mo for the first seat (annual \$720/yr), Business (annual \$1,350/yr), with annual billing running about 25% cheaper than monthly. It's a flat subscription instead of per-message credit burn, so cost stays predictable regardless of iteration volume.

---

## ECOSYSTEM & EXIT

**Low lock-in: it exports clean**, real Flutter/Dart from the Basic plan, and that native code can be downloaded and developed entirely outside the platform. Because Flutter is a mainstream open framework, you keep full ownership and can continue in any Dart toolchain.

---

## WHERE IT STANDS OUT

**The only pick here that produces genuine native mobile apps (real iOS/Android and web from one project) while still handing you clean**, fully-owned Flutter/Dart you can export from the Basic plan and develop anywhere. Native output plus code ownership is the standout.

---

## LIMITATIONS

**A steeper visual-development learning curve than prompt-only tools.** Prices rose in 2026 (Basic \$30 to \$39). Store deployment requires your own Apple and Google developer accounts. Getting full value also assumes comfort with Flutter/Dart to fix and extend the export.

---

Pricing and limits read from [FlutterFlow's own site](#) . Verified to the edition date above.

# v0

## BUILT FOR

Front-end developers and designers, technical, generating React and Tailwind UI components and pages from prompts within the Vercel ecosystem. Best fit: When you need polished UI and component code to drop into a Next.js project, not a full backend or standalone no-code app.

---

## CHOOSE THIS IF YOU ARE

Founders and developers who already own a backend and just need beautiful frontend UI fast.

## BEST WHEN

You want to generate clean, production-ready React + shadcn/ui + Tailwind components from prompts or screenshots and paste them into an existing Next.js codebase. The cleanest frontend output of this group, Plus \$30/user/mo (the old \$20 Premium is legacy-only).

## AVOID IF

You need a full-stack app with auth/DB/backend, native mobile, or a no-code database-backed tool; v0 hands you the front end and leaves the rest of the product to you.

## WHAT IT DOES

**Output is React/Next.js code**, exportable to GitHub, and it inherits Vercel's SOC 2 Type II posture, a front-end-first path into the Vercel stack.

**React/Next.js UI and application code**, synced natively to GitHub with branches and pull requests and deployable straight to Vercel. Code is exportable to GitHub, and the token/credit system that meters generation resets on a monthly cycle.

**Fast for front-end and UI generation**, especially inside a Next.js/Vercel workflow, and the permanent free tier (~\$5 of tokens/mo, roughly 3-8 prompts/day) generates immediately. Deploying to Vercel from the generated code is quick given the native integration.

**Moderate: it's best used by people comfortable with React/Next.js**, since both the generated output and the fixes live in that stack. Front-end work is approachable, but getting full value (and debugging) assumes familiarity with Next.js and the GitHub/Vercel flow.

**Web apps only (React/Next.js), with no native mobile.** It's a front-end/web builder tied to the Vercel stack. Mobile support means responsive web, with no native app.

**Strong for front-end and Next.js work, its home turf.** Full-stack depth leans on the surrounding stack. You extend via the exported React/Next.js code, native GitHub and Vercel deploy instead of an all-in-one backend, so how far it goes depends on that ecosystem.

---

## COST

**Permanent free tier worth ~\$5 of tokens/mo (roughly 3-8 prompts/day).** Paid plans from \$30/mo (Plus; the old \$20 Premium tier is closed to new signups). Generation is metered by a token/credit system that resets monthly, and the caps have moved over time, so the effective prompts-per-day there is can shift between periods.

---

## ECOSYSTEM & EXIT

**Code is exportable via native GitHub (branches and PRs).** The source itself is portable. The soft lock-in is deep Vercel coupling (the workflow assumes Vercel hosting and deploys). Moving off means giving up that tight integration even though the React/Next.js code travels with you.

---

## WHERE IT STANDS OUT

**The tightest Next.js/Vercel/GitHub workflow of the group:** native GitHub branches and PRs, direct Vercel deploy, a permanent free tier, and Vercel's SOC 2 Type II posture behind it. For a front-end team already on Vercel, the integration is the draw.

---

## LIMITATIONS

**Deep Vercel coupling is a soft lock-in even though the React/Next.js code exports to GitHub.** The token caps that meter generation move over time. It's also web only (React/Next.js) with no native mobile, a front-end-first tool instead of a full-stack or mobile one.

---

Pricing and limits read from [v0's own site](#) . Verified to the edition date above.

# Bubble

## BUILT FOR

Non-technical founders and operators building complex, database-backed web apps visually, without writing code, often for MVPs and internal tools.

Best fit: When you need a mature visual platform with workflows and a database and can invest in learning it, over lighter AI code generators.

---

## CHOOSE THIS IF YOU ARE

Non-technical founders wanting deep no-code control over complex web app logic and data.

### BEST WHEN

You're building a database-driven web app or marketplace with intricate workflows and want a mature visual builder with a huge plugin ecosystem, a managed no-code platform instead of AI-generated code you then maintain. Pricing runs on workload-unit tiers (Starter ~\$29/mo up through Growth/Team).

### AVOID IF

You want AI to write real, portable code you own, you need native mobile, or you expect to scale to heavy custom performance; Bubble's workload pricing and platform lock-in can bite at scale.

## WHAT IT DOES

**Launched in 2012, it is one of the oldest**, most established no-code tools.

**A hosted, running web application:** visual pages, a relational database, server-side workflows, and API connectors, all managed inside Bubble. It produces no source code, but your app lives entirely in Bubble's proprietary runtime.

**Fast for data/CRUD apps: a functional MVP in days with no code**, far quicker than hand-coding. Slower than AI prompt-to-app tools for a first screen, but much stronger once you need real backend logic and a database.

**No coding required, but a learning curve.** You must think in databases, workflows, and conditionals. It's genuine visual programming, not templates. Steeper to master than Softr or Glide, gentler than writing code.

**Primarily a responsive web-app platform;** that's its mature core. Native iOS/Android is now offered via a separate (added-cost) mobile plan track sharing the same backend, but the native product is newer and less proven than the web builder.

**Scales to real production SaaS and marketplaces with paying users,** but Workload-Unit costs and performance can strain under high traffic or heavy logic. Not built for compute-intensive apps or anything needing full code-level control.

---

## COST

**Free plan (50k Workload Units, dev-only, no live deploy).** Paid web plans (annual): Starter \$29, Growth \$119, Team \$349/mo. Native mobile is a separate track (\$42+). Costs meter on Workload Units. Overages ~\$0.30 per 1,000 WU. Inefficient apps burn fast.

---

## ECOSYSTEM & EXIT

**There is no exportable source code.** The app is locked to Bubble's proprietary runtime and hosting. Migrating off means a full rebuild elsewhere. Data is exportable (CSV/API). The UI and logic are not. Weakest exit ramp of the group.

---

## WHERE IT STANDS OUT

**The deepest no-code backend of the bunch:** a database, complex multi-step workflows, and a massive plugin and agency ecosystem. Best when a non-coder needs actual application logic, not just a polished front end.

---

## LIMITATIONS

**Proprietary lock-in (no code export),** a genuine learning curve, and unpredictable Workload-Unit costs at scale. Complex apps can hit performance ceilings, and the native-mobile track still trails the web product in maturity.

---

Pricing and limits read from **Bubble's own site** . Verified to the edition date above.



# Build

## BUILT FOR

Hand-code the product with an AI coding agent (Claude Code, Codex, Cursor) instead of a no-code builder. You own real, portable source from day one. The ceiling is your ability to read and fix what the AI writes. Founders with at least some engineering literacy who want a real, ownable product and custom logic. Wrong tool for a true non-coder. You must be able to review, debug, and deploy what the model generates.

---

## CHOOSE THIS IF YOU ARE

Technical founders who can ship code and want zero platform lock-in.

### BEST WHEN

You'd rather pair with an AI coding agent (Claude Code, Codex, Cursor) in a repo (full control over stack, architecture, dependencies, and hosting) than accept a generated-app platform's abstractions. You own portable code from day one and pay only for the AI subscription and your infra.

### AVOID IF

You're non-technical and can't debug generated output, you need a running prototype in an afternoon with no setup, or you need native mobile with a visual layer; DIY trades onboarding speed for control.

## WHAT IT DOES

**Not a builder platform.** It's you plus a very fast junior engineer in your terminal or IDE.

**Actual source code in real frameworks (Next.js, React Native, Python),** plus config, tests, and migrations. No black-box runtime; everything lands in your repo as files you can read, edit, and deploy anywhere.

**Fast for a clickable MVP.** Hours to a day for a CRUD app. Slower than no-code for the first screen, but pulls ahead once you need auth, payments, or custom logic that boxed builders can't express.

**Real and non-negotiable.** The AI writes 80%, but debugging, architecture calls, and "why is prod down" are yours. Non-engineers hit a wall fast. The tool amplifies skill, it doesn't replace it.

**Both, since it writes any framework.** Next.js for web, React Native/Expo or Swift for mobile. But mobile means real build tooling, signing, and app-store review. The AI writes code, not your Apple approval.

**Essentially none on capability.** It's real code. It scales as far as your architecture and skill do. The ceiling is human: your ability to review, test, and maintain what an eager AI happily over-generates.

---

## COST

**A time-vs-money trade, not a flat fee:** ~\$20-200/mo for the AI plan plus API/token burn that spikes with sloppy prompting. Token efficiency is a learned skill. Vague prompts on a big codebase drain credits fast.

---

## ECOSYSTEM & EXIT

**Total and immediate: it's your git repo,** no vendor lock-in, no export step. The strongest exit ramp of any option. The "platform" is just tooling. The asset is code you'd have written anyway, only faster.

---

## WHERE IT STANDS OUT

**High flexibility and ownership at near-no marginal cost per feature.** If you can code even a little, it collapses build time while leaving you a real, scalable, portable product, with no platform tax and no walls.

---

## LIMITATIONS

**Needs engineering judgment you can't fake;** AI writes plausible-but-wrong code and confident security holes, so unreviewed output is a liability. Burns tokens if undisciplined. Not a shortcut for non-technical founders.

---

No vendor page: this route is assembled from the components named above.

# Databases & Backend

Compared: Supabase · Firebase · Neon · Turso · MongoDB Atlas · Convex · Build

# At a glance

| TOOL                 | CHOOSE THIS IF YOU ARE                                                                                                                                                          | STARTING PRICE                                                                                                               | MAIN BENEFIT                                                                                                                                                                                                                                                                                                 |
|----------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Supabase</b>      | <b>Founders who want a relational Postgres backend plus auth</b><br>storage, realtime, and edge functions bundled into one \$0-to-\$25 stack.                                   | <b>A free tier covers early projects</b><br>with the Pro plan at \$25/mo and usage-based overages beyond the included quotas | <b>Delivers a full Postgres backend (database, auth, storage, realtime and edge functions) from one dashboard without stitching separate services together, while keeping the data in standard Postgres.</b><br>The combination of a familiar SQL core and a bundled feature set is its main differentiator. |
| <b>Firebase</b>      | <b>Mobile and realtime app builders who want Google's managed NoSQL</b><br>offline sync, push, and auth in one SDK.                                                             | <b>The Spark tier is free</b><br>while the Blaze plan is pay-as-you-go and billed by usage on a per-operation basis          | <b>Combines realtime data with deep native mobile SDK coverage across iOS, Android, web and Flutter, all inside the Google ecosystem alongside analytics and messaging.</b><br>For mobile-first teams already using Google services, that integration depth is what sets it apart.                           |
| <b>Neon</b>          | <b>Teams that want only serverless Postgres</b><br>with instant database branching per pull request, and will bring their own auth and hosting.                                 | <b>Free plan gives 100 compute-hours</b><br>(CU-hours) per month, 0.5GB storage per project and up to 100 projects           | <b>Instant, cheap branching plus scale-to-zero suits per-PR preview environments and AI/dev-tooling workflows that spin up many short-lived databases, forking is fast and idle branches consume little compute.</b><br>Being wire-compatible Postgres, this arrives without giving up standard tooling.     |
| <b>Turso</b>         | <b>Builders needing edge-local, low-latency reads or a separate SQLite database per customer</b><br>(multi-tenant isolation).                                                   | <b>Free tier covers 100 databases and roughly 5GB</b>                                                                        | <b>Low-latency edge reads via embedded SQLite replicas kept next to the application, which fits read-heavy edge and serverless apps.</b><br>It also has among the cheapest entry pricing here, with a \$4.99/mo Developer tier offering unlimited databases above the free tier.                             |
| <b>MongoDB Atlas</b> | <b>Teams whose data is genuinely document-shaped and want flexible schemas plus an expressive aggregation pipeline</b><br>managed globally.                                     | <b>The free M0 cluster provides 512MB with no backups</b>                                                                    | <b>A flexible document schema plus a large, well-established driver, tooling and hiring ecosystem that most stacks already know, backed by a mature public company across AWS, Azure and GCP.</b><br>Atlas Search and Vector Search extend it without leaving the managed platform.                          |
| <b>Convex</b>        | <b>TypeScript-first teams who want a reactive backend where queries</b><br>mutations, functions and the database are one typesafe system with real-time subscriptions built in. | <b>Free tier includes 1M function calls/mo</b><br>0.5GB database storage and 1GB file storage                                | <b>Provides reactive, realtime data out of the box through a code-first TypeScript model, so query results sync to clients as data changes without extra wiring.</b><br>For realtime and AI apps built in TypeScript, having reactivity as a default instead of an add-on is the standout.                   |
| <b>Build</b>         | <b>Founders with real ops skills or hard data-residency/compliance mandates who want zero vendor markup</b><br>total control of their database.                                 | <b>No tiers, no overage bills</b><br>flat VPS rent (~\$5-40/mo Hetzner/DO for small, more as you add RAM/CPU/disk)           | <b>Total control, no per-row/egress/seat fees, no lock-in, predictable flat cost, and free rein over extensions and tuning.</b><br>Cheapest at scale if you already have ops skill and want to own the whole stack.                                                                                          |

Full detail for every tool follows on its own page.

# Supabase

## BUILT FOR

Seed-to-Series-A teams and full-stack JS/TS developers who want a Postgres database with built-in auth, storage, and realtime instead of stitching separate services together, primarily to ship a relational-backed web or mobile app fast. Best fit: when you want open-source Postgres you can self-host later and need auth plus realtime from one console.

---

## CHOOSE THIS IF YOU ARE

Founders who want a relational Postgres backend plus auth, storage, realtime, and edge functions bundled into one \$0-to-\$25 stack.

### BEST WHEN

You need real SQL with joins, row-level-security policies, and a hosted Postgres you won't outgrow: the free tier gives 500MB DB, 50k MAU, and 2 projects (which auto-pause after ~1 week idle), while Pro at \$25/mo bumps you to 8GB DB, 100k MAU, and \$10 compute credit.

### AVOID IF

You're a mobile-first app leaning on Google's ecosystem and offline sync, or you only want raw serverless Postgres with per-branch previews.

---

## WHAT IT DOES

Because the core is ordinary Postgres rather than a proprietary store, **teams work directly in SQL instead of a vendor-specific data model.**

### Relational database running standard

**PostgreSQL**, so ordinary SQL, constraints and Postgres tooling all apply, with the pgvector extension available for storing and querying embeddings alongside application data, not in a separate vector service.

**One dashboard fronts Postgres plus auth, storage, realtime and edge functions**, so the pieces share a single database rather than separate services. pgvector is available in the same Postgres instance for storing and querying embeddings, keeping vector search alongside relational data instead of a bolt-on service.

**Free projects auto-pause after 7 days without API requests.** The data is retained but the project must be resumed manually before it serves traffic again. Paid projects do not auto-pause. This pause behavior is specific to the free tier.

**A hosted dashboard provisions Postgres plus auth,** storage and realtime. A working backend is available within minutes without assembling services. The managed Cloud path avoids the multi-service setup that self-hosting would otherwise require, shortening the time to a usable backend.

---

### COST

**A free tier covers early projects,** with the Pro plan at \$25/mo and usage-based overages beyond the included quotas. Costs then scale with consumption instead of a flat seat price, so heavier database, storage or function use pushes the bill up past the \$25/mo base.

---

### ECOSYSTEM & EXIT

**The data layer is standard Postgres.** The database itself is portable, but running Supabase yourself means operating a heavy multi-service Docker stack: Postgres, GoTrue for auth, PostgREST, Realtime, Storage and Kong. Portability of the data is high. The operational cost of self-hosting the full platform is not.

---

### TRACK RECORD

**Well-funded and widely adopted,** with a large community around the project. That backing and adoption give it more resources and third-party material than the smaller open-source options here, which lowers the risk of building on it.

**No published uptime SLA below the paid tiers appears in the sources used here.** Free-tier reliability is best-effort with no contractual guarantee. Founders below paid or Enterprise plans have no committed uptime percentage to point to.

---

### WHERE IT STANDS OUT

**Delivers a full Postgres backend (database, auth, storage, realtime and edge functions) from one dashboard without stitching separate services together,** while keeping the data in standard Postgres. The combination of a familiar SQL core and a bundled feature set is its main differentiator.

---

### LIMITATIONS

**Self-hosting the full platform is heavy,** requiring a multi-service Docker stack (Postgres, GoTrue, PostgREST, Realtime, Storage, Kong) to run yourself. On the managed side, free-tier projects pause after a period of inactivity, so the free option is not suited to always-on workloads.

---

Pricing and limits read from [Supabase's own site](#) . Verified to the edition date above.

# Firestore

## BUILT FOR

Solo developers and early startups, often mobile-first with light backend experience, who want a fully managed NoSQL backend with auth, hosting, and push notifications out of the box. Best fit: real-time mobile apps and MVPs where document data and Google's managed infrastructure beat running your own SQL layer.

---

## CHOOSE THIS IF YOU ARE

Mobile and realtime app builders who want Google's managed NoSQL, offline sync, push, and auth in one SDK.

## BEST WHEN

You're shipping iOS/Android or a realtime chat/presence app and want Firestore's client-side offline cache and instant listeners. Spark free covers 50k reads / 20k writes / 20k deletes per day, 1GiB stored, and 2M Cloud Functions calls/month, and it hard-blocks instead of bills you when quotas hit.

## AVOID IF

You need relational joins, complex queries, or transactions across tables, or you fear Blaze's pay-as-you-go read/write metering (\$0.06 per 100k reads) spiking unpredictably at scale.

---

## WHAT IT DOES

**It is tied into the wider Google ecosystem and reached through native mobile SDKs,** positioning it for teams building mobile-first products over SQL-centric backends.

**NoSQL document store organized as collections of documents,** not tables, with realtime sync so clients receive updates as data changes and can keep working offline. It is queried through the Firebase SDKs, not SQL.

**Bundles Firestore and the Realtime Database with Auth, Hosting, Messaging and Analytics,** and adds ML Kit alongside realtime data. The components are tightly bound to the Google ecosystem and consumed through Firebase's mobile SDKs, which deepens integration but also ties the stack to Google's platform.

**There is no pause or cold-start concept for the database itself.** Firestore is a pure request-billed serverless service, so an idle database costs only for stored data and incurs no compute to spin up. Activity simply resumes per-operation billing.

**Console-driven setup that is fast to reach a working state for mobile apps,** with the native SDKs handling client wiring. Because everything runs on Google's managed platform, there is no infrastructure to stand up. In return, setup happens entirely within Google's console.

---

### COST

**The Spark tier is free, while the Blaze plan is pay-as-you-go and billed by usage on a per-operation basis.** That model means cost tracks operation volume directly, which is flexible at low usage but is the source of the per-operation billing surprises flagged as a risk.

---

### ECOSYSTEM & EXIT

**Proprietary in its data model,** Security Rules and indexing, with no self-hosted equivalent to move to. Data can be exported via gcloud, but the rules, triggers and auth around it have no portable form and require a rewrite elsewhere, making this the highest-lock-in option here.

---

### TRACK RECORD

**Backed by Google with a long operating track record and a large surrounding ecosystem.** It is the most established option in this set, though that maturity comes packaged with the proprietary, no-self-hosting model that defines its lock-in.

**Firestore runs on Google Cloud infrastructure,** but no specific published availability SLA percentage appears in the sources used here. The signal is the platform's long operational track record under Google instead of a committed uptime number.

---

### WHERE IT STANDS OUT

**Combines realtime data with deep native mobile SDK coverage across iOS, Android, web and Flutter,** all inside the Google ecosystem alongside analytics and messaging. For mobile-first teams already using Google services, that integration depth is what sets it apart.

---

### LIMITATIONS

**No self-hosting and proprietary throughout,** giving it the highest lock-in here, and its Blaze plan bills per operation, so usage-based charges can surprise as volume grows. Both constraints stem from the closed, Google-hosted model, not being configurable away.

---

Pricing and limits read from [Firebase's own site](#) . Verified to the edition date above.



# Neon

## BUILT FOR

Developers and teams wanting serverless Postgres with database branching, comfortable with SQL and any modern Postgres client, primarily to spin up per-branch preview databases in CI/CD.

---

## CHOOSE THIS IF YOU ARE

Teams that want only serverless Postgres, with instant database branching per pull request, and will bring their own auth and hosting.

### BEST WHEN

You want a Postgres that scales to zero and forks a full copy of prod for every PR/preview like Git branches. The free plan gives up to 100 projects, 0.5GB storage and 100 compute-hours each, and paid compute is usage-based (\$0.106/CU-hour on Launch, \$0.222 on Scale) with no monthly minimum. Since Databricks acquired Neon in May 2025, compute prices have dropped 15-25%.

### AVOID IF

You want auth/storage/realtime included, or you need embedded/edge-local reads per tenant.

## WHAT IT DOES

**Databases can scale to zero and be forked cheaply while staying wire-compatible with standard Postgres.** Acquired by Databricks in May 2025, after which its pricing was restructured.

**Relational PostgreSQL that is wire-compatible with standard Postgres.** Existing drivers and connection strings work unchanged. Its distinctive traits are instant branching and scale-to-zero compute, and pgvector is available for embeddings.

**Database only: serverless PostgreSQL.** You bring your own backend/API layer, auth and storage.

**Compute scales to zero after an idle period,** so an inactive database stops consuming compute. The next query triggers a cold start to resume it. Storage keeps billing continuously while compute is suspended, since storage and compute are separated and priced independently.

**Standard Postgres connection, fast for anyone who knows Postgres.**

## COST

**Free plan gives 100 compute-hours (CU-hours) per month**, 0.5GB storage per project and up to 100 projects. Paid compute is billed per CU-hour (Launch \$0.106, Scale \$0.222) plus storage at \$0.35/GB-month. There is no monthly minimum, so cost is purely usage-based.

---

## ECOSYSTEM & EXIT

**Being wire-compatible standard Postgres**, connection strings and dumps port to any Postgres provider. The core database carries low lock-in. The branching feature is Neon-specific, so workflows built around instant branches would need rebuilding on a plain Postgres host.

---

## TRACK RECORD

**Acquired by Databricks in May 2025**, after which pricing was restructured in 2025-2026: storage fell from \$1.75 to \$0.35/GB-month, the free allowance doubled from 50 to 100 CU-hours, and the \$5/mo minimum was removed. The changes were largely founder-favorable, but they are recent.

**No published uptime SLA figure for Neon appears in the sources used here.** What is documented is the operational model (scale-to-zero compute with a cold start on resume), not a committed availability percentage.

## WHERE IT STANDS OUT

**Instant, cheap branching plus scale-to-zero suits per-PR preview environments and AI/dev-tooling workflows that spin up many short-lived databases**, forking is fast and idle branches consume little compute. Being wire-compatible Postgres, this arrives without giving up standard tooling.

---

## LIMITATIONS

**Cold-start latency on the first query after scale-to-zero is the main operational tradeoff**, and storage keeps billing while compute is suspended. The May 2025 Databricks acquisition, though it brought founder-favorable pricing, also adds some roadmap uncertainty for a recently-changed product.

---

Pricing and limits read from [Neon's own site](#) . Verified to the edition date above.

# Turso

## BUILT FOR

Developers building edge and local-first apps who want SQLite (libSQL) replicated close to users, primarily for low-latency reads at the edge and embedded or offline sync. Best fit: read-heavy, globally distributed apps needing many per-tenant databases, versus a single centralized Postgres or Mongo cluster.

---

## CHOOSE THIS IF YOU ARE

Builders needing edge-local, low-latency reads or a separate SQLite database per customer (multi-tenant isolation).

### BEST WHEN

You're doing database-per-tenant at scale or serving read-heavy globally-distributed users: libSQL replicates SQLite to the edge with no cold starts, and the free tier covers 100 databases, 5GB storage and 500M row reads/month, with Developer at \$4.99/mo and Scaler at \$24.92/mo.

### AVOID IF

You need heavy concurrent writes, Postgres extensions/features, or a single large shared database, or a document model.

---

## WHAT IT DOES

**It is SQLite-compatible via the libSQL fork and uses a single-writer model.**

**Relational and SQLite-compatible through the libSQL fork.** The SQLite dialect and tooling apply. It uses a single-writer model, which fits read-heavy edge workloads more than write-concurrent ones, with embedded replicas serving local reads.

**Database only: managed SQLite/libSQL at the edge with embedded replicas**, bring your own backend/auth/storage.

**There is no documented per-database pause or cold-start cycle**, but requests are served without an explicit resume step. Combined with edge-resident embedded replicas, reads are answered locally, not by waking a suspended primary.

**SQLite-simple to start.** Create a database in seconds via CLI or dashboard with no server to manage; the real learning curve is embedded/edge replication and per-tenant databases, which is new even to SQL-comfortable teams.

---

## COST

**Free tier covers 100 databases and roughly 5GB.** Developer is \$4.99/mo (unlimited databases, ~9GB). Paid: Scaler \$24.92/mo (24GB) and Pro \$416.58/mo (50GB, adds SSO/BYOK/HIPAA/SOC2), with custom Enterprise above, per Turso's current pricing page.

---

## ECOSYSTEM & EXIT

**libSQL is SQLite-compatible.** Local database files can be read with standard SQLite tooling, keeping the data itself portable. The edge sync layer and the hosted service are Turso-specific, so it is the embedded-replica setup, not the data that ties you to the platform.

---

## TRACK RECORD

**The smallest and newest company of the five,** with pricing tiers that have been introduced or revised comparatively recently. That leaves less of a track record for how stable the current Free, Developer, Scaler and Pro structure will prove over time.

**As the smallest and newest platform of the five,** Turso has less independently documented reliability history, and no widely published uptime SLA outside Enterprise appears in the sources used here: limited public incident history and no committed availability number below Enterprise.

---

## WHERE IT STANDS OUT

**Low-latency edge reads via embedded SQLite replicas kept next to the application,** which fits read-heavy edge and serverless apps. It also has among the cheapest entry pricing here, with a \$4.99/mo Developer tier offering unlimited databases above the free tier.

---

## LIMITATIONS

**The SQLite single-writer model constrains write concurrency.** It fits read-heavy workloads better than write-heavy ones. As the smallest and newest vendor here it also has the smallest ecosystem and hiring pool, and no widely published uptime SLA outside Enterprise.

---

Pricing and limits read from [Turso's own site](#) . Verified to the edition date above.

# MongoDB Atlas

## BUILT FOR

Teams from startup to enterprise and developers preferring document/NoSQL modeling at any technical level, primarily for flexible-schema apps, search, and analytics on managed MongoDB across clouds. Best fit: when data is document-shaped and evolving and you need mature scaling, sharding, and full-text or vector search.

---

## CHOOSE THIS IF YOU ARE

Teams whose data is document-shaped and want flexible schemas plus an expressive aggregation pipeline, managed globally.

### BEST WHEN

Your entities are nested/variable (catalogs, events, user profiles with sparse fields) and you'll lean on aggregation pipelines and horizontal sharding: the free M0 cluster gives 512MB, the Flex tier caps at \$30/mo, and dedicated M10 starts ~\$0.08/hr (~\$58/mo) climbing with RAM.

### AVOID IF

Your data is relational with joins and constraints, or you want a full BaaS with auth/realtime baked in.

---

## WHAT IT DOES

**It stores BSON documents with a dynamic schema and layers Atlas Search and Vector Search on top of the managed clusters.**

**NoSQL document store: data is held as BSON documents with a dynamic schema instead of fixed relational tables**, but document shape can vary within a collection. It suits applications whose data models change or nest, and uses the MongoDB wire protocol and drivers.

**Database only: managed NoSQL document DB;** bring your own backend/auth/storage (Atlas adds Search/Charts as add-ons).

**Free and shared clusters auto-pause after prolonged inactivity and must be resumed**, whereas dedicated M10+ clusters run continuously and do not pause. Idle behavior therefore tracks the tier. The free M0 can sleep, while paid dedicated capacity stays on.

**Quick to spin up, but needs NoSQL data-modeling discipline.**

## COST

**The free M0 cluster provides 512MB with no backups.** Above it sits a usage-based Flex tier capped at roughly \$30/mo, and dedicated clusters that start at the M10 size for \$0.08/hour billed by uptime. The step to dedicated is where production-grade features appear.

## ECOSYSTEM & EXIT

**mongodump/mongorestore and the official drivers work against any MongoDB-compatible deployment;** portability within the MongoDB ecosystem is high and moving between hosts or self-managed clusters is straightforward. Such lock-in as exists comes from the document model itself.

## TRACK RECORD

**The most mature offering in this set,** run by a public company (NASDAQ: MDB), with a long and comparatively stable pricing structure spanning the free M0, Flex and dedicated M10+ tiers. It is the least likely of the five to see abrupt free-tier or pricing changes.

**Dedicated M10+ clusters carry an uptime SLA with service credits,** while the free and shared tiers do not. The exact current SLA percentage is not documented in the sources used here, though the tiered-SLA structure itself is.

## WHERE IT STANDS OUT

**A flexible document schema plus a large,** well-established driver, tooling and hiring ecosystem that most stacks already know, backed by a mature public company across AWS, Azure and GCP. Atlas Search and Vector Search extend it without leaving the managed platform.

## LIMITATIONS

**The free M0 tier has no backups and no SLA.** Production-grade reliability requires stepping up to a dedicated M10+ cluster from \$0.08/hour. The document model also demands NoSQL modeling discipline, since schema flexibility can turn into inconsistency without deliberate design.

---

Pricing and limits read from [MongoDB Atlas's own site](#) . Verified to the edition date above.

# Convex

## BUILT FOR

Full-stack TypeScript developers at early-stage startups who want a reactive backend bundling database, server functions, and realtime into one strongly typed platform. Best fit: TS-first apps needing live-updating queries and server logic without managing a separate database, API layer, and websocket infrastructure.

---

## CHOOSE THIS IF YOU ARE

TypeScript-first teams who want a reactive backend where queries, mutations, functions and the database are one typesafe system with real-time subscriptions built in.

## BEST WHEN

You want end-to-end type safety and live-updating queries with zero websocket plumbing: write server functions in TS and the client re-renders automatically. Free covers up to 6 devs, 1M function calls, 20 GB-hours compute, 0.5GB database and 1GB file storage, Pro at \$25/mo.

## AVOID IF

You need raw SQL, existing Postgres tooling/ORMs, or portable standard queries, or a document store at scale.

---

## WHAT IT DOES

**It bundles reactive queries, file storage,** vector search and auth. Open-sourced in 2024 with self-hosting added, so the code-first model is not strictly tied to the vendor cloud.

**Its own reactive document/relational store with TypeScript queries and realtime sync.** It is not a standalone SQL/NoSQL engine you bring.

**TypeScript functions, reactive queries,** file storage, vector search and built-in auth are combined into one reactive platform. Query results propagate to clients as the underlying data changes. Realtime sync is part of the core, not a separate service you wire up yourself.



**Cold starts can occur after a deployment sits**

**idle**, but Convex does not aggressively scale to zero. A deployment is only treated as idle after 30 days with no function calls (deploy fee then waived). For most small apps the cold-start delay is minor.

**A code-first setup where backend logic is written as TypeScript functions and realtime sync works out of the box**, so reactive behavior does not need separate configuration. Time to value depends on comfort with the code-first workflow instead of clicking through a dashboard.

---

**COST**

**Free tier includes 1M function calls/mo**, 0.5GB database storage and 1GB file storage. Paid: Professional \$25/developer/mo (50GB database storage, 25M function calls); Business/Enterprise from ~\$2,500/mo minimum. Cost scales by usage (function calls, storage) plus a per-developer seat.

---

**ECOSYSTEM & EXIT**

**Open-sourced in 2024 with self-hosting available since early 2025**, running under Docker with data in either SQLite or Postgres. That combination hedges lock-in for a code-first platform. The backend logic and its store can be moved off the vendor cloud onto self-managed infrastructure.

---

**TRACK RECORD**

**Newer than the incumbents but funded**, and its codebase was open-sourced in 2024, which is a community hedge. If the company's direction changed, the self-hostable code remains available. Maturity is lower, but the open source acts as a fallback.

**Publishes a public status page**

**([status.convex.dev](https://status.convex.dev))**. Third-party monitors report ~99.16-99.73% 30-day uptime vs Convex's stated ~99.99% availability target (a target, not a contractual SLA), with 40+ mostly-minor incidents logged over the ~5 months to Jul 2026 (e.g. a 15 Jul 2026 failure pushing Node Actions). Fine for early apps, but watch the SLA gap.

---

**WHERE IT STANDS OUT**

**Provides reactive, realtime data out of the box through a code-first TypeScript model**, so query results sync to clients as data changes without extra wiring. For realtime and AI apps built in TypeScript, having reactivity as a default instead of an add-on is the standout.

---

**LIMITATIONS**

**It is a newer platform than the incumbents and carries a shorter track record**, and its code-first model differs from the dashboard-driven style of most BaaS tools. Teams expecting a click-through console instead of writing TypeScript functions face an adjustment.

---

Pricing and limits read from **Convex's own site** . Verified to the edition date above.



# Build

## BUILT FOR

The build-it-yourself lane: rent a raw Linux box, install Postgres yourself, hand-write SQL and your own API. Max control and zero per-seat fees. You are now the DBA, SRE, and on-call.

---

## CHOOSE THIS IF YOU ARE

Founders with real ops skills or hard data-residency/compliance mandates who want zero vendor markup and total control of their database.

### BEST WHEN

Regulation forces data to live on infrastructure you control, or you're cost-optimizing a predictable workload: a \$5-20/mo Hetzner/DigitalOcean VPS running Postgres costs a fraction of managed tiers, and you own tuning, extensions, and residency completely.

### AVOID IF

You're pre-seed and should be shipping product, not owning backups, patching, failover, and connection pooling, and 3am pages.

## WHAT IT DOES

**No dashboard, no managed magic.** Just Postgres and your own glue.

**Postgres by default: relational/SQL**, ACID, JSONB for semi-structured data, extensions (PostGIS, pgvector). Swap in MySQL, SQLite, or add Redis/Mongo yourself if the workload needs it. You choose the engine.

**Nothing bundled.** You assemble it. Supabase-OSS adds auth, storage, realtime and PostgREST if the whole stack is self-hosted. Connects to anything. It's your own API, open to any client, queue or service you wire up.

**No auto-pause, no cold starts:** the box runs 24/7 and you pay whether idle or busy. Upside: no first-query lag ever. Downside: you burn rent on an idle MVP and must script your own scale-to-zero if you want it.

**Slow: hours to days.** Provision VPS, secure SSH/firewall, install Postgres, configure backups, write migrations and the API layer. Docker-compose or Coolify shortcuts it, but time-to-first-query is far behind a managed one-click.

## COST

**No tiers, no overage bills: flat VPS rent (~\$5-40/mo Hetzner/DO for small, more as you add RAM/CPU/disk).** Cost is your TIME to set up and maintain. Beats managed pricing only long-term if you keep using it and don't value your hours highly.

## ECOSYSTEM & EXIT

**Zero lock-in: it IS self-hosting.** Plain Postgres dumps move anywhere. No proprietary API to unwind. Maximum portability is the whole point and the strongest reason to DIY over a managed vendor.

## TRACK RECORD

**Postgres is decades-mature with a massive community,** but support is DIY: Stack Overflow, docs, forums. No SLA, no support desk, no one to page at 3am but you. Fine for MVPs, risky once real users depend on uptime.

**Only as reliable as you make it.** Single VPS = single point of failure. No automatic failover, replication, or PITR unless you build it. Managed rivals give HA out of the box. Here reliability is whatever your engineering makes it.

## WHERE IT STANDS OUT

**Total control, no per-row/egress/seat fees,** no lock-in, predictable flat cost, and free rein over extensions and tuning. Cheapest at scale if you already have ops skill and want to own the whole stack.

## LIMITATIONS

**Breaks past small scale: no built-in HA,** backups, or scaling without heavy work, but security is on you (a misconfigured Postgres is a breach). No polished console. Ops burden compounds, the classic DIY trap as users grow.

---

No vendor page: this route is assembled from the components named above.

# Authentication

Compared: Clerk · Auth0 · Supabase Auth · WorkOS · Better Auth · Stytle · Build

# At a glance

| TOOL          | CHOOSE THIS IF YOU ARE                                                                                                                                                      | STARTING PRICE                                                                                           | MAIN BENEFIT                                                                                                                                                                                                                                                                                                   |
|---------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Clerk         | React/Next.js founders who want prebuilt drop-in auth UI organizations, and MFA working in an afternoon.                                                                    | Free tier covers 50,000 MRU (monthly retained users, a narrower unit than MAU) as of the Feb 2026 update | <b>Its edge is developer experience and pre-built UI.</b> Drop-in components and Next.js/React SDKs get hosted auth live faster than assembling flows yourself. That speed-to-integration is what distinguishes it within its target ecosystem.                                                                |
| Auth0         | Teams needing a mature, compliance-heavy identity platform with dozens of social enterprise connections and audit-grade features.                                           | Free tier covers 25,000 MAU and notably includes SSO SCIM and AI features                                | <b>Stands out for the breadth of its enterprise identity feature set and protocol support, spanning B2C and B2B with adaptive MFA, attack protection, and SSO/SCIM.</b> It covers the widest range of identity requirements of this group.                                                                     |
| Supabase Auth | Teams that want authorization enforced in the database itself users stored as Postgres rows governed by row-level-security policies, with auth bundled at no separate cost. | Free tier covers 50,000 MAU                                                                              | <b>The cheapest integrated auth at scale for Supabase users.</b> 100K MAU is included in the \$25/mo Pro plan (\$0.00325/MAU only above 100K), versus roughly \$1,000/mo on Clerk, while keeping users in your own Postgres with RLS integration. Cost and database integration are its distinguishing points. |
| WorkOS        | B2B SaaS founders who need to close enterprise deals gated on SAML SSO SCIM directory sync.                                                                                 | AuthKit is free up to 1,000,000 MAU an outlier among these tools                                         | <b>Pairs a free auth tier up to 1,000,000 MAU with specialist enterprise SSO and SCIM priced per connection.</b> The combination of a high free auth ceiling and dedicated enterprise connectors is what sets it apart for B2B teams.                                                                          |
| Better Auth   | TypeScript teams who want to fully own their auth storing users in their own DB with no per-MAU fees or vendor lock-in.                                                     | Free and open source with commercial use allowed and no per-MAU fee at any volume                        | <b>Full control with zero per-user fees: an open-source, self-hosted library that keeps users in your own database and charges nothing per MAU.</b> Ownership of the runtime and the absence of metered pricing are its defining traits.                                                                       |
| Stytch        | Teams wanting API-first, headless passwordless biometric flows plus built-in fraud and device-fingerprinting.                                                               | Free up to 10,000 MAU (both B2C and B2B), then ~\$0.05/MAU with volume discounts and no hard caps        | <b>The auth + fraud combo: passwordless-native login plus built-in device fingerprinting and bot protection in one API, paired with a 10k-MAU free tier and usage-based pay-as-you-go pricing.</b>                                                                                                             |
| Build         | Founders who want maximum control and no lock-in SuperTokens self-hosted, Auth.js for Next, or Lucia primitives to learn the internals.                                     | No per-MAU pricing and no cliff which is the exact reason to DIY                                         | <b>Kills MAU billing, keeps user data in your DB, and bends to any custom flow.</b> Cheapest and most flexible for MVPs and cost-sensitive products where you control the stack and don't need enterprise features day one.                                                                                    |

Full detail for every tool follows on its own page.

# Clerk

## BUILT FOR

Startups and product teams on React/Next.js who want prebuilt auth UI, user management, and organizations dropped in fast with low-to-mid backend effort. Best fit: B2C and B2B SaaS wanting polished drop-in components and multi-tenant orgs quickly, versus hand-wiring flows or configuring a heavier identity platform.

---

## CHOOSE THIS IF YOU ARE

React/Next.js founders who want prebuilt drop-in auth UI, organizations, and MFA working in an afternoon.

### BEST WHEN

You're on React/Next and want prebuilt SignIn components, user profiles, and B2B org management without building screens. The free tier is 50k MRU (monthly retained users, a narrower unit than the MAU rivals meter, so bills run lower), Pro is \$25/mo, Business \$300/mo.

### AVOID IF

You're not in the React/Next ecosystem, need turnkey enterprise SSO/SCIM reselling, or want to fully own user data with no per-user fees.

---

## WHAT IT DOES

**It runs as a managed identity layer instead of a library**, so the user store and auth flows live on Clerk's infrastructure and integrate through its Next.js/React SDKs.

**The fastest to get live here**, since drop-in pre-built components handle the auth UI, not requiring you to build flows. That speed is the main reason teams reach for it, and it pairs with the Next.js/React SDKs where the components are best supported.

**Supports passwords, social login**, passwordless, MFA, and device and session management, with bot and abuse protection included. The method coverage is delivered through the hosted service and its pre-built components instead of requiring you to assemble each flow yourself.

**Available on higher plans but not the product's primary focus**, which is developer-oriented hosted auth and pre-built UI. Enterprise SSO is present, not the specialty. Teams whose main need is SAML/SCIM at scale are not its core audience.

**Ships with built-in organizations and multi-tenant primitives as part of the hosted service**, so team and tenant structures are provided instead of modeled from scratch. This aligns with its developer-oriented, drop-in positioning.

---

## COST

**Free tier covers 50,000 MRU (monthly retained users, a narrower unit than MAU) as of the Feb 2026 update.** The Pro plan is \$25/mo and also includes 50,000 MRU, with additional users at \$0.02/MRU (volume discounts higher up), at 100K MRU that works out to roughly \$1,025/mo, the steepest per-user scaling among these five as volume grows.

---

## ECOSYSTEM & EXIT

**Users live in Clerk's hosted store.** Export is supported, but because the auth flows and user store run on Clerk's infrastructure, moving off requires re-integration into a new provider instead of just relocating data.

---

## TRACK RECORD

**Well-funded and widely adopted**, with developer support that matches its developer-first positioning. It is an established hosted vendor, though newer than Okta-owned Auth0, and its standing rests on adoption within the React/Next.js community.

---

## WHERE IT STANDS OUT

**Its edge is developer experience and pre-built UI.** Drop-in components and Next.js/React SDKs get hosted auth live faster than assembling flows yourself. That speed-to-integration is what distinguishes it within its target ecosystem.

---

## LIMITATIONS

**MAU pricing scales steeply as you grow (roughly \$1,000/mo at 100K MAU, the steepest here)**, and the product is at its best within the React/Next.js ecosystem. Teams on other stacks, or scaling to high user counts, feel the cost and fit constraints most.

---

Pricing and limits read from [Clerk's own site](#) . Verified to the edition date above.

# Auth0

## BUILT FOR

Mid-market to enterprise teams and security-conscious engineers needing broad protocol support, extensible rules/actions, and compliance across any stack or language. Best fit: complex identity needs (many connections, enterprise SSO, fine-grained extensibility) where breadth and maturity outweigh higher per-user cost.

---

## CHOOSE THIS IF YOU ARE

Teams needing a mature, compliance-heavy identity platform with dozens of social and enterprise connections and audit-grade features.

### BEST WHEN

You need proven SOC2/HIPAA posture, many identity providers, and features like audit-log streaming to Datadog/Splunk and per-org RBAC: the free plan covers 25k MAU (now bundling one enterprise SSO connection and SCIM) across Auth0's public-cloud regions (US, EU, UK, AU, JP, CA), B2C Essentials starts \$35/mo for 500 MAU and B2B Essentials \$150/mo.

### AVOID IF

You're cost-sensitive and early.

---

## WHAT IT DOES

It offers broad protocol and SDK support with **Universal Login and a Rules/Actions extensibility model**, positioning it as enterprise identity infrastructure, not a lightweight developer tool.

**Universal Login gets a basic setup running quickly**, but the platform's depth means deeper configuration through Rules/Actions takes more effort than Clerk's drop-in approach. It is framework-agnostic, so it fits most stacks at the cost of more upfront wiring.

**Broad authentication method support backed by adaptive MFA and attack protection at the platform level.** Its enterprise-grade security posture and wide protocol/SDK support are core to the product, aimed at organizations with varied identity requirements.

**Enterprise connections are a core capability:**

SSO and SCIM are included even in the free 25,000-MAU tier and expand at the Professional plans.

Broad protocol support makes enterprise identity central to the product, not an add-on bolted onto a consumer tool.

**B2B plans include organizations for multi-tenant apps**, aligning with its enterprise identity focus across B2C and B2B. Multi-tenancy is a planned platform feature, not something you build in your own data layer.

---

**COST**

**Free tier covers 25,000 MAU and notably includes SSO**, SCIM and AI features. Paid plans step up sharply: Essentials B2C \$35/mo, Essentials B2B \$150/mo, Professional B2C \$240/mo, Professional B2B \$800/mo, jumps that startups describe as a 'growth penalty'.

---

**ECOSYSTEM & EXIT**

**Standard protocols aid portability.** Basic identity data and flows are movable. The friction is that deep customizations built with Actions and Rules are Auth0-specific and would need to be rebuilt on another platform.

---

**TRACK RECORD**

**The most mature option here, owned by Okta**, carrying enterprise trust and support. Its long track record and enterprise-grade platform make it the established choice where organizational credibility and support depth matter.

---

**WHERE IT STANDS OUT**

**Stands out for the breadth of its enterprise identity feature set and protocol support**, spanning B2C and B2B with adaptive MFA, attack protection, and SSO/SCIM. It covers the widest range of identity requirements of this group.

---

**LIMITATIONS**

**Steep tier jumps that startups call a 'growth penalty' (Essentials at \$35-\$150/mo up to Professional at \$240-\$800/mo)**, and it requires more configuration than a drop-in tool like Clerk. Cost escalation and setup effort are the main tradeoffs.

---

Pricing and limits read from [Auth0's own site](#) . Verified to the edition date above.



# Supabase Auth

## BUILT FOR

Teams already on Supabase Postgres wanting authentication tied directly to their database with row-level security, in a JS/TS-friendly workflow. Best fit: when you're inside the Supabase stack and want auth, RLS, and data living in one project, not adding a separate identity vendor.

---

## CHOOSE THIS IF YOU ARE

Teams that want authorization enforced in the database itself: users stored as Postgres rows governed by row-level-security policies, with auth bundled at no separate cost.

## BEST WHEN

You want auth integrated with RLS policies so authorization lives in the database: free covers 50k MAU (bundled into the same \$0/\$25 Supabase plan), no separate vendor or bill.

## AVOID IF

You're not on Supabase (the integration is the whole point), need polished prebuilt auth components and org management, or require enterprise SAML/SCIM connections sold per-customer.

## WHAT IT DOES

**It ships inside Supabase, never as a standalone product**, and is used through Supabase's SDKs.

**Straightforward to enable if you are already on Supabase**, but it ships with less pre-built UI than Clerk. You build more of the interface yourself. You get tighter integration with Postgres and RLS in return for more front-end work.

**Supports passwords, social and passwordless sign-in**, integrated with Postgres row-level security so authorization is enforced in the database. Security is tied to the RLS model, not a separate policy layer, keeping auth and data rules in one place.

**Available but not the platform's enterprise-SSO strength**, which lies in Postgres and RLS integration, not in enterprise connectors. Teams needing SAML/SCIM as a primary requirement fall outside its main design focus.

**Multi-tenancy is modeled in your own Postgres schema and enforced with row-level security.** There is no built-in organizations feature. You design the tenant structure yourself, consistent with its database-centric approach.

---

## COST

**Free tier covers 50,000 MAU.** The \$25/mo Pro plan includes 100,000 MAU, then \$0.00325/MAU beyond, so 100K MAU is covered by the \$25 Pro base and 150K runs about \$187/mo. That per-user rate makes it the cheapest at scale of this group. 100K users cost roughly \$1,000/mo on Clerk, a gap driven by the per-user rate.

---

## ECOSYSTEM & EXIT

**Users live in your own Supabase Postgres database,** which keeps the data close to you, though the auth itself remains tied to Supabase. Portability of the underlying user records is better than a fully external store, but you are still bound to the Supabase platform.

---

## TRACK RECORD

**Backed by Supabase's funding and a large community,** giving it established institutional and ecosystem support. Trust here derives from the broader Supabase platform, not from a standalone auth track record.

---

## WHERE IT STANDS OUT

**The cheapest integrated auth at scale for Supabase users.** 100K MAU is included in the \$25/mo Pro plan (\$0.00325/MAU only above 100K), versus roughly \$1,000/mo on Clerk, while keeping users in your own Postgres with RLS integration. Cost and database integration are its distinguishing points.

---

## LIMITATIONS

**Ships with less pre-built UI than Clerk,** so you build more of the interface yourself, and it is most valuable only if you are already on Supabase. Outside that platform, the integration advantages that justify it largely disappear.

---

Pricing and limits read from [Supabase Auth's own site](#) . Verified to the edition date above.

# WorkOS

## BUILT FOR

B2B SaaS companies moving upmarket whose engineers need enterprise features (SSO/SAML, SCIM directory sync, audit logs) bolted onto an existing auth setup. Best fit: when your app already has login and you now need SAML and SCIM to close enterprise deals, versus replacing your whole auth stack.

---

## CHOOSE THIS IF YOU ARE

B2B SaaS founders who need to close enterprise deals gated on SAML SSO and SCIM directory sync.

### BEST WHEN

A big customer says they need SSO with Okta/Entra and SCIM provisioning and you must ship it fast: AuthKit is free up to 1M MAU for core auth, and enterprise SSO uses per-connection volume tiers (\$125/connection at 1-15, sliding to \$50 at 101-200), with Directory Sync a separate identically-priced SKU.

### AVOID IF

You're building consumer B2C auth, or you'll have many tiny enterprise customers. The per-connection pricing (doubled if you need SCIM too) becomes brutal.

## WHAT IT DOES

**The design separates everyday auth from the enterprise connectors**, which are provisioned and billed one connection at a time instead of by user volume.

**AuthKit provides a hosted UI**, and the platform is oriented toward standing up enterprise SSO quickly for B2B teams. Setup centers on provisioning per-connection SSO/SCIM, not assembling consumer auth flows from scratch.

**Core authentication runs through AuthKit**, extended by enterprise SSO/SAML and SCIM directory sync provisioned per connection. The emphasis is the enterprise connectors that larger B2B buyers require, not a broad consumer method set.

**The specialist here: SSO is \$125 per connection for the first 15**, tiering down to \$50 at 101-200 connections, and SCIM follows the same ladder. A connection needing both SSO and SCIM is billed twice at the entry tier (\$125 + \$125).

**Organization-centric by design, oriented around B2B SaaS where each customer is an organization with its own SSO/SCIM**

**connections.** The org model matches its per-connection enterprise-readiness focus.

---

#### COST

**AuthKit is free up to 1,000,000 MAU**, an outlier among these tools. Enterprise SSO is not priced per MAU but per connection: \$125 each for the first 15 connections. Cost scales with the number of enterprise customers connected instead of with total user volume.

---

#### ECOSYSTEM & EXIT

**Standards-based on SAML and SCIM**, which aids portability, with the enterprise side billed per connection. Reliance on open standards keeps migration friction lower than proprietary flows would, consistent with its connection-based model.

---

#### TRACK RECORD

**An established B2B-focused vendor trusted for enterprise readiness**, with a product built around the SSO/SCIM needs of enterprise buyers. Its reputation is tied to serving that B2B enterprise-connection use case.

---

#### WHERE IT STANDS OUT

**Pairs a free auth tier up to 1,000,000 MAU with specialist enterprise SSO and SCIM priced per connection.** The combination of a high free auth ceiling and dedicated enterprise connectors is what sets it apart for B2B teams.

---

#### LIMITATIONS

**Enterprise SSO and SCIM are billed separately per connection (\$125 each at the entry tier, and a connection needing both is billed twice)**, so costs add up as the number of enterprise customers grows instead of staying flat.

---

Pricing and limits read from [WorkOS's own site](#) . Verified to the edition date above.

# Better Auth

## BUILT FOR

TypeScript developers who want a framework-agnostic, open-source, self-hosted auth library with full ownership of code and user data. Best fit: teams that want auth living in their own database and codebase with no per-MAU billing, versus a hosted identity service.

---

## CHOOSE THIS IF YOU ARE

TypeScript teams who want to fully own their auth, storing users in their own DB with no per-MAU fees or vendor lock-in.

### BEST WHEN

You want a framework-agnostic, open-source, TypeScript-native auth library you self-host: free forever, users in your own database, plugins for MFA/passkeys/orgs, and no billing that scales with success.

### AVOID IF

You don't want to own security patching, session hardening, and breach risk, or you need turnkey enterprise SSO/SCIM connections out of the box.

## WHAT IT DOES

It is a distinct project from Lucia (the older library that was sunset in 2025, with its maintainer redirecting users toward Better Auth), not a rename of it. Your app and database own the runtime behavior, with functionality extended through plugins.

**More setup than the hosted options.** You wire the library into your own app and database and own the runtime, since it runs inside your application, not as a separate identity server. The payoff is full code control. The cost is more integration work upfront.

**Authentication methods come from an open plugin ecosystem configured inside your own app.** Because it is a self-hosted library, method coverage depends on the plugins added, and the wiring into application and database is yours.

**SSO/SAML is available through plugins that you implement and host yourself**, consistent with its self-hosted library model. There is no managed enterprise-connection service. The capability exists but the integration and operation are your responsibility.

**Organization and multi-tenant support comes via plugins configured in your own app.** As a self-hosted library, the tenancy model is something you add and run yourself, not a hosted platform feature.

---

## COST

**Free and open source, with commercial use allowed and no per-MAU fee at any volume.**

The cost is operational, not licensed (your own hosting, database and ongoing maintenance), so spend tracks infrastructure and engineering time instead of user count.

---

## ECOSYSTEM & EXIT

**No vendor lock-in.** Users sit in your own database and the library runs inside your application. Because there is no hosted store to leave, migration is a matter of your own infrastructure, with no provider to extract data from.

---

## TRACK RECORD

**A fast-growing open-source project,** not a hosted vendor, so support comes from the community and its ecosystem. You carry operational responsibility for security updates and hosting, which shifts the trust question onto your own team.

---

## WHERE IT STANDS OUT

**Full control with zero per-user fees:** an open-source, self-hosted library that keeps users in your own database and charges nothing per MAU. Ownership of the runtime and the absence of metered pricing are its defining traits.

---

## LIMITATIONS

**You own hosting, security updates,** and maintenance, since it runs inside your own app and database, and as a fast-growing open-source project it is less proven at very high scale than the established hosted vendors. Operational burden is the core tradeoff.

---

Pricing and limits read from [Better Auth's own site](#) . Verified to the edition date above.

# Stytch

## BUILT FOR

Developers building passwordless and fraud-resistant flows who want APIs for magic links, passkeys, OTP, and device fingerprinting over prebuilt UI. Best fit: when auth is a core product surface needing embedded, API-first passwordless and bot/fraud protection, versus drop-in component kits.

---

## CHOOSE THIS IF YOU ARE

Teams wanting API-first, headless passwordless and biometric flows plus built-in fraud and device-fingerprinting.

### BEST WHEN

You're building custom auth UX around passwordless magic links, OTP, passkeys, or WebAuthn and want fraud/device intelligence via API instead of prebuilt screens. Free covers 10k MAU (plus 5 SSO/SCIM connections), then roughly \$0.05/MAU.

### AVOID IF

You want prebuilt drop-in UI components and to avoid building auth screens, you're optimizing purely for lowest cost, or you need enterprise SSO reselling.

## WHAT IT DOES

**Fast for developers via clean SDKs plus pre-built (and headless) UI components.** Docs are strong and integration is quick, though the drop-in UI is less turnkey/polished than Clerk's out-of-the-box components.

**Broad: magic links, email/SMS/WhatsApp OTP, passkeys/WebAuthn, passwords, OAuth, TOTP, and M2M tokens.** Standout protection is built-in device fingerprinting/fraud (10k fingerprints free, then \$0.005 each).

**Yes: SAML and OIDC SSO plus SCIM directory sync in the B2B product.** Five connections are included free, then \$125 per additional connection. Genuinely enterprise-ready for B2B SaaS deals.

**Strong: unlimited Organizations in the B2B product with per-org settings, RBAC, JIT provisioning, and multi-tenant session handling.** Built specifically for B2B SaaS multi-tenancy from the ground up.

## COST

**Free up to 10,000 MAU (both B2C and B2B)**, then ~\$0.05/MAU with volume discounts and no hard caps. The cliff: SSO/SCIM connections are \$125 each beyond the 5 free, and per-MAU cost compounds at consumer scale.

---

Pricing and limits read from [Stytch's own site](#) . Verified to the edition date above.

## ECOSYSTEM & EXIT

**Moderate.** Standards-based OAuth/SAML/OIDC aid portability, but sessions, user records, and the fraud layer are Stytch-specific. Migration APIs/support exist. The per-MAU model can pressure economics once large enterprise contracts land.

## TRACK RECORD

**Well-funded independent company (founded 2020) with strong docs**, SOC 2 / ISO / GDPR compliance, a 99.99% enterprise uptime SLA, and growing adoption. Younger than Auth0 but mature and actively shipping.

## WHERE IT STANDS OUT

**The auth + fraud combo: passwordless-native login plus built-in device fingerprinting and bot protection in one API**, paired with a 10k-MAU free tier and usage-based pay-as-you-go pricing.

## LIMITATIONS

**Per-MAU pricing gets expensive at consumer scale and can misalign once big enterprise contracts sign.** Drop-in UI is less polished than Clerk, and it has fewer legacy turnkey integrations/extensions than Auth0's older ecosystem.



# Build

## BUILT FOR

Skip the auth SaaS: wire authentication yourself with a library (Lucia, Auth.js/NextAuth, SuperTokens self-host, Passport) against your own DB. You own the sessions, tokens, and user table, and every edge case.

---

## CHOOSE THIS IF YOU ARE

Founders who want maximum control and no lock-in: SuperTokens self-hosted, Auth.js for Next, or Lucia primitives to learn the internals.

### BEST WHEN

Compliance demands fully self-hosted identity, or you're deliberately learning auth deeply. All are free/open-source, keep users entirely in your infrastructure, and impose no per-MAU tax. Pick SuperTokens when you want a self-hostable full solution, Auth.js when you just need session handling glued into Next.

### AVOID IF

You're pre-seed; roll-your-own only when self-hosting is a hard requirement, not a preference.

---

## WHAT IT DOES

**A library handles hashing, sessions/JWTs, and OAuth flows** while you store users in your Postgres/MySQL. No hosted service. Auth logic lives in your app and runs on your infra.

**Slower and mostly UI-less.** Libraries give primitives, but polished login screens are on you. Forms, reset flows and verification emails are yours to build. SuperTokens ships some pre-built UI. Most others leave the UX entirely to you.

**Email/password and OAuth (Google/GitHub) come easy.** Magic links, TOTP/2FA, rate-limiting, CSRF, and secure cookies are on you to add correctly. The library helps, but protection depth equals your security diligence.

**The weak spot.** SAML/OIDC enterprise SSO is complex and mostly unsupported out-of-box (SuperTokens/self-host has some). Rolling your own SAML is a security minefield. If you need enterprise SSO soon, DIY is the wrong call.

**No built-in orgs, roles, or invitations.** You design the tenant schema, membership, and RBAC in your own tables. Doable but real work, but the SaaS rivals hand you this, whereas here every multi-tenant feature is custom.

---

## COST

**No per-MAU pricing and no cliff,** which is the exact reason to DIY. You pay only compute/DB you already run. Escapes the brutal auth-SaaS jump where free tiers end and MAU billing balloons. Cost is engineering time, mostly one-time.

---

## ECOSYSTEM & EXIT

**Zero vendor lock-in.** Users and password hashes sit in your DB, portable anywhere. The flip side is you own migrations, hash-scheme upgrades, and breaking-change maintenance on the library yourself.

---

## TRACK RECORD

**Libraries are mature and widely used,** but auth is security-critical and DIY means no audits, no compliance certs, no support line. A subtle session or token bug is a breach. Trust rests entirely on your implementation.

---

## WHERE IT STANDS OUT

**Kills MAU billing, keeps user data in your DB,** and bends to any custom flow. Cheapest and most flexible for MVPs and cost-sensitive products where you control the stack and don't need enterprise features day one.

---

## LIMITATIONS

**Security burden is heavy and unforgiving.** One mistake leaks accounts. No enterprise SSO, orgs, or compliance out-of-box, but UI and edge cases (email deliverability, account recovery) are all yours. Breaks down as security/enterprise needs rise.

---

No vendor page: this route is assembled from the components named above.

# Hosting

Compared: [Vercel](#) · [Render](#) · [Railway](#) · [Cloudflare Pages](#) · [Fly.io](#) · [Netlify](#) · [Build](#)

# At a glance

| TOOL                    | CHOOSE THIS IF YOU ARE                                                                                                                              | STARTING PRICE                                                                                                                                                                                                                                                            | MAIN BENEFIT                                                                                                                                                                                                                                                                                                          |
|-------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Vercel</b>           | <b>Frontend and Next.js teams who want zero-config deploys</b><br>instant preview URLs per PR, and a polished edge/serverless developer experience. | <b>Hobby is \$0 with hard usage caps</b><br>Pro is \$20/user/mo, though independent reporting puts typical Pro bills near \$67/mo once real usage is included; Enterprise starts around \$3,500+/mo. Fluid Compute bills only during code execution, not during I/O waits | <b>The strongest host for Next.js/React, pairing zero-config Git deploys with Fluid Compute, which bills only during code execution rather than during I/O waits, cutting the idle-function spend that weighs on conventional serverless billing.</b>                                                                 |
| <b>Render</b>           | <b>Full-stack founders who want a simple PaaS with fixed</b><br>predictable per-service pricing and a genuinely permanent free tier.                | <b>Flat per - service pricing</b><br>each paid service carries a fixed cost and stays warm, so the bill is predictable and does not swing with traffic                                                                                                                    | <b>The most complete free tier of the group (static sites, 750 web-service instance-hours/month and a free Postgres with no card), combined with flat per-service pricing and warm, always-on paid instances that give a predictable bill.</b>                                                                        |
| <b>Railway</b>          | <b>Solo devs and small teams wanting a slick PaaS that meters per-second</b><br>low-traffic backends stay cheap.                                    | <b>Usage - based</b>                                                                                                                                                                                                                                                      | <b>The strongest modern developer experience for container deploys, paired with usage-based, scale-to-zero economics.</b><br>Services idle down to no compute cost and you pay for consumption, which favors bursty and low-traffic workloads.                                                                        |
| <b>Cloudflare Pages</b> | <b>Teams shipping static sites</b><br>SPAs, or JAMstack frontends plus edge Workers who never want to pay egress.                                   | <b>Free static hosting with unlimited sites</b><br>unlimited bandwidth and 500 builds/month, carrying no commercial-use restriction                                                                                                                                       | <b>The cheapest global static hosting here (unlimited sites and unlimited bandwidth with no egress fees and no commercial-use restriction), served from Cloudflare's global edge, so traffic-driven bandwidth bills essentially do not apply.</b>                                                                     |
| <b>Fly.io</b>           | <b>Teams wanting Docker containers running close to users across many regions</b><br>with persistent volumes and pure pay-as-you-go.                | <b>Pure usage - based with no seat fee</b><br>each running machine is billed per second                                                                                                                                                                                   | <b>The finest-grained global control in this group.</b><br>Docker containers run as micro-VMs placed per region and billed per second with no seat fee, so you tune exactly where compute runs and pay only for the machine time used.                                                                                |
| <b>Netlify</b>          | <b>JAMstack and static-frontend teams who want a rich build-plugin ecosystem plus built-in forms</b><br>identity, and edge functions.               | <b>Flat + credit - based</b>                                                                                                                                                                                                                                              | <b>Framework flexibility plus batteries included: the best pick for non-React stacks (Astro, SvelteKit, Hugo), with built-in forms, identity/auth, and A/B split-testing that replace separate paid tools.</b><br>Its free tier allows commercial use (unlike Vercel), and background functions run up to 15 minutes. |
| <b>Build</b>            | <b>Technically confident founders who want a self-hosted Heroku experience with zero platform markup</b><br>full control.                           | <b>Flat monthly VPS rent</b><br>(~\$4-6/mo entry Hetzner, more for RAM/CPU/disk), not usage-metered                                                                                                                                                                       | <b>Rock-bottom flat cost, no cold starts, no bill spikes, runs literally anything, zero lock-in.</b><br>Best value for a founder with ops skill running an MVP or steady-traffic app who wants predictable spend.                                                                                                     |

Full detail for every tool follows on its own page.

# Vercel

## BUILT FOR

Frontend and full-stack JS/TS teams, especially Next.js users, wanting zero-config Git deploys, preview environments, and edge/serverless functions. Best fit: Next.js and modern frontend apps where developer experience, per-PR previews, and a global CDN matter more than raw infra control.

---

## CHOOSE THIS IF YOU ARE

Frontend and Next.js teams who want zero-config deploys, instant preview URLs per PR, and a polished edge/serverless developer experience.

### BEST WHEN

You're on Next.js (or a modern frontend) and want git-push deploys, per-PR preview environments, and edge functions with a smooth workflow: Hobby is free, Pro is \$20/user/mo including 1TB fast data transfer and 10M edge requests, with overage at \$0.15/GB.

### AVOID IF

You need always-on stateful backends, containers, or databases, or you serve heavy bandwidth where \$0.15/GB overages hurt.

## WHAT IT DOES

**It serves static sites alongside serverless and edge functions running on Fluid Compute,** which bills only during actual code execution instead of while a function sits idle waiting on I/O, the platform's main efficiency lever.

**Static sites plus serverless and edge functions on Fluid Compute.** The compute layer is framework-optimized around Next.js, and functions are billed only during code execution (not during I/O waits), which is the platform's stated way of cutting wasted function spend.

**Largely mitigated on Fluid Compute.** Functions are billed during execution and kept responsive instead of paying an idle penalty. The classic serverless cold-start hit is not the platform's main concern.

**Zero-config for Next.js.** Connect a Git repo and a push triggers a build and deploy with no pipeline to assemble. The tailored path is strongest inside the Next.js/React lane the platform is optimized around.

**Auto-scales front-end and function workloads across a global edge network**, so distribution is handled by the platform, not configured per region. Reach is built into the edge model, not something you provision.

---

## COST

**Hobby is \$0 with hard usage caps.** Pro is \$20/user/mo, though independent reporting puts typical Pro bills near \$67/mo once real usage is included, but Enterprise starts around \$3,500+/mo. Fluid Compute bills only during code execution, not during I/O waits.

**The highest here.** There are no automatic spend limits by default (you must configure spend management manually), and bandwidth past the 1TB allowance costs \$0.15/GB. Documented cases include a \$23,000 bill after a DDoS and \$700-\$1,100 charges from ordinary traffic spikes.

**Hobby is free but ships hard caps (100GB bandwidth, 1M invocations, 4h Active CPU and 360 GB-hr memory per month)**, and hitting any one cap stops the project serving until the next month, not auto-billing overage. Hobby is also restricted to non-commercial use.

---

## ECOSYSTEM & EXIT

**A standard Next.js app ports to other hosts.** The application itself is not locked in. What does not transfer is the tightest Vercel-specific optimizations, so a moved app can lose some of the platform-tuned behavior.

---

## WHERE IT STANDS OUT

**The strongest host for Next.js/React**, pairing zero-config Git deploys with Fluid Compute, which bills only during code execution rather than during I/O waits, cutting the idle-function spend that weighs on conventional serverless billing.

---

## LIMITATIONS

**Hobby bans commercial use and carries hard monthly caps that stop the project serving when hit.** It also holds the biggest bill-shock risk here: no default spend cap, \$0.15/GB over the 1TB allowance, and documented cases like a \$23,000 DDoS bill and \$700-\$1,100 traffic-spike charges.

---

Pricing and limits read from [Vercel's own site](#) . Verified to the edition date above.

# Render

## BUILT FOR

Small-to-mid teams wanting a Heroku-style PaaS for web services, background workers, cron jobs, and managed Postgres without touching Kubernetes. Best fit: full-stack apps needing long-running services and databases with simple config, versus frontend-first or function-only platforms.

---

## CHOOSE THIS IF YOU ARE

Full-stack founders who want a simple PaaS with fixed, predictable per-service pricing and a permanent free tier.

### BEST WHEN

You want web services, managed Postgres, cron jobs, and background workers with pricing you can predict and no surprise metering.

### AVOID IF

You want cheapest metered pay-per-second for low-traffic apps, or you're frontend-only.

---

## WHAT IT DOES

**Pricing is flat per service and paid instances stay warm.** It targets steady always-on multi-service stacks, not bursty workloads.

**Web services, static sites and managed PostgreSQL and Redis**, deployed from containers or native runtimes. The managed-database support alongside long-running web services is what makes it a full-stack host instead of a static or front-end-only one.

**The main free-tier catch.** Free web services spin down after 15 minutes of inactivity and then take 30-60 seconds to cold-start on the next request. Paid instances stay warm and avoid this entirely. The spin-down is specific to the free tier.

**Very easy Git-based deploys with automatic runtime detection**, so a standard app builds and ships without much configuration. It runs from containers or native runtimes, keeping the setup step light for common full-stack stacks.

**Auto-scales services and offers deployment across multiple regions.** Scaling is handled at the service level on always-on instances, fitting steady multi-service stacks better than bursty scale-to-zero patterns.

---

## COST

**Flat per-service pricing: each paid service carries a fixed cost and stays warm,** so the bill is predictable and does not swing with traffic. That trades away usage-based savings for cost certainty, which is the point for steady always-on stacks.

**Low and predictable.** Flat per-service pricing means the bill is set by how many services you run, not by traffic, so a usage spike does not translate into a surprise charge the way a metered model can.

**A genuine full-stack free tier requiring no card:** static sites, web services with 750 instance-hours/month, and a free Postgres. The limits are that free web services cold-start after spinning down and the free Postgres expires 30 days after creation (with a 14-day grace period to upgrade).

---

## ECOSYSTEM & EXIT

**Low lock-in.** Services run as standard containers or native runtimes, so moving to another host is mainly a matter of redeploying the same artifacts elsewhere instead of unwinding platform-specific code.

---

## WHERE IT STANDS OUT

**The most complete free tier of the group (static sites, 750 web-service instance-hours/month and a free Postgres with no card),** combined with flat per-service pricing and warm, always-on paid instances that give a predictable bill.

---

## LIMITATIONS

**The free tier's web services spin down after 15 minutes idle and take 30-60 seconds to cold-start,** and the free Postgres expires 30 days after creation. Avoiding both means moving to paid, warm instances. The free tier is genuine but time- and latency-limited.

---

Pricing and limits read from **Render's own site** . Verified to the edition date above.



# Railway

## BUILT FOR

Indie developers and early startups wanting the fastest path from repo to running services and databases, with usage-based pricing and a clean dashboard. Best fit: quick full-stack deploys and prototypes where you want containers, databases, and env vars managed together with minimal setup.

---

## CHOOSE THIS IF YOU ARE

Solo devs and small teams wanting a slick PaaS that meters per-second so low-traffic backends stay cheap.

## BEST WHEN

You run a web service plus database plus worker at modest traffic and want to pay only for what actually runs.

## AVOID IF

You need a meaningful free tier, you're hosting a static frontend, or you want globally-distributed containers with persistent volumes.

## WHAT IT DOES

**It runs any container plus managed databases and meters consumption per resource.** Idle services stop billing, an economic model that suits bursty or low-traffic workloads over flat always-on ones.

**Any container, plus managed databases,** deployed from Git or Docker. Consumption is metered per resource with scale-to-zero. Services can idle down to no compute cost, the model Railway leans on instead of flat per-service plans.

**A consequence of scale-to-zero:** a service that has idled down cold-starts on its next request. Keeping a service always-on avoids the delay but forgoes the idle-cost savings, but the choice is between cost and first-request latency.

**Among the best modern developer experience in this group.** Deploy from Git or Docker in minutes. The workflow is built around getting a container running with minimal setup, which is central to Railway's positioning.

**Scales compute up and down, including to zero when idle,** within a regional footprint. The strength is elastic, consumption-following compute, not the broad multi-region distribution some peers emphasize.

## COST

**Usage-based.** A permanent Free plan exists at \$0/mo with \$1 of monthly usage credits, alongside a one-time \$5 trial credit (30 days, no card). The entry paid plan is Hobby at \$5/mo with \$5 of monthly credits, and Pro is \$20/mo per workspace with \$20 included. Consumption beyond the included credits is billed as used.

**Usage-based.** The bill scales directly with consumption and a busy period costs more. Scale-to-zero caps idle spend by stopping services with no traffic, but there is no flat ceiling, so cost tracks how much compute you actually draw.

**A permanent Free plan exists:** \$0/mo with \$1 of monthly usage credits, plus a separate one-time \$5 trial credit (30 days, no card required). The ongoing entry paid plan is Hobby at \$5/mo with its \$5 monthly credit allowance.

---

## ECOSYSTEM & EXIT

**Low lock-in.** Workloads are standard containers. The app is portable and migrating away is largely redeploying the same container image on another host without proprietary rework.

---

## WHERE IT STANDS OUT

**The strongest modern developer experience for container deploys**, paired with usage-based, scale-to-zero economics. Services idle down to no compute cost and you pay for consumption, which favors bursty and low-traffic workloads.

---

## LIMITATIONS

**The Free plan's \$1 of monthly credits is tiny.** Real workloads land on the \$5/mo Hobby plan quickly. Being usage-based, bills scale with consumption, so there is no flat ceiling on a busy month.

---

Pricing and limits read from [Railway's own site](#) . Verified to the edition date above.

# Cloudflare Pages

## BUILT FOR

Developers deploying static sites and edge apps who want tight pairing with Workers on Cloudflare's global network, in a JS/TS-friendly workflow. Best fit: JAMstack and edge-first apps that lean on Workers, KV, R2, and D1 at the edge, versus origin-based serverless hosts.

---

## CHOOSE THIS IF YOU ARE

Teams shipping static sites, SPAs, or JAMstack frontends plus edge Workers who never want to pay egress.

### BEST WHEN

You're serving a high-traffic static site or SPA and want unlimited free bandwidth and requests at any scale, with Workers for edge logic. The free tier is uniquely uncapped on bandwidth (500 builds/mo), so a viral frontend costs nothing in egress.

### AVOID IF

You need long-running stateful backends, containers, or a traditional Node server with persistent processes, or deep Next.js SSR features.

## WHAT IT DOES

**It runs on Cloudflare's global edge network and is distinguished less by features than by an unusually generous free tier:** unlimited sites and unlimited bandwidth.

**Static assets plus edge compute via Workers,** with R2 object storage and KV available for state. Compute runs within the edge runtime, not as long-lived containers, which shapes both what it can run and how portable that code is.

**Effectively none for static assets,** which are served straight from the edge. Workers run on an always-warm edge runtime instead of spinning up per request. The edge-compute path avoids the cold-start penalty seen on scale-to-zero container hosts.

**Easy for static and Jamstack projects.** Connect a Git repo and deploys run against the edge. The straightforward path applies to static builds; adding Workers edge compute brings the edge runtime's own constraints into the setup.

**The strongest global reach in this group,** served from Cloudflare's global edge network. Static assets and Workers run from that edge footprint, which is the platform's core distribution advantage over container-based hosts.

---

## COST

**Free static hosting with unlimited sites,** unlimited bandwidth and 500 builds/month, carrying no commercial-use restriction. Beyond the free allowances, Workers edge compute is metered. The paid layer sits on compute; bandwidth and static hosting stay free.

**Low.** Bandwidth is unlimited with no egress fees on the free tier, so the usual source of hosting bill-shock (traffic-driven data-transfer charges) does not apply. Metered cost only appears if you push Workers compute past its free allowances.

**The most generous static free tier here:** unlimited sites, unlimited bandwidth and 500 builds/month, with no commercial-use restriction. The main boundary is on compute (Workers usage is metered once past its free allowances), not on hosting or bandwidth.

---

## ECOSYSTEM & EXIT

**Static sites and standard code are portable and move easily.** The exception is Workers-specific code written for the edge runtime, which is less portable and would need rework to run on a conventional host.

---

## WHERE IT STANDS OUT

**The cheapest global static hosting here (unlimited sites and unlimited bandwidth with no egress fees and no commercial-use restriction),** served from Cloudflare's global edge, so traffic-driven bandwidth bills essentially do not apply.

---

## LIMITATIONS

**The edge model constrains compute.**

Workloads run within the Workers edge runtime rather than as full long-lived containers, and Workers-specific code is less portable to a conventional host, the price of the platform's cheap, wide edge reach.

---

Pricing and limits read from [Cloudflare Pages's own site](#) . Verified to the edition date above.

# Fly.io

## BUILT FOR

Engineers who want to run full Docker containers and VMs close to users in multiple regions, comfortable managing more infrastructure themselves. Best fit: latency-sensitive or stateful apps needing multi-region containers and persistent volumes, versus higher-level PaaS abstractions that hide the machine.

---

## CHOOSE THIS IF YOU ARE

Teams wanting Docker containers running close to users across many regions, with persistent volumes and pure pay-as-you-go.

### BEST WHEN

You need real containers deployed to multiple global regions with persistent storage (even a database or stateful service), billed per-second with no base plan fee (a shared-cpu-1x/256MB runs ~\$2/mo at 24/7. No free tier, just a short trial).

### AVOID IF

You want managed hand-holding and a free tier, a static frontend, or you'd rather not think about regions and volumes at all.

---

## WHAT IT DOES

**Billing is per second per running machine with no seat fee**, so cost tracks actual machine runtime, not a fixed monthly plan.

**Any Docker container, run as a micro-VM (a Machine)**, plus managed Fly Postgres. Because the unit is a standard container under full region control, it runs largely whatever you can package in Docker, not a fixed set of supported runtimes.

**You control the tradeoff per machine.** Keep a machine warm for instant response, or let it auto-stop to save cost and accept a start-up delay on the next request. The behavior is a per-machine configuration choice, not a platform default.

**The most involved of this group.** You define machines and regions explicitly via a `fly.toml` configuration instead of getting a zero-config Git flow, which is the cost of the per-region, per-machine control the platform offers.

### **Purpose-built for global, multi-region**

**deployment.** You place machines per region and run compute close to users. Reach is explicit and under your control. You decide the regions instead of relying on an automatic edge network.

---

### **COST**

**Pure usage-based with no seat fee.** Each running machine is billed per second. A minimal always-on machine (shared-cpu-1x, 256MB) runs about \$1.94/mo, while typical production usage lands around \$13-20/mo. Cost scales with how much machine time you actually run.

**Per-second billing means cost rises with load and running-machine count,** and there is no spend cap by default. Because there is no fixed ceiling, a heavier workload (more machines or more runtime) adds up directly against the per-second rate.

**No free tier for new users.** New accounts get a \$5 one-time trial credit, after which a card is required to continue. Unlike Render, there is no ongoing free option. The trial is a short evaluation window, not a standing plan.

---

### **ECOSYSTEM & EXIT**

**Low app lock-in.** The deployed unit is a standard Docker container, but the application moves to any container host. What is Fly-specific is the machine and region orchestration, not the app itself.

---

### **WHERE IT STANDS OUT**

#### **The finest-grained global control in this group.**

Docker containers run as micro-VMs placed per region and billed per second with no seat fee, so you tune exactly where compute runs and pay only for the machine time used.

---

### **LIMITATIONS**

**No free tier (only a \$5 one-time trial credit before a card is required),** and it has the steepest setup here, defining machines and regions via fly.toml. Per-second billing has no default cap. Cost adds up under sustained load.

---

Pricing and limits read from [Fly.io's own site](#) . Verified to the edition date above.

# Netlify

## BUILT FOR

Frontend teams and agencies building JAMstack sites who want Git-based deploys, deploy previews, and serverless functions with minimal config. Best fit: static and hybrid marketing or content sites where build pipeline and CDN simplicity lead, versus app-platform hosts running long-lived servers.

---

## CHOOSE THIS IF YOU ARE

JAMstack and static-frontend teams who want a rich build-plugin ecosystem plus built-in forms, identity, and edge functions.

### BEST WHEN

You're deploying a static site or SPA and value Netlify's mature plugin marketplace, form handling, and add-ons: free covers 300 credits/month (bandwidth draws them down at 20 credits/GB, so ~15GB at most), Pro is \$20/mo flat with unlimited seats (Netlify removed per-seat pricing in April 2026).

### AVOID IF

You're on Next.js and want first-class SSR/ISR, you serve heavy bandwidth and want free egress, or you need real always-on backends, containers, or databases.

## WHAT IT DOES

**Batteries-included: built-in forms, identity/auth**, and A/B split-testing that replace separate tools. No traditional backend or provisioned database.

**Static sites + serverless functions + edge functions + background functions (up to 15 minutes) + scheduled/cron functions + Netlify Blobs (key-value storage)**. Next.js runs via the @netlify/next adapter. You can build APIs, process background jobs, and handle form submissions, but there's no persistent always-on backend or provisioned relational database.



**Slower than Vercel.** Serverless cold starts run ~3s+ (versus Vercel's ~1s), edge function cold starts ~28ms (vs ~12ms), and ~90ms average TTFB (vs ~70ms). It also has 16+ edge locations versus Vercel's 100+, so distant users can feel slightly more latency. Fine for static/Jamstack; noticeable for SSR-heavy apps.

**Excellent and framework-agnostic: Git-based deploys,** preview URLs, a mature plugin ecosystem, and vercel.json-equivalent config via netlify.toml. Its built-in forms, identity/auth, and split-testing remove entire categories of extra setup. The standout for Astro/Hugo/SvelteKit and any multi-framework or static workflow.

**Auto-scales on its CDN across 16+ edge locations,** solid but fewer points of presence than Vercel (100+) or Cloudflare (300+), so users far from a POP can feel slightly more latency. Perfectly fine for most static/Jamstack traffic. Less ideal if global edge latency is a hard requirement.

---

## COST

**Flat + credit-based.** Free (300 credits/mo) → Pro \$20/mo flat with unlimited seats → Enterprise. Netlify moved to credit-based pricing and, on 2026-04-14, dropped per-seat charges. Pro includes 3,000 credits/month. Surface price is close to Vercel's, but the credit model changes the real cost. Heavy builds and function usage draw down credits, so model your actual usage, not trusting the sticker price.

**Moderate: the credit model can surprise you on heavy build minutes or function usage,** but it's generally more predictable than Vercel's bandwidth/invoke overages. The main thing to watch is drawing down the included monthly credit allowance (300 on Free, 3,000 on Pro) faster than expected on frequent deploys or busy functions.

**Free tier: 300 credits/month, custom domains, and HTTPS.** Crucially, it allows commercial use, unlike Vercel's Hobby tier. Everything draws from the same credit pool (bandwidth at 20 credits/GB, so ~15GB if spent entirely on bandwidth), which can throttle teams that deploy frequently or run busy functions.

---

## ECOSYSTEM & EXIT

**Lower lock-in than Vercel.** It's framework-agnostic; standard apps move cleanly to another host or self-hosting. You'll hit some friction only if you depend on Netlify-specific features (Blobs storage, identity/auth, split-testing), which have no drop-in equivalent elsewhere.

---

## WHERE IT STANDS OUT

**Framework flexibility plus batteries included:** the best pick for non-React stacks (Astro, SvelteKit, Hugo), with built-in forms, identity/auth, and A/B split-testing that replace separate paid tools. Its free tier allows commercial use (unlike Vercel), and background functions run up to 15 minutes.



## LIMITATIONS

**Slower serverless cold starts (~3s) and far fewer edge locations (16+) than Vercel.** The credit-based model (300 credits/month on Free) can confuse and throttle heavy users. No persistent backend or provisioned database, and SSR support is clunkier, so the latest Next.js features tend to land on Vercel first.

---

Pricing and limits read from [Netlify's own site](#) . Verified to the edition date above.

# Build

## BUILT FOR

Own the metal: rent a VPS (Hetzner/DigitalOcean) and run your app with Docker, Coolify, or Caddy instead of a managed PaaS. You get a cheap, unopinionated server and full root, plus all the sysadmin that implies. Founders and devs comfortable with the terminal who want low, predictable cost and full control. Not built for those who want git-push-and-forget. Here you own uptime, patches, and firewall, not a platform.

---

## CHOOSE THIS IF YOU ARE

Technically confident founders who want a self-hosted Heroku experience with zero platform markup and full control.

### BEST WHEN

You want to run Coolify (open-source PaaS) or plain Docker Compose on a \$5-20/mo Hetzner/DigitalOcean box: git-push deploys, databases, and SSL you fully own, at a fraction of managed pricing for predictable workloads.

### AVOID IF

You're pre-seed and every hour on uptime, TLS renewal, scaling, and security patching is an hour not spent on product.

## WHAT IT DOES

### Coolify/Dokku add a Heroku-like layer on top.

Bare Docker+Caddy is even more hands-on. No platform abstracting the box away.

**Anything that runs on Linux: any language,** long-lived servers, background workers, cron, databases, websockets, with no runtime restrictions or function timeouts. More flexible than serverless PaaS, which constrains what you can run.

**None.** The server runs continuously so every request hits a warm process, with no serverless cold-start latency. What it costs you is paying 24/7 whether traffic comes or not, and scaling manually, not to zero.

**Hardest factor.** Bare VPS means manual deploys, SSH, and config unless you add Coolify/Dokku for push-to-deploy. Even then it's more work than a managed PaaS. Great once wired, painful to get there and to keep running.

**Manual and single-region by default.** You resize the box (vertical) or hand-build load balancers and multi-region (horizontal). No auto-scale, no global edge/CDN unless you assemble it. This is where DIY clearly breaks at scale.

---

## COST

**Flat monthly VPS rent (~\$4-6/mo entry Hetzner, more for RAM/CPU/disk),** not usage-metered. No per-request, bandwidth (within cap), or build-minute billing. Cost is fixed and low. Your TIME is the real spend, mostly upfront.

**Almost none, which is the draw.** Flat rent means no surprise autoscale bill. Bandwidth is generous and capped. Hidden cost is your labor and outage risk, plus paying for idle capacity you provisioned but don't use.

**No free tier: you pay from day one (~\$4/mo).** Cheaper than most PaaS paid plans, but there's no \$0 hobby lane. Catch: the low price hides the setup/maintenance time, which is the actual cost of going DIY.

---

## ECOSYSTEM & EXIT

**Very portable: it's just Docker/Linux,** movable to any provider with a container and a DNS change. No proprietary platform to escape. Owning standard tooling is the strongest long-term argument for self-managing.

---

## WHERE IT STANDS OUT

**Rock-bottom flat cost, no cold starts,** no bill spikes, runs literally anything, zero lock-in. Best value for a founder with ops skill running an MVP or steady-traffic app who wants predictable spend.

---

## LIMITATIONS

**You are the ops team.** Security patches, backups, monitoring, and 3am outages are yours. No auto-scaling, edge, or managed HA without heavy work, but unpolished unless you invest. Ops load grows with users, which is DIY's ceiling.

---

No vendor page: this route is assembled from the components named above.

# Product Analytics & Experimentation

Compared: PostHog · Amplitude · Mixpanel · Google Analytics · Heap · Statsig · Build

# At a glance

| TOOL                    | CHOOSE THIS IF YOU ARE                                                                                                                                                 | STARTING PRICE                                                                                                                                                                                          | MAIN BENEFIT                                                                                                                                                                                                                |
|-------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>PostHog</b>          | <b>Founders who want product analytics</b><br>session replay, feature flags, experiments, and surveys unified under one usage-based bill (and optionally self-hosted). | <b>Usage-based with a very generous free tier</b><br>1M events, 5K web + 2.5K mobile recordings, 1M feature-flag requests (experiments billed with flags), no card                                      | <b>The only major tool that is open-source and self-hostable while bundling analytics + replay + flags + experiments together under one unusually generous free plan.</b>                                                   |
| <b>Amplitude</b>        | <b>Product teams that need deep behavioral analytics</b><br>(detailed cohorts, funnels, and retention) at enterprise scale.                                            | <b>Free plan forever: 2M events/month with 10K monthly session replays</b><br>no card required                                                                                                          | <b>Deep analytics and cohorting with enterprise governance controls, plus a free tier that bundles session replay and basic experimentation.</b>                                                                            |
| <b>Mixpanel</b>         | <b>Teams that want fast, interactive event analytics with predictable</b><br>event-based billing.                                                                      | <b>Free plan: 1M events/mo</b><br>no card, including 10K session replays/mo and core reports; Growth is free for the first 1M events then ~\$0.00028/event, with unlimited seats                        | <b>Fast, self-serve analytics UX paired with a large free tier and simple, transparent per-event pricing.</b>                                                                                                               |
| <b>Google Analytics</b> | <b>Founders tracking top-of-funnel web traffic</b><br>acquisition channels, and marketing/SEO performance, for free.                                                   | <b>Free for standard GA4</b><br>(subject to sampling and quota limits)                                                                                                                                  | <b>Free, widely used, and strong for marketing and acquisition analytics and Google Ads attribution, with a large community and documentation.</b>                                                                          |
| <b>Heap</b>             | <b>Teams that don't want to manually instrument events and prefer to autocapture everything</b><br>then define metrics retroactively.                                  | <b>Free plan up to ~10K sessions/mo with core analytics and 6-month history</b><br>Growth, Pro, and Premier tiers are custom-quoted (Growth reportedly ~\$3,600/yr)                                     | <b>Retroactive autocapture means no lost data from un-instrumented events, plus integration with Contentsquare's experience-analytics suite.</b>                                                                            |
| <b>Statsig</b>          | <b>Teams whose core need is experimentation and feature flags with a rigorous stats engine</b><br>with analytics riding alongside.                                     | <b>Large free tier: 2M events/month with unlimited feature-flag and config checks</b><br>50K session replays and core experimentation; advanced stats (CUPED, sequential testing) sit on Pro/Enterprise | <b>The most generous free experimentation offering paired with a statistically rigorous engine, and it scales to very high volume more cheaply than most rivals.</b>                                                        |
| <b>Build</b>            | <b>Privacy-conscious or data-sovereignty-driven teams who want to own every analytics event with no per-event fees</b>                                                 | <b>Software is free/open-source</b>                                                                                                                                                                     | <b>Full data ownership, GDPR-friendly, no event-volume billing, and PostHog-OSS gives analytics+replay+flags in one self-hosted stack.</b><br><br>Best for privacy-driven or high-volume products willing to run the infra. |

Full detail for every tool follows on its own page.

# PostHog

## BUILT FOR

Product and engineering teams wanting analytics, session replay, feature flags, and experiments in one open-source, self-hostable platform. Best fit: startups wanting an all-in-one product suite they can self-host and keep data in, versus buying separate analytics and feature-flag tools.

---

## CHOOSE THIS IF YOU ARE

Founders who want product analytics, session replay, feature flags, experiments, and surveys unified under one usage-based bill (and optionally self-hosted).

## BEST WHEN

You want one tool covering funnels, retention, replays, and A/B tests instead of stitching several. The free tier is a fat 1M events/month with no base fee, and everything's in one dashboard.

## WHAT IT DOES

**Autocaptures clicks and pageviews via a JS snippet**, plus manual/custom events; supports server-side ingestion and reverse-proxy setups for first-party data.

**Full funnels, retention, user paths**, trends, and static/dynamic behavioral cohorts are included on every plan, including the free tier, with no gating on core analysis.

**Native web and mobile (iOS/Android/React Native/Flutter) replay is bundled**; 5K web + 2.5K mobile recordings free per month, then usage-priced, linked to analytics events.

**Feature flags, multivariate flags**, and A/B/n experiments with a built-in stats engine are included free (1M flag requests/mo). Experiment cost is billed via flag request volume.

**JavaScript/Web SDK plus native iOS**, Android, React Native, and Flutter SDKs, and many server SDKs (Python, Node, Go, Ruby, PHP, etc.).

---

## COST

### Usage-based with a very generous free tier:

1M events, 5K web + 2.5K mobile recordings, 1M feature-flag requests (experiments billed with flags), no card. Self-host is MIT-licensed and free.

---

Pricing and limits read from [PostHog's own site](#) . Verified to the edition date above.

---

## ECOSYSTEM & EXIT

### CDP-style pipelines with 60+

**sources/destinations**, warehouse sync (BigQuery, Snowflake, Postgres), Segment import, Zapier, and reverse-ETL between products.

---

---

## WHERE IT STANDS OUT

The only major tool that is open-source and self-hostable while bundling analytics + replay + flags + experiments together under one unusually generous free plan.

---

---

## LIMITATIONS

Breadth over depth: individual modules (replay, warehouse) are less mature than **specialists**, and 4x event-overage math plus self-host ops can surprise growing teams.

# Amplitude

## BUILT FOR

Product managers and growth teams at scaling companies needing deep behavioral analytics, funnels, retention, and cohorting across web and mobile. Best fit: data-driven product orgs doing serious funnel and retention analysis where analytical depth and governance outweigh lighter, faster tools.

---

## CHOOSE THIS IF YOU ARE

Product teams that need deep behavioral analytics (detailed cohorts, funnels, and retention) at enterprise scale.

## BEST WHEN

Understanding user behavior is central and you want extensive cohorting/funnel/retention analysis, now with feature flags and session replay bundled. The free plan gives 2M events/month (10K session replays included), Plus starts at \$0 usage-based with the first 2M events free, and pricing is event-based with unlimited seats.

## AVOID IF

High event volumes could make usage-based bills climb, or you want an all-in-one cheap tool with replay+flags.

## WHAT IT DOES

**Supports autocapture / no-code event streaming plus SDK-based custom events,** with a data-governance layer for taxonomy, naming, and schema control.

**Deep funnels, retention, user journeys,** and behavioral cohorts, its historical core strength and a common reason enterprises adopt it.

**Native Session Replay is included on all plans (Free tier 10,000 sessions/mo, Growth 20,000/mo, Enterprise 50,000/mo),** tightly tied to captured analytics events for qualitative context.

**Basic feature flags and web experimentation are now included on all plans,** including free. Advanced, server-side experimentation is a paid add-on module.



**Browser JavaScript plus native iOS, Android, React Native, Flutter, and Unity SDKs, and multiple server-side SDKs.**

---

Pricing and limits read from [Amplitude's own site](#) . Verified to the edition date above.

---

## COST

**Free plan forever: 2M events/month with 10K monthly session replays**, no card required. Plus starts at \$0 and is usage-based (first 2M events/month free, then event-based pricing up to 70M events). Growth and Enterprise are custom-quoted, also event-based.

---

## ECOSYSTEM & EXIT

**100+ integrations, Segment, warehouse connectors (Snowflake, BigQuery, Redshift), reverse-ETL, and an included customer data platform (CDP).**

---

## WHERE IT STANDS OUT

**Deep analytics and cohorting with enterprise governance controls**, plus a free tier that bundles session replay and basic experimentation.

---

## LIMITATIONS

**Event-based pricing can still get expensive at volume.** Advanced experimentation and higher volumes get pricey. Heavier taxonomy/setup than autocapture-first tools for small teams.

# Mixpanel

## BUILT FOR

Product and marketing teams wanting event-based analytics with fast, self-serve reports, funnels, and dashboards at a low learning curve.

---

## CHOOSE THIS IF YOU ARE

Teams that want fast, interactive event analytics with predictable, event-based billing.

## BEST WHEN

You value forecastable costs and snappy self-serve reports; since switching fully to event-based pricing (Feb 2025), bills scale cleanly with volume: free covers 1M events, then Growth is usage-based at about \$0.28 per 1,000 events above the free 1M (e.g. ~\$140/mo at 1.5M events), with 20K session replays included.

## WHAT IT DOES

**Added Autocapture in Feb 2025**, available on all plans; it captures pageviews, clicks, and form fills with no code, alongside traditional SDK-defined custom events.

**Strong funnels, retention, and Flows plus behavioral cohorts**, though cohorts and some advanced properties are gated off the free plan and require Growth+.

**Native web and mobile session replay is**

**bundled**, with 10K replays per month even on the free plan, linked to the underlying event stream.

**No native feature flags or experiment delivery.**

Offers A/B report analysis but relies on warehouse/integrations or third-party flag tools to run tests.

**JavaScript/Web SDK plus native iOS, Android, React Native, and Flutter SDKs**, and server-side SDKs (Python, Node, etc.).

---

## COST

**Free plan: 1M events/mo, no card**, including 10K session replays/mo and core reports. Growth is free for the first 1M events then ~\$0.00028/event, with unlimited seats.

---

Pricing and limits read from [Mixpanel's own site](#) . Verified to the edition date above.

---

## ECOSYSTEM & EXIT

**Segment, mParticle, warehouse connectors (BigQuery, Snowflake)**, Slack, and 100+ integrations for ingestion and downstream syncing.

---

## WHERE IT STANDS OUT

**Fast, self-serve analytics UX paired with a large free tier and simple**, transparent per-event pricing.

---

## LIMITATIONS

**Feature flags and experimentation exist but sit high up**: built into Enterprise, a paid add-on on Growth. Cohorts and custom properties are gated off the free plan, and event-based pricing can spike with high-volume tracking.

# Google Analytics

## BUILT FOR

Marketers and site owners of any size tracking web and app traffic, acquisition channels, and conversions for free. / Best fit: marketing-driven site analytics and attribution, versus product event tools; a common choice when budget is zero and reach matters.

---

## CHOOSE THIS IF YOU ARE

Founders tracking top-of-funnel web traffic, acquisition channels, and marketing/SEO performance, for free.

## BEST WHEN

You need to know where visitors come from, which campaigns convert, and how your marketing site and Google Ads perform. GA4 is free and integrates natively with the Google Ads/marketing stack for attribution. Choose it over the product-analytics tools when the question is where your traffic and acquisition come from, not how users behave inside your app.

## AVOID IF

You need real product analytics (user-level event streams, precise funnels, cohorts, and retention), because GA4 samples data and isn't built for in-app behavioral analysis.

## WHAT IT DOES

### "Enhanced measurement" auto-tracks

pageviews, scrolls, outbound clicks, site search, and file downloads. Deeper events require gtag.js or Google Tag Manager configuration.

### Funnel, path, cohort, and retention

Explorations exist but are less flexible and product-focused than dedicated tools, and data sampling kicks in at scale.

**No native session replay.** GA4 does not record user sessions. Teams pair it with dedicated tools (Microsoft Clarity, Hotjar, FullStory) for qualitative playback.

**No native A/B testing since Google Optimize sunset on Sep 30, 2023.** It requires third-party tools (Optimizely, VWO) or Firebase for app-side experiments.

## **gtag.js and Google Tag Manager for web.**

Mobile is captured via Firebase SDKs (iOS/Android) that feed data into GA4.

---

Pricing and limits read from [Google Analytics's own site](#) . Verified to the edition date above.

---

### **COST**

**Free for standard GA4 (subject to sampling and quota limits).** The enterprise GA360 tier is custom six-figure pricing. No product-team-specific pricing exists.

---

### **ECOSYSTEM & EXIT**

**Deep Google ecosystem ties (Google Ads, Search Console, and free BigQuery export),** plus a wide range of third-party connectors.

---

### **WHERE IT STANDS OUT**

**Free, widely used, and strong for marketing and acquisition analytics and Google Ads attribution,** with a large community and documentation.

---

### **LIMITATIONS**

**Not a true product-analytics tool:** clunky behavioral analysis, data sampling, no replay, no native experimentation, and a notably steep GA4 learning curve.

# Heap

## BUILT FOR

Product and growth teams wanting autocapture of every interaction so events can be analyzed retroactively without pre-instrumenting the app. Best fit: teams that don't want to define events upfront and value retroactive analysis, versus manual event tracking that misses what you forgot to log.

---

## CHOOSE THIS IF YOU ARE

Teams that don't want to manually instrument events and prefer to autocapture everything, then define metrics retroactively.

### BEST WHEN

You're moving fast and don't want to pre-plan a tracking plan; Heap autocaptures every click, pageview, and interaction so you can define funnels and events after the fact, analyzing behavior that happened before you thought to track it. Choose it when retroactive autocapture and never missing an event is the priority.

### AVOID IF

You want that same autocapture with transparent, cheaper, all-in-one pricing, or you prefer predictable event-based billing and mature reporting; Heap's differentiation is thinner now that rivals autocapture too.

## WHAT IT DOES

**Its defining feature: retroactive autocapture of every click**, pageview, and form interaction, letting you define events after the fact with no code changes or replays lost.

**Full funnels, retention, paths**, and segments built on retroactively captured data, plus strong journey analysis across the full interaction history.

**Native session replay is available**, typically as an add-on rather than bundled by default, alongside heatmaps and Contentsquare experience analytics.

**No native feature flags or experiments.** It pairs with Contentsquare experience analytics or third-party experimentation tools to run and analyze tests.

**Web JavaScript SDK plus native iOS, Android, and React Native SDKs** for autocapture across surfaces.

---

Pricing and limits read from [Heap's own site](#) . Verified to the edition date above.

---

## COST

**Free plan up to ~10K sessions/mo with core analytics and 6-month history**, but Growth, Pro, and Premier tiers are custom-quoted (Growth reportedly ~\$3,600/yr). Pricing is opaque.

---

## ECOSYSTEM & EXIT

**Warehouse connectors (Snowflake, BigQuery, Redshift)**, Segment, Salesforce, HubSpot, and native Contentsquare experience-data integration.

---

## WHERE IT STANDS OUT

**Retroactive autocapture means no lost data from un-instrumented events**, plus integration with Contentsquare's experience-analytics suite.

---

## LIMITATIONS

**Autocapture creates data noise and governance overhead.** Session replay is an add-on. Pricing is opaque and enterprise-leaning; no native experimentation or flags.

# Statsig

## BUILT FOR

Engineering and product teams running feature flags, experiments, and A/B tests at scale, with a warehouse-native deployment option. Best fit: experimentation-heavy orgs needing a rigorous stats engine and gradual rollouts, versus analytics-first tools that treat testing as an add-on.

---

## CHOOSE THIS IF YOU ARE

Teams whose core need is experimentation and feature flags with a rigorous stats engine, with analytics riding alongside.

## BEST WHEN

A/B testing and controlled rollouts are central and you want a warehouse-grade experimentation platform with product analytics attached. The free tier is large: 2M events/month, unlimited flag and config checks, 50K session replays and unlimited seats; Pro (\$150/mo) includes 5M events with \$0.05 per additional 1k.

---

## WHAT IT DOES

**Autocapture for web analytics plus SDK-based event logging**, with automatic exposure logging tied to every flag check and experiment assignment.

**Product analytics with funnels**, retention, and metric explorers, tightly coupled to experiment metrics so tests and analysis share the same definitions.

**Native session replay is included and integrated with experiments and analytics**, letting you tie qualitative playback back to test exposures.

**Its core strength: Unlimited feature gates and core experimentation (multi-variate, Bayesian and frequentist) available even on the free tier**; advanced statistical techniques like CUPED, sequential testing and holdouts sit on Pro and Enterprise.



**30+ SDKs (JavaScript/Web, iOS, Android, React Native, Flutter)**, plus many server-side SDKs (Node, Python, Go, Java, etc.).

---

Pricing and limits read from **Statsig's own site** . Verified to the edition date above.

---

## COST

**Large free tier: 2M events/month with unlimited feature-flag and config checks**, 50K session replays and core experimentation; advanced stats (CUPED, sequential testing) sit on Pro/Enterprise. Pro starts \$150/mo with 5M events included, then \$0.05 per additional 1K.

---

## ECOSYSTEM & EXIT

**Warehouse-native mode (Snowflake, BigQuery, Databricks)**, Segment, Slack, and analytics/CDP connectors for both ingestion and metric sourcing.

---

## WHERE IT STANDS OUT

**The most generous free experimentation offering paired with a statistically rigorous engine**, and it scales to very high volume more cheaply than most rivals.

---

## LIMITATIONS

**Analytics and replay are less mature than dedicated specialists**. Ownership is in flux. OpenAI acquired Statsig (Sep 2025) and the brand/platform moved to Amplitude (2026).

# Build

## BUILT FOR

Own your data, dodge event-volume billing. Privacy- and cost-conscious founders who want data on their own servers and are fine doing ops. Great for MVPs and EU/GDPR needs. Not for teams wanting a zero-setup hosted console and instant dashboards.

---

## CHOOSE THIS IF YOU ARE

Privacy-conscious or data-sovereignty-driven teams who want to own every analytics event with no per-event fees.

### BEST WHEN

GDPR/data-ownership rules or principle demand analytics on infrastructure you control. Choose this when ownership and privacy outweigh convenience.

### AVOID IF

The cloud free tiers more than cover early-stage needs with zero ops burden.

## WHAT IT DOES

**PostHog-OSS gives near-full product analytics on your infra;** Umami/Plausible give lightweight web stats. The rawest path is an events table plus SQL and a Metabase dashboard you build.

**PostHog-OSS keeps autocapture (clicks/pageviews auto-tracked) plus manual events via SDK.** Plausible/Umami are pageview-focused with light custom events. Pure DB approach = fully manual instrumentation you code for every event.

**Self-hosted PostHog delivers funnels, retention, and cohorts like the cloud version.** Umami/Plausible don't. They're traffic stats, not product analytics. Roll-your-own means writing every funnel/cohort as SQL yourself.

**Only PostHog-OSS offers it self-hosted,** and replay is storage- and CPU-heavy. It can strain a modest VPS. Umami/Plausible have none; a DIY events table can't replay sessions. Budget real infra if you want this at scale.

**PostHog-OSS includes feature flags and A/B experiments self-hosted,** a genuine edge. Umami/Plausible and the raw-SQL route have none. You'd bolt on a separate flag library or build gating logic yourself.

**PostHog ships web + mobile (iOS/Android/RN) SDKs that also point at self-hosted.** Plausible/Umami are web-first with thin mobile support. DIY logging works anywhere but you write and maintain each client integration.

---

## COST

**Software is free/open-source.** You pay only the VPS it runs on (flat, low). No per-event pricing, the whole point versus metered SaaS that bills as you grow. Cost is setup/maintenance time, one-time-ish, cheaper long-term at high volume.

---

---

No vendor page: this route is assembled from the components named above.

---

## ECOSYSTEM & EXIT

**Fewer than hosted.** Self-hosted PostHog has many but some cloud integrations/apps lag. Umami/Plausible are minimal. Raw-DB approach connects to anything via SQL but every pipe (warehouse, CRM) is yours to build.

---

---

## WHERE IT STANDS OUT

**Full data ownership, GDPR-friendly,** no event-volume billing, and PostHog-OSS gives analytics+replay+flags in one self-hosted stack. Best for privacy-driven or high-volume products willing to run the infra.

---

---

## LIMITATIONS

**Ops-heavy: PostHog self-hosting is resource-hungry and can be fragile at scale (its own team nudges you to cloud);** Umami/Plausible are shallow. Raw SQL means building everything. Breaks when data volume or feature depth grows.

# Incorporation & Formation

Compared: Stripe Atlas · Clerky · Firstbase · doola · Northwest Registered Agent · Carta · Build

# At a glance

| TOOL                              | CHOOSE THIS IF YOU ARE                                                                                                                     | STARTING PRICE                                                                                                                                                                                                                                                                                | MAIN BENEFIT                                                                                                                                                                                                                                                                                                                                      |
|-----------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Stripe Atlas</b>               | US founders on the venture track who want a Delaware C-corp a bank in one weekend.                                                         | A single \$500 one-time fee covers formation EIN and the first-year registered agent                                                                                                                                                                                                          | The breadth of the bundle: formation, EIN, a US bank account and Stripe payments delivered together within about two business days under one \$500 fee, with \$2,500 in Stripe credits and \$50k+ in partner perks.<br>No other provider here folds banking and payments into the base.                                                           |
| <b>Clerky</b>                     | Founders who will run their own SAFEs option grants, and priced rounds and want lawyer-trusted paperwork.                                  | Company setup starts at \$427 with post-incorporation setup at \$299 and the Lifetime Package at \$819                                                                                                                                                                                        | It stands apart on document depth, founder stock with vesting, bylaws, IP assignment, board consents, and unlimited SAFEs, convertible notes and equity grants under the \$819 Lifetime Package.<br>For companies expecting to raise, the equity paperwork is more complete than the other providers offer.                                       |
| <b>Firstbase</b>                  | Non-US founders who need a US Delaware or Wyoming entity EIN, and address without a US SSN.                                                | Firstbase Start is \$399 for formation recurring costs add separately: Agent Autopilot \$299/year, Mailroom from ~\$35/month, tax filing about \$1,799/year (or ~\$899/year for a non-resident single-member LLC), and Firstbase One at \$2,388/year bundling bookkeeping, tax and compliance | Combining formation with ongoing compliance, bookkeeping and tax in one platform, with Firstbase One at \$2,388/year consolidating those services.<br>It suits founders who want the post-formation operating stack from the same provider that filed the entity, and it's tuned for non-resident founders.                                       |
| <b>doola</b>                      | Non-resident solo founders and ecommerce sellers who want formation plus ongoing compliance handled as a yearly service                    | Plans run from Starter at \$297/year to Tax and Compliance at \$1,999/year and Business-in-a-Box at \$2,999/year (\$329/month)                                                                                                                                                                | Its main advantage is forming a US LLC, C-Corp or DAO LLC without a US SSN, aimed at non-resident founders, with tax and compliance attached to the plan.<br>The DAO LLC option and the no-SSN process set it apart from the more US-resident-oriented providers in this group.                                                                   |
| <b>Northwest Registered Agent</b> | Bootstrapped US founders who want the cheapest most private LLC and nothing they won't use.                                                | The registered agent business address, mail scanning and privacy come to \$125/year with no hidden or per-document fees; once you hold five or more registered-agent services with Northwest, the rate drops to \$100/year each                                                               | Its strength is the registered-agent and privacy service at \$125/year with no hidden or per-document fees, dropping to \$100/year each once you hold five or more registered-agent services.<br>For founders whose priority is agent coverage and a private business address rather than a formation bundle, it is a focused, lower-cost option. |
| <b>Carta</b>                      | Founders whose day-one priority is clean equity cap table, 409A readiness, and electronic share issuance for investors and an option pool. | Carta Launch is free for startups with under 25 stakeholders and under \$1M raised (start with just an email + certificate of incorporation); beyond that you upgrade to unpublished paid early-stage plans adding 409A and scale                                                             | Owens the founder equity layer no filer does: a free cap table for early startups, leading SAFE/round tooling and benchmarking, and a clean upgrade path to 409A and full equity management as you raise capital.                                                                                                                                 |
| <b>Build</b>                      | Founders with the time and legal literacy to file directly keep maximum control over structure and spend.                                  | Mostly one-time DE filing ~\$89-200 + registered agent ~\$50-300/yr + optional lawyer review ~\$1.5-3k                                                                                                                                                                                        | Lowest hard cost, total control, and you actually learn your own corporate structure, useful for a technical solo/duo founder pre-fundraise who wants to understand every doc they're signing.                                                                                                                                                    |

Full detail for every tool follows on its own page.

# Stripe Atlas

## BUILT FOR

Pre-seed and seed founders, often solo or a small technical team, forming a Delaware C-corp to raise venture money and issue equity. Best fit: the right call when you want a fast, standardized C-corp with a bank account and Stripe payments wired in, versus law-firm-grade or state-flexible setups.

---

## CHOOSE THIS IF YOU ARE

US founders on the venture track who want a Delaware C-corp and a bank in one weekend.

### BEST WHEN

You value the bundled ecosystem. \$500 one-time covers formation, the \$90 DE state fee, EIN, founder stock issuance with an 83(b) workflow, and year-one registered agent (\$100/yr after), plus ~\$2,500 Stripe credits and \$50K+ in Mercury/AWS/Xero perks.

### AVOID IF

You need attorney-grade, negotiable financing docs or the cheapest private LLC (go Northwest); Atlas now forms both a C-corp and an LLC, but only in Delaware.

## WHAT IT DOES

**Runs on a single \$500 one-time fee covering formation,** EIN and the first-year registered agent, now supports both a Delaware C-Corp and an LLC (chosen at onboarding), and is open to founders based anywhere.

**Supports both a Delaware C-Corp and an LLC:** founders pick the entity type during onboarding at the same \$500 one-time fee. Delaware is the state of formation and there is no other-state option, but the entity choice is no longer C-Corp-only.

**The guided flow generates the standard formation documents plus founder stock paperwork as part of the \$500 package.** Document scope is oriented toward standing up a C-Corp quickly rather than maintaining the broader post-incorporation legal library that a documents-first competitor keeps on hand.

**After formation, the main ongoing elements are Stripe payments plus the partner perks,** \$2,500 in Stripe credits and \$50k+ in partner perks. Continuing legal document tooling is lighter than a documents-focused platform. The value tilts toward payments and credits rather than ongoing paperwork.

**EIN filing and a US bank account are included in the \$500 package,** alongside \$2,500 in Stripe credits and \$50k+ in partner perks. Among these providers it offers the most complete banking-and-EIN bundle, folding account setup and payments into the single one-time formation fee.

---

## COST

**A single \$500 one-time fee covers formation,** EIN and the first-year registered agent. The recurring cost afterward is the \$100/year agent renewal from year two onward. There are no tiered service plans layered on, and the fee includes \$2,500 in Stripe credits and \$50k+ in partner perks.

---

## WHERE IT STANDS OUT

**The breadth of the bundle: formation,** EIN, a US bank account and Stripe payments delivered together within about two business days under one \$500 fee, with \$2,500 in Stripe credits and \$50k+ in partner perks. No other provider here folds banking and payments into the base.

---

## LIMITATIONS

**The Delaware-only state of formation is the main structural constraint,** though founders can now choose either a C-Corp or an LLC. Ongoing cost is modest but not zero (a \$100/year agent renewal applies from year two), and the continuing legal document library is lighter than a documents-first platform.

---

Pricing and limits read from [Stripe Atlas's own site](#) . Verified to the edition date above.

# Clerky

## BUILT FOR

First-time startup founders and their lawyers handling C-corp formation, 83(b) elections, SAFEs, and equity paperwork with legal precision. Best fit: pick it over Atlas when clean, VC-ready cap-table and financing documents matter more than speed or bundled banking.

---

## CHOOSE THIS IF YOU ARE

Founders who will run their own SAFEs, option grants, and priced rounds and want lawyer-trusted paperwork.

## BEST WHEN

Legal correctness beats onboarding polish. \$427 per-use incorporation (+ state fees) or an \$819 Company Lifetime Package bundling expedited filing, year-one registered agent, annual-report reminders and lifetime document workflows.

## AVOID IF

You want the banking-and-perks bundle (go Atlas), a non-resident package with a US address, or just cheap registered-agent service (go Northwest).

## WHAT IT DOES

**Company setup starts at \$427**, with no banking or payments layer bundled in. Oriented toward companies that expect to raise.

**Concentrates on the Delaware C-Corp**, the structure typically used by companies raising venture financing. It does not target LLC or other-state formations. The platform's scope is the C-Corp and the equity and governance paperwork that accompanies it, rather than a broad menu of entity types.

**Generates founder stock with vesting**, bylaws, IP assignment and board consents, a deeper set of incorporation and governance documents than the other providers here assemble. Document coverage, not banking or payments, is the product's focus and the reason raising-stage teams choose it.

**The \$819 Lifetime Package covers unlimited SAFEs**, convertible notes, hiring paperwork, equity grants and maintenance filings, so fundraising and compliance documents can be generated on an ongoing basis without per-document charges. This continuing document library is the platform's core post-formation function.



**Does not bundle banking or EIN handling.** The product scope is legal documents only. Founders using Clerky arrange banking and their EIN separately, so the platform sits alongside a bank and a payment provider rather than replacing them the way a fully bundled formation product would.

---

Pricing and limits read from [Clerky's own site](#) . Verified to the edition date above.

---

## COST

**Company setup starts at \$427**, with post-incorporation setup at \$299 and the Lifetime Package at \$819. Pricing is document-package based rather than subscription: the \$819 tier removes per-document charges on future SAFEs, notes and filings, but no banking or agent cost is bundled in.

---

## WHERE IT STANDS OUT

**It stands apart on document depth**, founder stock with vesting, bylaws, IP assignment, board consents, and unlimited SAFEs, convertible notes and equity grants under the \$819 Lifetime Package. For companies expecting to raise, the equity paperwork is more complete than the other providers offer.

---

## LIMITATIONS

**There is more legal detail for founders to work through**, and no banking, EIN or payments layer is bundled, so those must be arranged separately. Pricing is per document package (from \$427) instead of an all-in service, and only the Delaware C-Corp structure is supported.

# Firstbase

## BUILT FOR

Non-US and remote founders incorporating a US LLC or C-corp remotely, plus small teams wanting ongoing compliance handled. Best fit: choose it when you need US entity plus EIN, registered agent, and tax filings managed for a foreign or distributed founder, not just the initial filing.

---

## CHOOSE THIS IF YOU ARE

Non-US founders who need a US Delaware or Wyoming entity, EIN, and address without a US SSN.

## AVOID IF

You're US-based and VC-track or you just want the cheapest private US LLC (Northwest at ~\$575+ over 5 years).

---

## WHAT IT DOES

**Formation runs \$399 one-time with lifetime support**, while compliance and tax functions are sold as separate recurring add-ons rather than included in the base.

**Forms an LLC or a C-Corp, but initial formation is Delaware or Wyoming only (foreign qualification and registered agent are available in all 50 states)**. The LLC-or-C-Corp choice suits founders still deciding on entity type, and the platform is built for non-resident founders as well as US-based ones.

**Produces the standard formation documents and includes lifetime support with the \$399**

**Start plan**. The document set is geared to getting the entity formed; deeper equity and governance paperwork is not the platform's emphasis, which leans instead toward ongoing compliance and bookkeeping.

**Post-formation, Firstbase One at \$2,388/yr (\$199/mo billed annually) bundles bookkeeping**, tax filing and compliance tracking, while lighter needs can be met by individual add-ons such as Agent Autopilot at \$299/yr. The ongoing services, not the one-time filing, are where recurring cost concentrates.

**Includes EIN handling and offers banking,** mailroom and tax as add-ons, not a single bundle. The mailroom starts at \$35/month and tax filing runs about \$1,799/year, so the full banking-and-compliance stack is assembled from separately priced components layered on the \$399 base.

---

Pricing and limits read from [Firstbase's own site](#) . Verified to the edition date above.

---

## COST

**Firstbase Start is \$399 for formation.** Recurring costs add separately: Agent Autopilot \$299/year, Mailroom from ~\$35/month, tax filing about \$1,799/year (or ~\$899/year for a non-resident single-member LLC), and Firstbase One at \$2,388/year bundling bookkeeping, tax and compliance. Total ongoing cost depends on which add-ons you take.

---

## WHERE IT STANDS OUT

**Combining formation with ongoing compliance,** bookkeeping and tax in one platform, with Firstbase One at \$2,388/year consolidating those services. It suits founders who want the post-formation operating stack from the same provider that filed the entity, and it's tuned for non-resident founders.

---

## LIMITATIONS

**The recurring services are priced separately from the \$399 formation,** Agent Autopilot at \$299/year, mailroom and tax add-ons each billed on their own, so the headline formation price understates the ongoing cost for founders who want the full compliance and bookkeeping stack.

# doola

## BUILT FOR

Solopreneurs, e-commerce sellers, and international founders forming a US LLC on a budget with bookkeeping and tax add-ons. Best fit: the low-cost pick for a simple LLC and compliance when you are not raising venture capital and want an affordable all-in-one.

---

## CHOOSE THIS IF YOU ARE

Non-resident solo founders and ecommerce sellers who want formation plus ongoing compliance handled as a yearly service.

### BEST WHEN

You'd rather pay \$297/yr (+ state fees) to keep formation, EIN, registered agent, a US business address and a Mercury banking intro bundled and renewing automatically, and may add a \$1,999 tax tier.

### AVOID IF

You want state fees included up front and cheaper tax filing, or you're a US founder raising VC and need a C-corp with a real cap table.

## WHAT IT DOES

**Sold on annual plans starting at \$297/yr**, with state filing fees and franchise taxes billed on top of the subscription, not inside it.

**Forms an LLC, C-Corp or DAO LLC**, and completes the filing without requiring a US SSN, which is the main draw for non-resident founders. The DAO LLC option is unusual among these providers, extending coverage beyond the standard LLC and C-Corp entity types offered elsewhere in this group.

**Formation documents are oriented to non-resident founders and produced as part of the process that forms the US entity without a US SSN.** The paperwork focus is on completing the filing and enabling downstream tax and compliance work; startup equity and governance documents are outside its lane.

**Ongoing tooling covers bookkeeping**, tax and compliance, plus e-commerce analytics, sold through tiered annual plans, from Pulse bookkeeping at \$300/yr up to Tax and Compliance at \$1,999/yr. The recurring services are packaged into the plan structure instead of being offered piecemeal per document.

**Provides EIN and banking assistance geared to non-residents**, completing setup without a US SSN. Banking help is part of the non-resident onboarding (there is no bundled US bank account), and the service is delivered within the annual plan starting at \$297/yr instead of as a one-time add.

---

Pricing and limits read from [doola's own site](#) . Verified to the edition date above.

---

## COST

**Plans run from Starter at \$297/year to Tax and Compliance at \$1,999/year and Business-in-a-Box at \$2,999/year (\$329/month)**. State filing fees (\$50-\$500+) and annual franchise taxes (\$0-\$800/year) are billed on top, so the real cost sits above the plan price depending on the state of formation.

---

## WHERE IT STANDS OUT

**Its main advantage is forming a US LLC**, C-Corp or DAO LLC without a US SSN, aimed at non-resident founders, with tax and compliance attached to the plan. The DAO LLC option and the no-SSN process set it apart from the more US-resident-oriented providers in this group.

---

## LIMITATIONS

**Costs are structured as annual plans with recurring fees rather than a one-time formation**, and state filing fees (\$50-\$500+) and franchise taxes (\$0-\$800/year) fall outside the plan price. Founders who want only a one-off filing without an ongoing subscription may find the model less suitable.

# Northwest Registered Agent

## BUILT FOR

Founders and small businesses in any US state needing registered-agent service and LLC or corp filing with strong privacy and human support. Best fit: best when registered agent and privacy across states is the priority, rather than a startup-specific equity or fundraising toolkit.

---

## CHOOSE THIS IF YOU ARE

Bootstrapped US founders who want the cheapest, most private LLC and nothing they won't use.

## BEST WHEN

Cost and privacy dominate: roughly \$39 + state fee to form and \$125/yr for registered agent (about \$575+ all-in over five years), with Northwest listing its own address to keep yours off public records.

## AVOID IF

You're raising venture money and need a Delaware C-corp with 83(b) and a cap table, or you want banking credits and perks bundled (go Atlas).

## WHAT IT DOES

**Company formation is an optional add-on**, not part of the base service, so incorporation itself is not included unless it is purchased separately.

**Handles LLC and Corporation filings across US states**, but only through the formation add-on, not the base agent service. The registered-agent product itself does not include incorporation, so entity filing is a separate purchase layered onto the \$125/year agent plan.

**Formation documents are available only through the paid formation add-on.** The provider's document strength lies in registered-agent and privacy paperwork, the business address and mail handling, not in equity or governance documents aimed at financing or cap-table work.

**The ongoing product is the registered-agent service itself:** a business address and mail scanning maintained for \$125/year. There is no bookkeeping, tax or equity tooling layered on. Post-formation value is limited to the agent and mail-handling functions of the base plan.

**The bundle centers on a business address and mail scanning at \$125/year**, not banking or EIN.

Those functions sit outside the core registered-agent product, so founders needing an EIN or a bank account handle them elsewhere or through the separate formation add-on.

---

Pricing and limits read from Northwest Registered Agent's own site . Verified to the edition date above.

---

## COST

**The registered agent, business address**, mail scanning and privacy come to \$125/year with no hidden or per-document fees. Once you hold five or more registered-agent services with Northwest, the rate drops to \$100/year each. Incorporation is not part of this figure, since the formation add-on is priced separately from the agent service.

---

## WHERE IT STANDS OUT

**Its strength is the registered-agent and privacy service at \$125/year with no hidden or per-document fees**, dropping to \$100/year each once you hold five or more registered-agent services. For founders whose priority is agent coverage and a private business address rather than a formation bundle, it is a focused, lower-cost option.

---

## LIMITATIONS

**Incorporation is not part of the base \$125/year agent service.** Formation is a separate add-on. A founder wanting a single all-in formation bundle assembles it in pieces. The product is agent- and privacy-first, with banking, EIN and equity tooling outside its scope.

# Carta

## BUILT FOR

Funded startups and their finance leads managing cap tables, 409A valuations, equity grants, and investor reporting after incorporation. Best fit: adopt it once you have real equity to administer and investors to report to, not for the day-one filing itself.

---

## CHOOSE THIS IF YOU ARE

Founders whose day-one priority is clean equity: cap table, 409A readiness, and electronic share issuance for investors and an option pool.

## BEST WHEN

You've already formed and need equity managed properly. Carta Launch is free up to 25 stakeholders and \$1M raised, then custom-quoted paid early-stage tiers (no public flat monthly price like "\$599/mo"; standalone 409As from ~\$3,500) as you grow; choose it over the pure formation tools when your real problem is tracking ownership, SAFEs converting, and vesting, not filing the certificate.

## AVOID IF

You just need to incorporate and get an EIN or you're a tiny cap table where a spreadsheet suffices; Carta's median real cost climbs to ~\$14K+/yr as stakeholders grow.

## WHAT IT DOES

**Guides founders to a Delaware C-corp (the VC standard) or LLC via partner filers instead of filing with the state itself.** Cap-table/equity tooling serves US startups. The free Launch tier now also covers APAC & MENA.

**Offers incorporation resources plus a vendor/law-firm handoff instead of filing docs in-house.** Its real strength is auto-generating and tracking SAFEs, equity grants, board consents, and 83(b)/409A guidance.



**Deep: cap-table management, electronic securities issuance,** SAFE auto-generation and benchmarking, option grants, 409A valuations, board consents, scenario modeling and investor/stakeholder reporting.

---

Pricing and limits read from [Carta's own site](#) . Verified to the edition date above.

**Not a banking or EIN provider;** those come from the incorporation partner or Stripe/Mercury-type accounts; Carta focuses on equity infrastructure and, at paid tiers, fund administration and compliance.

---

## COST

**Carta Launch is free for startups with under 25 stakeholders and under \$1M raised (start with just an email + certificate of incorporation);** beyond that you upgrade to unpublished paid early-stage plans adding 409A and scale.

---

## WHERE IT STANDS OUT

**Owns the founder equity layer no filer does:** a free cap table for early startups, leading SAFE/round tooling and benchmarking, and a clean upgrade path to 409A and full equity management as you raise capital.

---

## LIMITATIONS

**Not a true formation or registered-agent service (relies on partners);** no EIN/banking; paid-tier pricing is opaque and gets pricey as you scale; overkill if you only need a filing; and it carries past data-privacy/secondaries scrutiny.

# Build

## BUILT FOR

Cheapest cash cost, highest time and error cost.

---

## CHOOSE THIS IF YOU ARE

Founders with the time and legal literacy to file directly and keep maximum control over structure and spend.

### BEST WHEN

You'll submit the Delaware/Wyoming certificate yourself (state fee ~\$90-\$110), get an EIN free from the IRS, and use startup-lawyer templates (or a flat-fee firm) for bylaws, stock purchase agreements and 83(b).

### AVOID IF

You're time-poor or unsure about 83(b) timing and founder-stock mechanics; a missed 83(b) window or malformed cap table costs far more than Atlas's \$500.

## WHAT IT DOES

### Any structure/state you're willing to research:

DE C-corp, LLC, S-corp, home-state entities. No guardrails. You pick and bear the consequences. Delaware C-corp is the standard VC path. Get the wrong one and re-domesticating later is costly.

**Templates only, no built-in counsel.** Fine for a clean single-founder cap table; risky for vesting, IP assignment, 83(b) elections and multi-founder splits. Budget a few hours of a startup lawyer (~\$1.5-3k) to review before signing, which is cheaper than fixing later.

**None bundled.** You self-assemble. Carta/Pulley for cap table, a calendar for DE franchise tax and annual reports, manual board-consent tracking. Easy to miss a filing deadline. A lapsed registered agent or unpaid franchise tax can void good standing.

**EIN is free directly from the IRS (SS-4 online, ~15 min);** SaaS charging for it adds nothing. Bank account you open yourself (Mercury/Brex). No bundled cash perks or partner credits. You forgo the ~\$5k+ in deals platforms broker.

## COST

**Mostly one-time: DE filing ~\$89-200 + registered agent ~\$50-300/yr + optional lawyer review ~\$1.5-3k.** Beats Atlas (\$500) or Clerky (\$400+) only if you value your setup hours cheaply and won't re-do it. Recurring: franchise tax + annual report you file yourself.

---

No vendor page: this route is assembled from the components named above.

## WHERE IT STANDS OUT

**Lowest hard cost, total control,** and you actually learn your own corporate structure, useful for a technical solo/duo founder pre-fundraise who wants to understand every doc they're signing.

## LIMITATIONS

**Breaks the moment complexity rises:** multiple founders, advisors, SAFEs, option pools or diligence. One botched 83(b) (30-day IRS deadline) or missing IP assignment can tank a financing. Doesn't scale; switch to Clerky/counsel before your first priced round.

# Bookkeeping & Accounting

Compared: QuickBooks Online · Xero · Wave · Puzzle · Zoho Books · Pilot · Build

# At a glance

| TOOL                     | CHOOSE THIS IF YOU ARE                                                                                                                          | STARTING PRICE                                                                                                                                                                                                              | MAIN BENEFIT                                                                                                                                                                                                                                                                                         |
|--------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>QuickBooks Online</b> | <b>SMBs and founders who want the software every US bookkeeper CPA already knows.</b>                                                           | <b>After the Aug 1, 2026 increase</b><br>paid plans run Simple Start \$38/mo, Essentials \$85, Plus \$140 and Advanced \$340 (the hike raised Essentials, Plus and Advanced), so rising prices are the main cost caveat     | <b>The accountant-default position (most US bookkeepers and CPAs already use it), combined with built-in payroll (add-on) and the widest integration ecosystem here.</b><br>For founders who want the tool their accountant already knows plus payroll under one roof, that combination is the draw. |
| <b>Xero</b>              | <b>Small teams and international founders who want multiple people in the books without per-seat fees</b>                                       | <b>No free tier</b>                                                                                                                                                                                                         | <b>Unlimited users on every plan (so team and accountant access adds no per-seat cost), paired with 1,000+ integrations and full double-entry accounting available across all tiers rather than gated to the top plan.</b><br>An established global vendor with enterprise-grade security.           |
| <b>Wave</b>              | <b>Solo founders, freelancers and pre-revenue startups who want real accounting and invoicing for \$0.</b>                                      | <b>Starter is \$0/mo and covers core accounting and invoicing</b>                                                                                                                                                           | <b>Core accounting and unlimited invoicing free at the Starter tier (\$0/mo), and it's the easiest tool here to set up.</b><br>For freelancers and very small businesses under \$500K revenue that want working books at no software cost, that's the standout.                                      |
| <b>Puzzle</b>            | <b>VC-backed US startups that must show investors accurate burn runway, and GAAP-clean books from day one.</b>                                  | <b>Starter is free for the first 2 months</b><br>then \$30/mo. Paid tiers step up to Core \$72/mo, Complete \$120/mo and Scale \$360/mo (roughly 17% off billed annually), so cost scales with seats, automation and volume | <b>AI-native automation as the core of the product, built around accrual, investor-ready books for tech startups, with connections to the startup stack (Stripe, banks).</b><br>For a founder who wants modern, automated books aligned to investor expectations, that focus is the differentiator.  |
| <b>Zoho Books</b>        | <b>Bootstrapped founders who need real inventory tracking</b><br>a full back-office suite on a genuinely cheap plan (free under \$50K revenue). | <b>Free forever for businesses under \$50K revenue</b><br>including 1 user plus 1 accountant, up to 1,000 invoices/year, bank reconciliation and 50+ reports                                                                | <b>A free-forever plan (under \$50K revenue) that already includes bank reconciliation and 50+ reports, plus deep integration across the wider Zoho app suite.</b><br>For micro-businesses wanting real accounting at no cost within one ecosystem, that combination stands out.                     |
| <b>Pilot</b>             | <b>Funded founders who want to hand bookkeeping to humans</b><br>never open accounting software.                                                | <b>No free tier</b><br>annual billing                                                                                                                                                                                       | <b>Purpose-built for VC-backed startups: real US bookkeepers, accrual/GAAP + ASC 606 from day one, investor-grade reporting, R&amp;D credit help, deep startup-tool integrations.</b>                                                                                                                |
| <b>Build</b>             | <b>Extremely early or pre-revenue founders with a handful of transactions a month</b>                                                           | <b>Effectively free software</b><br>(Sheets) + your time + a CPA at tax time (~\$500-2k/yr)                                                                                                                                 | <b>Best for the earliest stage: a pre-revenue or few-transactions startup that wants near-zero burn and full visibility into its own numbers, with a founder who's financially literate and disciplined.</b>                                                                                         |

Full detail for every tool follows on its own page.

# QuickBooks Online

## BUILT FOR

Small businesses and their bookkeepers or accountants who want the US-standard general ledger, invoicing, payroll, and tax-ready books. Best fit: the safe default when you need an accountant who already knows the tool and deep third-party integrations, over leaner or newer options.

---

## CHOOSE THIS IF YOU ARE

SMBs and founders who want the software every US bookkeeper and CPA already knows.

## BEST WHEN

Universal accountant support and ecosystem depth matter more than startup-specific features: paid tiers Simple Start \$38, Essentials \$85, Plus \$140 and Advanced \$340/mo after the Aug 1, 2026 hike (Intuit's QuickBooks Free is a separate \$0 product), with the deepest app marketplace and payroll/payments integrations.

## AVOID IF

You're VC-backed and need daily runway, simultaneous cash+accrual GAAP books and SaaS rev-rec, or you're pre-revenue and want free.

---

## WHAT IT DOES

**It offers add-on payroll and a wide integration ecosystem**, with paid tiers running from Simple Start (\$38/mo) to Advanced (\$340/mo) after the Aug 1, 2026 increase. Intuit also offers QuickBooks Free, a separate \$0/mo product distinct from these QuickBooks Online paid tiers.

**The US accountant default.** Most bookkeepers and CPAs already work in it, which simplifies handoffs, reviews and later migrations (its export is accountant-familiar). Available across web and mobile. The most broadly supported option among US practices in this set.

**Produces standard accounting reports from a platform accountants and diligence teams already recognize**, which reduces friction during reviews. Deep accounting features across tiers support the reporting depth investors expect as a company grows.

**Full accounting with deep feature depth across Simple Start, Essentials, Plus and Advanced.** Capabilities expand as you move up tiers. Built-in payroll is available as an add-on. Feature breadth is the reason it anchors many US practices.

**Automation covers transaction categorization and report generation.** It's practical, workflow-level automation aimed at reducing manual bookkeeping rather than an AI-native rebuild of the accounting process.

---

### COST

**After the Aug 1, 2026 increase,** paid plans run Simple Start \$38/mo, Essentials \$85, Plus \$140 and Advanced \$340 (the hike raised Essentials, Plus and Advanced), so rising prices are the main cost caveat. Separately, Intuit offers QuickBooks Free (\$0/mo, with monthly caps), a distinct product rather than a tier of QuickBooks Online. Payroll is a separate add-on.

---

### ECOSYSTEM & EXIT

**Scales across tiers from Simple Start (\$38/mo) to Advanced (\$340/mo) after the Aug 1, 2026 increase,** and its standard export plus accountant familiarity make moving data or switching preparers straightforward. The main scaling caveat is rising prices, with increases effective Aug 1, 2026.

**The widest integration ecosystem in this group,** and payroll is built in as an add-on rather than requiring a separate third-party tool. For founders who want the accountant-default plus payroll under one vendor, integration breadth is a primary draw.

---

### WHERE IT STANDS OUT

**The accountant-default position (most US bookkeepers and CPAs already use it),** combined with built-in payroll (add-on) and the widest integration ecosystem here. For founders who want the tool their accountant already knows plus payroll under one roof, that combination is the draw.

---

### LIMITATIONS

**Paid-tier pricing is high and rising.** Increases took effect Aug 1, 2026 for Essentials, Plus and Advanced, with Advanced now \$340/mo (up from \$200). Intuit's separate QuickBooks Free product exists but is distinct from these QuickBooks Online plans, and payroll costs extra, so cost and its trajectory are the main drawbacks.

---

Pricing and limits read from [QuickBooks Online's own site](#) .  
Verified to the edition date above.

# Xero

## BUILT FOR

Small businesses and accounting firms, strong outside the US, wanting clean double-entry books, unlimited users, and a broad app ecosystem. Best fit: choose it over QuickBooks for better multi-user pricing and international workflows, especially outside North America.

---

## CHOOSE THIS IF YOU ARE

Small teams and international founders who want multiple people in the books without per-seat fees.

## BEST WHEN

Collaboration and global basics drive the choice. Plans from ~\$25/mo with unlimited users at no extra charge (versus QBO's tier-gated user counts), strong multi-currency and bank-feed support on higher tiers.

## AVOID IF

You specifically need the widest US CPA familiarity (QBO wins there) or startup-native accrual/runway/rev-rec built in; Xero's US bookkeeper pool is thinner than QBO's.

## WHAT IT DOES

US tiers run **Early (\$25/mo), Growing (\$55/mo) and Established (\$90/mo, which adds multi-currency)**. Aimed at growing businesses.

**Full double-entry accounting backed by an established global vendor and enterprise-grade security.** Every plan includes unlimited users. Giving an accountant or bookkeeper access adds no per-seat cost, useful when several reviewers touch the books. Available on web and mobile.

**Full double-entry accounting produces complete financial statements suitable for diligence.** Enterprise-grade security and an established global vendor add credibility, and the reporting depth is consistent across every plan, with nothing gated to higher tiers.

**Full double-entry accounting on every plan,** with bank reconciliation automation built in. Feature depth is consistent across tiers rather than gated, and 1,000+ integrations extend functionality. Established adds multi-currency for cross-border work.



**Automation centers on bank reconciliation.**

Matching imported transactions to reduce manual coding. It's targeted at the highest-volume bookkeeping task, not a broad AI layer across the product.

---

**COST**

**No free tier.** Early \$25/mo, Growing \$55/mo and Established \$90/mo, with Established adding multi-currency. Every plan includes unlimited users. Unlike per-seat tools the headline price is the full team cost regardless of headcount.

---

**ECOSYSTEM & EXIT**

**Scales without per-seat friction since every plan includes unlimited users**, and US tiers run Early (\$25/mo) through Established (\$90/mo), where Established adds multi-currency for cross-border growth. Standard data export keeps later migration straightforward.

**1,000+ integrations, one of the larger app marketplaces here**, so you extend the core accounting with third-party tools. Payroll is added through those integrations or partners, not bundled into the product itself.

---

**WHERE IT STANDS OUT**

**Unlimited users on every plan (so team and accountant access adds no per-seat cost)**, paired with 1,000+ integrations and full double-entry accounting available across all tiers rather than gated to the top plan. An established global vendor with enterprise-grade security.

---

**LIMITATIONS**

**No free tier (the entry point is Early at \$25/mo)**, so there is no zero-cost option for the smallest users. Setup carries a moderate learning curve, and multi-currency is reserved for the top Established tier (\$90/mo); lower plans go without.

---

Pricing and limits read from [Xero's own site](#) . Verified to the edition date above.

# Wave

## BUILT FOR

Freelancers, sole proprietors, and very small businesses wanting free invoicing and basic accounting with minimal bookkeeping knowledge.

---

## CHOOSE THIS IF YOU ARE

Solo founders, freelancers, and pre-revenue startups who want real accounting and invoicing for \$0.

## BEST WHEN

You're keeping costs at zero and volume is low. Core accounting and unlimited invoicing are free, with Pro at \$19/mo adding bank feeds and multi-user access (receipt scanning is a separate add-on at \$8-11/mo).

## AVOID IF

You need burn/runway and accrual GAAP for a raise, inventory tracking, or the CPA-universal standard for a growing team (go QuickBooks); Wave's reporting and integrations are thin.

## WHAT IT DOES

**Built for freelancers, sole proprietors and very small businesses under \$500K revenue.** A Pro tier (\$19/mo) adds bank auto-import and multi-user access; receipt scanning is a separate add-on (\$8/mo on Pro, \$11/mo on Starter).

**Suited to simple books for freelancers and sole proprietors.** It covers core accounting only. It fits fewer complex practices. Access is web and mobile, and the Starter tier (\$0/mo) allows one user, with multi-user access requiring the Pro tier. Standard data export eases handoff.

**Reporting is basic and matched to core accounting only, adequate for pre-raise freelancers and very small businesses,** but thinner than the double-entry reporting investors and diligence typically expect. Most founders migrate up before raising.

**Covers core accounting and invoicing.** Bank auto-import and multi-user access now sit behind the Pro tier (\$19/mo) instead of the free Starter plan, and receipt scanning is a separate paid add-on (\$8/mo on Pro, \$11/mo on Starter), a limit for anyone needing automatic bank feeds or receipt capture at no cost.

**Automation is limited to receipt scanning,** which extracts data from captured receipts. Beyond that the automation footprint is minimal, consistent with its core-accounting-only scope and free positioning.

## COST

**Starter is \$0/mo and covers core accounting and invoicing.** Pro is \$19/mo and adds bank auto-import and multi-user access. Receipt scanning is a separate add-on (\$8/mo on Pro, \$11/mo on Starter). Several features that were previously free (including receipt scanning and bank import) moved behind payment in early 2024.

## ECOSYSTEM & EXIT

**Best suited under \$500K revenue.** Its core-accounting-only scope means growing businesses typically outgrow it and migrate up. Standard data export supports that move, and the Pro tier (\$19/mo) extends runway with bank auto-import and multi-user access.

**Core integrations, with payments and payroll offered as paid add-ons outside the free Starter plan.** The integration footprint is narrower than the incumbents', matching its core-accounting-only, small-business scope.

## WHERE IT STANDS OUT

**Core accounting and unlimited invoicing free at the Starter tier (\$0/mo),** and it's the easiest tool here to set up. For freelancers and very small businesses under \$500K revenue that want working books at no software cost, that's the standout.

## LIMITATIONS

**Scope is core accounting only and best fit under \$500K revenue,** so growing businesses outgrow it. Several previously-free features moved behind payment. Bank auto-import and multi-user access need Pro (\$19/mo), and receipt scanning is a separate add-on (\$8/mo on Pro, \$11/mo on Starter), narrowing what the free Starter plan covers.

---

Pricing and limits read from **Wave's own site** . Verified to the edition date above.

# Puzzle

## BUILT FOR

Startup founders and finance-savvy operators who want accrual-based, real-time books built for venture metrics and runway. Best fit: pick it over QuickBooks when you are a tech startup wanting modern, automated accounting tied to burn and runway instead of legacy small-business books.

---

## CHOOSE THIS IF YOU ARE

VC-backed US startups that must show investors accurate burn, runway, and GAAP-clean books from day one.

## BEST WHEN

Starter is free for the first 2 months, then \$30/mo (1 seat), Core \$72 (5 seats), Complete \$120 and Scale \$360 (unlimited seats; ~17% off annually).

## AVOID IF

You're a non-venture SMB where a bookkeeper's QBO familiarity matters more (go QuickBooks), you want free forever, or you'd rather outsource the work entirely to humans.

## WHAT IT DOES

**Plans are Starter \$30/mo (first 2 months free),** Core \$72/mo, Complete \$120/mo and Scale \$360/mo (about 17% off billed annually). It targets founders who want modern books over accountant-default tools.

**A newer, startup-oriented vendor accessed via web.** Its books are accrual and investor-ready, aligned with how startups and their reviewers expect accounts to be kept, though as a newer product fewer accountants have direct familiarity with it. Export is available but the tooling is newer.

**Built specifically around accrual,** investor-ready books, the accounting basis startups are expected to present to investors. This is its core design intent rather than an add-on, aligning the default output with what diligence and boards look for.

**Accrual accounting with startup-specific automation as the core of the product.** The feature set centers on investor-ready books and connections to the startup stack (Stripe, banks), not broad general-business accounting modules.

**AI-native automation is the core differentiator.**

The product is built around automating startup accounting. This is the main reason a founder would reach for it over the incumbent tools.

---

**COST**

**Starter is free for the first 2 months**, then \$30/mo. Paid tiers step up to Core \$72/mo, Complete \$120/mo and Scale \$360/mo (roughly 17% off billed annually), so cost scales with seats, automation and volume. There is no permanent free tier.

---

**ECOSYSTEM & EXIT**

**Scales for startups from Starter (\$30/mo, first 2 months free) up through Scale (\$360/mo) as volume and automation needs grow.** Export is available, though as a newer product the migration tooling is less established than the incumbents'.

**Integrates with the modern startup stack (notably Stripe and bank connections)**, not a broad third-party marketplace. The focus is the tools startups already run, which suits its target users but is narrower than the incumbents' ecosystems.

---

**WHERE IT STANDS OUT**

**AI-native automation as the core of the product**, built around accrual, investor-ready books for tech startups, with connections to the startup stack (Stripe, banks). For a founder who wants modern, automated books aligned to investor expectations, that focus is the differentiator.

---

**LIMITATIONS**

**A newer vendor with a shorter track record and export tooling less established than the incumbents'.** Its free entry is time-limited to the first 2 months, after which Starter is \$30/mo and paid tiers reach Scale at \$360/mo. Fewer accountants are familiar with it.

---

Pricing and limits read from **Puzzle's own site** . Verified to the edition date above.

# Zoho Books

## BUILT FOR

Small and mid-size businesses already using or considering the Zoho suite, wanting affordable accounting with automation and inventory. Best fit: best when you want a low-cost, feature-rich ledger that plugs into Zoho CRM and apps, over standalone incumbents.

---

## CHOOSE THIS IF YOU ARE

Bootstrapped founders who need real inventory tracking and a full back-office suite on a genuinely cheap plan (free under \$50K revenue).

### BEST WHEN

Value and inventory drive the pick: free under \$50K annual revenue, Standard \$15/mo, Professional \$40/mo (unlocks inventory), scaling to \$120 Elite and \$240 Ultimate, with tight Zoho CRM/Inventory/Mail integration.

### AVOID IF

Your accountant expects QuickBooks (go QBO), you're a venture startup needing accrual burn/runway, or you want a truly zero-cost invoicing-first tool.

---

## WHAT IT DOES

**Even the free plan carries bank reconciliation and 50+ reports;** paid tiers run from Standard (\$15/mo) up to ~\$240/mo. Aimed at micro-businesses and founders.

**Full accounting accessed via web and mobile,** and the free plan explicitly allows 1 user plus 1 accountant, so accountant access is built into even the no-cost tier. Fewer US accountants use it than QuickBooks, which can add friction when handing books to an outside CPA.

**Full accounting with 50+ reports available even on the free plan,** giving complete financial reporting suitable as a business grows. Bank reconciliation is included from the free tier, supporting the clean, reconciled books diligence expects.

**Full accounting features, and unusually the free plan already includes bank reconciliation and 50+ reports.** The free tier does carry explicit limits (1 user plus 1 accountant and up to 1,000 invoices/year), beyond which you move to paid tiers.

**Automation is delivered across the broader Zoho suite**, so its value grows as you adopt other Zoho apps. Within Books it supports standard accounting workflows. The differentiator is ecosystem-wide automation, not a standalone AI feature.

---

### COST

**Free forever for businesses under \$50K revenue**, including 1 user plus 1 accountant, up to 1,000 invoices/year, bank reconciliation and 50+ reports. Paid plans start at Standard \$15/mo and run up to ~\$240/mo.

---

### ECOSYSTEM & EXIT

**Scales within the Zoho ecosystem**, moving from the free-forever plan (under \$50K revenue) to Standard (\$15/mo) and up to ~\$240/mo. Standard data export supports migration, though the value is strongest for users staying inside the Zoho suite.

**Integration is deepest inside the Zoho ecosystem**, connecting Books to the wider suite of Zoho apps. That makes it strong for teams already on Zoho, but the ecosystem is more self-contained than a broad open marketplace.

---

### WHERE IT STANDS OUT

**A free-forever plan (under \$50K revenue) that already includes bank reconciliation and 50+ reports**, plus deep integration across the wider Zoho app suite. For micro-businesses wanting real accounting at no cost within one ecosystem, that combination stands out.

---

### LIMITATIONS

**Fewer US accountants use it than QuickBooks**, which can add friction when handing off to an outside CPA, and its value is strongest inside the Zoho ecosystem rather than as a standalone tool. The free plan is also capped at 1 user, 1 accountant and 1,000 invoices/year.

---

Pricing and limits read from [Zoho Books's own site](#) . Verified to the edition date above.

# Pilot

## BUILT FOR

Funded startups and growing companies that want to outsource bookkeeping to a human team rather than run software themselves. Best fit: choose it when you would rather hand off monthly close, tax, and CFO support than do the books in-house.

---

## CHOOSE THIS IF YOU ARE

Funded founders who want to hand bookkeeping to humans and never open accounting software.

## BEST WHEN

You'd rather buy the outcome than run the tool: managed US-based bookkeepers on accrual basis, with an Essentials tier around \$99/mo for very low-expense companies stepping up to Core at ~\$499+/mo (rising with monthly spend), plus optional tax and CFO services.

## AVOID IF

You want to DIY cheaply, or you specifically want a live self-serve burn/runway dashboard you control day to day.

---

## WHAT IT DOES

**Pilot's US team runs your books monthly on QuickBooks Online**, with optional tax, CFO and payroll add-ons.

**Fully US-based dedicated bookkeeper; you need no accountant.** In-app messaging on all plans, phone+email from Core. Books live in QuickBooks so any CPA can view/export.

**Strong: accrual/GAAP by default**, ASC 606 revenue recognition, and a reporting pack covering burn, runway, ARR and gross margin that VCs and auditors expect.

**Monthly close, reconciliation, cash or accrual books**, custom chart of accounts; reports by the 10th business day. Bill mgmt capped at 10/mo on Core; AI-only Essentials has no human bookkeeper.

**AI categorizes transactions, reconciles and closes books monthly**, and answers questions on cash flow; on Core/Custom AI augments a human bookkeeper instead of replacing them.

---

## COST

**No free tier, annual billing.** Essentials (AI-only) \$99/mo up to \$100K monthly expenses; Core (human) scales with expenses; Custom quoted. Tax from ~\$750-2,450+/yr.



---

## ECOSYSTEM & EXIT

**High: all books are in QuickBooks Online and exportable.** No proprietary-ledger lock-in; Custom plan scales to multi-entity, AR/AP, rev-rec and payroll.

**Connects to Stripe, Brex, Mercury, Ramp, Gusto, Expensify and more;** built on QuickBooks. Payroll admin via the Custom plan / COO add-on.

---

---

Pricing and limits read from **Pilot's own site** . Verified to the edition date above.

---

## WHERE IT STANDS OUT

**Purpose-built for VC-backed startups:** real US bookkeepers, accrual/GAAP + ASC 606 from day one, investor-grade reporting, R&D credit help, deep startup-tool integrations.

---

---

## LIMITATIONS

**Premium pricing that climbs with monthly expenses;** tax sold separately; AI-only tier lacks human review. Bill-pay capped at 10/mo on Core; AR/AP and payroll need the Custom plan.

# Build

## BUILT FOR

Roll-your-own books: a spreadsheet ledger you maintain plus a part-time bookkeeper or CPA pulled in at tax time. Near-zero software cost, works only at low transaction volume and while you're disciplined.

---

## CHOOSE THIS IF YOU ARE

Extremely early or pre-revenue founders with a handful of transactions a month.

## BEST WHEN

Volume is trivial and cash is precious. A simple spreadsheet ledger plus a part-time CPA for quarterly cleanup and taxes costs almost nothing and keeps you close to the numbers.

## AVOID IF

You're doing more than ~20-30 transactions a month, taking investor money, or need accrual GAAP and audit-ready books; reconciliations and rev-rec bury a spreadsheet fast.

## WHAT IT DOES

**No QuickBooks/Bench/Pilot.**

**You source your own CPA: any local firm or freelancer**, no platform lock-in. Total geographic freedom, but no vetting or continuity. A solo CPA disappearing at tax season is your problem to solve. Referrals from other founders beat cold-hiring.

**Weak.** Spreadsheets aren't GAAP accrual books; diligence teams and auditors expect real financials. Fine for a pre-seed check, but you'll likely pay a firm to rebuild 1-2 years of clean books before a Series A, a pay-twice risk.

**You build A/R, A/P, categorization and reports by hand.** No enforced audit trail, no automatic reconciliation, no locking closed periods. Errors compound silently. Adequate for a simple cash-in/cash-out MVP, not for accrual, inventory or revenue recognition.

**None unless you wire it: e.g. a script or LLM prompt to auto-categorize CSV exports.** Doable for a technical founder but you own the accuracy. A miscategorized deductible is a tax cost. Off-the-shelf tools already do this reliably.

## COST

**Effectively free software (Sheets) + your time + a CPA at tax time (~\$500-2k/yr).** Cheaper than Pilot/Bench (\$200-1k+/mo) only while volume is tiny. The crossover comes fast once your hours or clean-up costs are counted.

---

No vendor page: this route is assembled from the components named above.

## ECOSYSTEM & EXIT

**Data is 100% yours in a file:** maximally portable, with no export fees. But it doesn't scale: past ~50-100 monthly transactions, payroll, or multi-entity, manual reconciliation collapses into hours of error-prone work. Plan to migrate to accounting software early.

**No native feeds.** You manually import bank/Stripe CSVs. Payroll must be a separate service (Gusto ~\$40+/mo); DIY payroll tax filing is a compliance minefield you should not self-build. This is where DIY leaks the most hours.

## WHERE IT STANDS OUT

**Best for the earliest stage: a pre-revenue or few-transactions startup that wants near-zero burn and full visibility into its own numbers,** with a founder who's financially literate and disciplined.

## LIMITATIONS

**No audit trail, no controls, no scalability,** and easy to get wrong in ways that are legally/financially risky (tax filings, sales tax nexus, 409A). Manual books are a stopgap; move to real accounting software before revenue or your first raise.

# Customer Support & Helpdesk

Compared: Intercom · Help Scout · Freshdesk · Crisp · Chatwoot · Zendesk · Build

# At a glance

| TOOL       | CHOOSE THIS IF YOU ARE                                                                          | STARTING PRICE                                                                                                                                                                                                                                                                                                             | MAIN BENEFIT                                                                                                                                                                                                                                                                |
|------------|-------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Intercom   | Product-led SaaS teams that want in-app messaging plus an AI agent deflecting tickets           | <b>Per seat on three annual tiers</b><br>(Essential \$29, Advanced \$85 and Expert \$132 per seat/mo), with Fin AI billed separately at \$0.99 per outcome                                                                                                                                                                 | <b>Its distinguishing element is the Fin AI agent operating across chat, email, in-app and social, billed per outcome at \$0.99.</b><br>The standout is a messaging-first suite where an AI agent handles front-line conversations, backed by deep customization.           |
| Help Scout | Small teams who want email-first human, uncluttered support that feels personal.                | <b>Per seat across Standard \$25</b><br>Plus \$45 and Pro \$75 per user/mo (monthly billing; annual is roughly 16% less), with AI Answers added at \$0.75 per resolution and a 3-month free trial                                                                                                                          | <b>It stands out on simplicity.</b><br>It is the easiest to set up here, built around a shared inbox with an easy import and metered AI Answers at \$0.75 per resolution. The appeal is a low-configuration shared-inbox desk for small teams.                              |
| Freshdesk  | Cost-conscious teams that want a full ticketing suite with the cheapest per-seat entry          | <b>A free program for up to 2 agents</b><br>(currently a 6-month term under Freshworks' terms), then paid tiers from Growth at \$19/agent/mo up to Enterprise at \$89/agent/mo. Pricing is per agent, with a low-cost entry point for teams of up to two and a range that extends from small teams to enterprise ticketing | <b>Its standout is classic ticketing depth (SLAs, routing and reporting with deep configuration) combined with a free program for up to 2 agents (currently a 6-month term).</b><br>It is full ticketing that a very small team can start on at no cost for that period.    |
| Crisp      | Small-to-mid teams of ~10-20 who get penalized by per-seat pricing want everything flat.        | <b>Flat per workspace rather than per seat</b><br>Free \$0, Mini \$45/mo for 4 seats, Essentials \$95/mo for 10 seats and Plus \$295/mo, each including bundled AI credits (\$5, \$25, \$75)                                                                                                                               | <b>The differentiator is flat per-workspace pricing (Free \$0, Mini \$45, Essentials \$95, Plus \$295) with AI credits bundled into each tier and a Hugo agent resolving up to ~60% of repetitive questions.</b><br>Cost is predictable per workspace rather than per seat. |
| Chatwoot   | Technical teams that want an open-source self-hostable helpdesk they fully control.             | <b>Self-hosting the open-source version is free beyond the infrastructure you run it on; paid self-hosted and cloud plans run from \$19 to \$99/agent/mo. Cost is therefore either infrastructure-only or a per-agent plan, depending on whether you self-host the free build or take a paid tier</b>                      | <b>Its standout is open-source (MIT) data ownership.</b><br>Self-hosting is free beyond infrastructure, AI is included in the base, and it positions as an alternative to Intercom, Zendesk and Salesforce Service Cloud for teams that can self-host.                      |
| Zendesk    | Scaling mid-market and enterprise support orgs that need deep routing SLAs, reporting, and WFM. | <b>Per-agent/month, billed annually</b><br>Suite Team \$55, Suite Professional \$115, plus Suite Enterprise (quote; Copilot add-on \$50/agent/mo)                                                                                                                                                                          | <b>Mature, scalable omnichannel ticketing with extensive routing/SLA/reporting and a large app marketplace.</b><br>A capable choice for teams expecting complex, high-volume, multi-channel support as they grow.                                                           |
| Build      | Pre-launch or tiny teams handling a trickle of support with near-zero budget                    | <b>No per-seat fees</b><br>which is the big win                                                                                                                                                                                                                                                                            | <b>Ideal for an early team with a technical member.</b><br>Near-zero per-seat cost, full data ownership (good for privacy-sensitive users), and no vendor lock-in while volume is low.                                                                                      |

Full detail for every tool follows on its own page.

# Intercom

## BUILT FOR

Product-led SaaS companies with support and success teams wanting in-app messaging, AI resolution, and proactive engagement. Best fit: the pick when live chat, in-product onboarding, and an AI agent like Fin matter more than classic ticket queues, and you can pay for it.

---

## CHOOSE THIS IF YOU ARE

Product-led SaaS teams that want in-app messaging plus an AI agent deflecting tickets.

### BEST WHEN

You want Fin wired into the messenger: seats at \$29/\$85/\$132 per agent/mo (Essential/Advanced/Expert, annual) plus Fin at \$0.99 per resolution.

### AVOID IF

Costs must be predictable and you're volume-heavy; Fin's per-resolution meter makes bills climb (8 agents + ~2,100 resolutions ≈ \$2,800/mo).

---

## WHAT IT DOES

**It covers chat, email, in-app messaging and social in one workspace**, and is aimed at teams wanting a messaging-first support setup with an AI agent handling front-line conversations.

**Chat, email, in-app messaging and social.** The in-app channel is central to its messaging-first positioning, letting conversations run inside the product as well as over email and social, all managed from a single support workspace.

**The Fin AI agent, priced at \$0.99 per outcome**, where an outcome is a resolution, handoff or disqualification, billed on top of seats. Alongside Fin it offers workflows for automation, so AI cost scales with usage rather than being bundled into the seat price.

**Self-service is driven by the Fin AI agent**, which handles front-line queries and is billed at \$0.99 per outcome (resolution, handoff or disqualification). Deflection comes from Fin resolving conversations across the channels, with each automated resolution counting as a paid outcome.

**Intercom offers deep customization within a conversation-oriented messaging suite**, not a classic ticket system. Routing and reporting sit around Fin-handled conversations. Ticketing is shaped by the messaging model, not by explicit SLA tooling.

**Collaboration runs through a team workspace where conversations across chat**, email, in-app and social are assigned and worked together. Because it is a messaging suite, shared handling is organized around conversations rather than a traditional email inbox.

---

## COST

**Per seat on three annual tiers (Essential \$29, Advanced \$85 and Expert \$132 per seat/mo)**, with Fin AI billed separately at \$0.99 per outcome. Total cost is a seat base plus usage. Spend rises with both headcount and the volume of Fin outcomes.

---

## ECOSYSTEM & EXIT

**A large app store provides a broad catalogue of prebuilt integrations alongside its APIs**. This breadth suits teams that need to connect the support suite to many surrounding tools without custom build work.

**Setup is moderate and migration is a standard import from other help desks**. Getting live takes some configuration for the messaging suite, but moving existing data in follows a conventional import path rather than a custom project.

---

## WHERE IT STANDS OUT

**Its distinguishing element is the Fin AI agent operating across chat**, email, in-app and social, billed per outcome at \$0.99. The standout is a messaging-first suite where an AI agent handles front-line conversations, backed by deep customization.

---

## LIMITATIONS

**Cost is the main limitation**. Seats run \$29 to \$132 per seat/mo and Fin AI adds \$0.99 per outcome on top, so costs stack across seats plus usage and can climb as both headcount and Fin outcomes grow.

---

Pricing and limits read from [Intercom's own site](#) . Verified to the edition date above.

# Help Scout

## BUILT FOR

Small to mid-size teams wanting a shared inbox and simple help center without the complexity or price of enterprise tools. Best fit: right when email-centric, human support with light automation is the goal, over heavier chat-first or enterprise suites.

---

## CHOOSE THIS IF YOU ARE

Small teams who want email-first, human, uncluttered support that feels personal.

### BEST WHEN

You value a clean shared inbox over a sprawling ticketing suite: Standard \$25, Plus \$45, Pro \$75 per user/mo (monthly; roughly 16% less billed annually) with an AI Answers add-on at \$0.75 per resolution, staying affordable up to ~25 agents.

### AVOID IF

You need the strongest AI deflection, enterprise routing/CSAT/WFM depth at scale, a free self-hostable option, or flat non-per-seat billing for a big team.

---

## WHAT IT DOES

**It targets small teams that want a shared inbox plus a knowledge base without heavy ticketing**, covering email, chat and a knowledge base as its core channels.

**Email, live chat and a knowledge base.** The channel set is narrower than the phone-and-social desks, matching its email-first, shared-inbox focus; the knowledge base doubles as a self-service channel alongside email and chat.

**AI Answers, priced at \$0.75 per resolution**, with a 3-month free trial on new accounts, billed on top of seats. It is paired with saved replies and automations, so AI cost is metered per resolved conversation, not included in the plan.

**The knowledge base is one of its listed channels**, and self-service is reinforced by AI Answers, which resolves questions at \$0.75 per resolution. A customer can self-serve through documentation or the metered AI answering layer before reaching the shared inbox.



**Customization is simple, and the desk is organized around a shared inbox instead of heavy SLA configuration.** Ticketing exists but is lighter than a full ticket system, matching the email-first, small-team positioning.

**The shared inbox is the core of the product.** It is a shared-inbox help desk for small teams. Collaboration is organized around that single inbox for email and chat, with the knowledge base attached, not around ticket ownership.

---

## COST

**Per seat across Standard \$25**, Plus \$45 and Pro \$75 per user/mo (monthly billing; annual is roughly 16% less), with AI Answers added at \$0.75 per resolution and a 3-month free trial. Cost is a per-user base plus metered AI. It scales with both seats and resolved conversations.

---

## ECOSYSTEM & EXIT

**Common integrations and an API connect it to the usual surrounding tools**, with programmatic extension available. The integration surface is standard rather than an extensive marketplace, matching its simpler positioning.

**Setup is the easiest in this group**, with an easy import for migration. The path to going live is fast and data migration is low-friction, consistent with its simple, small-team shared-inbox design.

---

## WHERE IT STANDS OUT

**It stands out on simplicity.** It is the easiest to set up here, built around a shared inbox with an easy import and metered AI Answers at \$0.75 per resolution. The appeal is a low-configuration shared-inbox desk for small teams.

---

## LIMITATIONS

**It is a simple product.** Customization is light and ticketing is thinner than a full SLA-driven system. Teams needing deep ticketing, routing or SLA configuration may find it thinner, and AI Answers is a metered \$0.75-per-resolution add-on.

---

Pricing and limits read from **Help Scout's own site** . Verified to the edition date above.

# Freshdesk

## BUILT FOR

Growing support teams wanting a full ticketing helpdesk with multichannel, automations, and tiered pricing that scales.

---

## CHOOSE THIS IF YOU ARE

Cost-conscious teams that want a full ticketing suite with the cheapest per-seat entry.

### BEST WHEN

Budget per agent is the deciding factor: free for 1-2 agents for 6 months, then Growth \$19, Pro \$55, Enterprise \$89 per agent/mo (annual), with a Freddy AI Copilot add-on at ~\$29/agent.

### AVOID IF

You want stronger AI deflection and product messaging, a truly free self-hosted stack, or flat workspace pricing instead of per-agent.

---

## WHAT IT DOES

**It is aimed at teams wanting classic ticket-based support with SLA and reporting depth,** spanning email, chat, phone and social, and it offers automations and the Freddy AI assistant.

**Email, chat, phone and social.** It carries the widest channel set here, including a phone channel, which fits its role as a full ticketing desk handling support across email, chat, phone and social from one system.

**Automations plus Freddy AI.** Freddy Copilot is a paid add-on (about \$29/agent/mo, on Pro and Enterprise) and the customer-facing Freddy AI Agent is billed per session. AI is a paid layer sitting on top of the base plan.

**Self-service pairs with the desk's automation tooling and Freddy AI.** The customer-facing Freddy AI Agent is billed per session and Freddy Copilot is a ~\$29/agent/mo add-on, so deflection runs through automations plus a metered/paid AI layer rather than a fully bundled assistant.

**Ticketing is the center of the product.** SLAs, routing and reporting are core features, and customization runs to deep ticketing config. Support here is ticket-driven with explicit SLA and reporting capability, not a lighter shared-inbox model.

**Collaboration is organized around tickets.** With deep ticketing config, agents work shared tickets across email, chat, phone and social with ownership and internal handling. Agent collaboration centers on the ticket record instead of a lightweight inbox.

## COST

**A free program for up to 2 agents (currently a 6-month term under Freshworks' terms),** then paid tiers from Growth at \$19/agent/mo up to Enterprise at \$89/agent/mo. Pricing is per agent, with a low-cost entry point for teams of up to two and a range that extends from small teams to enterprise ticketing.

## ECOSYSTEM & EXIT

**A large marketplace, part of the wider Freshworks ecosystem,** sits alongside an API. This gives a sizeable catalogue of prebuilt connectors on top of programmatic access, fitting its role as a full ticketing platform.

**Setup is moderate because of the depth of its ticketing configuration,** and migration is a standard import. Expect more upfront configuration than the simplest desks, offset by a conventional import for existing ticket data.

## WHERE IT STANDS OUT

**Its standout is classic ticketing depth (SLAs, routing and reporting with deep configuration) combined with a free program for up to 2 agents (currently a 6-month term).** It is full ticketing that a very small team can start on at no cost for that period.

## LIMITATIONS

**The limitation is configuration effort.** Setup is moderate and customization is deep ticketing config, so it carries a more classic, heavier feel than the simplest shared-inbox desks. Getting the SLA, routing and reporting depth working takes upfront setup.

---

Pricing and limits read from [Freshdesk's own site](#) . Verified to the edition date above.

# Crisp

## BUILT FOR

Startups and small businesses wanting an affordable all-in-one chat, shared inbox, and chatbot widget on their site.

---

## CHOOSE THIS IF YOU ARE

Small-to-mid teams of ~10-20 who get penalized by per-seat pricing and want everything flat.

## BEST WHEN

Predictable flat-rate billing wins: Mini at \$45/mo flat for 4 seats and Essentials at \$95/mo for 10 seats, bundling live chat, shared inbox, chatbot and a CRM without per-agent fees, cutting costs 50-67% versus per-seat rivals for teams that size.

---

## WHAT IT DOES

**It bundles AI credits into each plan and targets small teams that want flat per-workspace pricing**, covering chat, email and social.

**Chat, email and social.** The channels are chat-led, matching its chat-first design, with email and social also supported. Conversations across these channels feed the same workspace where the Hugo AI agent operates.

**The native Hugo AI agent, which is reported to resolve up to ~60% of repetitive questions**, drawing on AI credits bundled into each plan (\$5 Mini, \$25 Essentials, \$75 Plus). AI capacity draws down against the plan's credit allowance. There is no per-outcome bill.

**Self-service centers on the native Hugo AI agent**, which is reported to resolve up to ~60% of repetitive questions using each plan's bundled AI credits (\$5, \$25, \$75). Repetitive queries are deflected by Hugo, drawing down credits instead of incurring per-outcome charges.

**Customization is chat-first, so ticketing and routing sit around live conversations.** There is no dedicated SLA engine. SLA and reporting tooling is not documented in the sources used here. This is a chat-led workspace suited to smaller teams.

**The workspace provides a shared team view across chat,** email and social. With chat-first customization, collaboration is organized around live conversations in one per-workspace space, consistent with its small-team, flat-priced positioning.

---

## COST

**Flat per workspace rather than per seat:** Free \$0, Mini \$45/mo for 4 seats, Essentials \$95/mo for 10 seats and Plus \$295/mo, each including bundled AI credits (\$5, \$25, \$75). Cost is fixed per plan regardless of usage up to the seat and credit limits.

---

## ECOSYSTEM & EXIT

**Crisp connects to standard tools through common integrations and can be extended via its API.** The surface is modest, with no large marketplace behind it, in line with its chat-first, small-team focus.

**Setup is easy and migration standard.** With flat per-workspace pricing and a chat-first design, going live is quick, with data migration following a standard process, not a custom one.

---

## WHERE IT STANDS OUT

**The differentiator is flat per-workspace pricing (Free \$0, Mini \$45, Essentials \$95, Plus \$295) with AI credits bundled into each tier and a Hugo agent resolving up to ~60% of repetitive questions.** Cost is predictable per workspace rather than per seat.

---

## LIMITATIONS

**The bundled AI credits (\$5, \$25, \$75) draw down with use.** Heavy Hugo usage can exhaust a plan's allowance, making heavy usage a real cost risk. Integrations are modest rather than an extensive marketplace, and customization is chat-first.

---

Pricing and limits read from [Crisp's own site](#) . Verified to the edition date above.

# Chatwoot

## BUILT FOR

Engineering-minded teams wanting an open-source, self-hostable support inbox for chat, email, and social channels. Best fit: pick it when data control, self-hosting, or avoiding per-seat SaaS fees outweighs turnkey convenience.

---

## CHOOSE THIS IF YOU ARE

Technical teams that want an open-source, self-hostable helpdesk they fully control.

## BEST WHEN

Data ownership and cost-at-scale matter: free open-source Community edition to self-host, or cloud tiers Startups \$19, Business \$39, Enterprise \$99 per agent/mo, with an omnichannel inbox and Captain AI credits.

## AVOID IF

You don't want to run and maintain servers/updates or you need turnkey AI deflection out of the box.

## WHAT IT DOES

**It runs self-hosted or on cloud and suits teams with engineering capacity that want data ownership and low cost.**

**Live chat, email and omni-channel.** It presents as an omnichannel desk covering live chat and email plus additional channels, all run on infrastructure you self-host or take on its cloud, keeping channel data under your control.

**AI is included in the base, whether self-hosted or on cloud.** AI ships in the base price, not metered per-resolution usage. There is no separate per-outcome AI charge to account for on top of the plan or infrastructure cost.

**Self-service runs on the same self-hosted or cloud deployment as the rest of the platform,** with AI included in the base, not metered. There is no separate per-resolution charge for automated answers, so self-service cost is folded into the plan or the infrastructure you run.

As an open-source, fully customizable platform, **ticketing and routing are whatever you configure and self-host**. Configuration is open and self-managed with no prebuilt SLA tooling, so the ticketing model depends on how you set up the deployment.

**Collaboration happens in a shared**, open-source team workspace that you self-host or run on cloud, spanning live chat, email and its omni-channel inbox. Being fully customizable, how the shared inbox and collaboration behave depends on your own configuration.

---

## COST

**Self-hosting the open-source version is free beyond the infrastructure you run it on; paid self-hosted and cloud plans run from \$19 to \$99/agent/mo. Cost is therefore either infrastructure-only or a per-agent plan**, depending on whether you self-host the free build or take a paid tier.

---

## ECOSYSTEM & EXIT

**An open API and omnichannel integrations are the core**, and because the platform is open-source and self-hostable you can extend or modify it directly. Integration is therefore both API-based and open at the code level for teams with engineering capacity.

**Setup requires engineering to self-host**, so going live is an infrastructure task, not a sign-up. Migration benefits from open-source portability, meaning your data and deployment stay under your control, but standing it up and maintaining it needs technical capacity.

---

## WHERE IT STANDS OUT

**Its standout is open-source (MIT) data ownership**. Self-hosting is free beyond infrastructure, AI is included in the base, and it positions as an alternative to Intercom, Zendesk and Salesforce Service Cloud for teams that can self-host.

---

## LIMITATIONS

**Self-hosting is the trade-off**. It requires engineering to self-host, and you own the running, updating and maintenance of the deployment. The low-cost, data-ownership benefits come with the operational burden of infrastructure you manage yourself.

---

Pricing and limits read from [Chatwoot's own site](#) . Verified to the edition date above.

# Zendesk

## BUILT FOR

Larger support organizations and enterprises needing detailed ticketing, omnichannel, reporting, and many integrations at scale. / Best fit: Suited to a big team with complex workflows and compliance needs that outgrow lighter helpdesks.

---

## CHOOSE THIS IF YOU ARE

Scaling mid-market and enterprise support orgs that need deep routing, SLAs, reporting, and WFM.

## BEST WHEN

Breadth and enterprise features justify the premium: Suite Team \$55, Professional \$115 per agent/mo, Enterprise custom, with a Copilot AI add-on at ~\$50/agent. A realistic mid-market all-in lands \$165-\$265/agent once QA/WFM/AI stack up.

---

## WHAT IT DOES

**Unifies support requests from every channel into one agent workspace with automation,** knowledge base, AI and analytics.

**True omnichannel: email, web/mobile chat and messaging,** social (WhatsApp, Instagram, Facebook, X), voice/phone, SMS and a customer help-center portal, all converging into a single unified ticket queue and agent workspace.

**Zendesk AI (and AI/Copilot agents) auto-resolve tickets,** suggest replies, triage/route and summarize. Priced separately. A Copilot add-on (~\$50/agent/mo) plus outcome-based per-resolution AI fees on top of Suite seats.

**Guide powers a branded, multi-language help center with SEO-friendly articles,** community forums, and an AI answer bot/help-center search that deflects tickets. Content-block management and workflow approvals sit on higher tiers.



**Skills/omnichannel routing, triggers/automations, SLA policies with breach alerts,** and Explore analytics for dashboards. Custom SLAs, advanced routing and richer reporting unlock on Growth/Professional and above.

**Agents collaborate via internal notes,** side conversations (loop in colleagues/vendors by email/Slack), collision detection and light agents. It is a full ticketing system, not a simple shared-inbox tool, adding power but also overhead.

---

### COST

**Per-agent/month, billed annually: Suite Team \$55,** Suite Professional \$115, plus Suite Enterprise (quote; Copilot add-on \$50/agent/mo). A cheaper Support-only tier starts at \$19. AI features and per-resolution charges stack on top, inflating real cost.

---

### ECOSYSTEM & EXIT

**The Zendesk Marketplace lists 1,000+ prebuilt apps (Salesforce, Jira, Shopify, Slack, HubSpot) plus a documented REST API,** webhooks and the Sunshine platform for custom data, one of the larger integration ecosystems in support software.

**Fast to start but non-trivial to configure well (triggers, SLAs, routing).** Import via CSV, API, or third-party movers (Help Desk Migration). Data is exportable, but heavy customization and macros create real switching friction later.

---

### WHERE IT STANDS OUT

**Mature, scalable omnichannel ticketing with extensive routing/SLA/reporting and a large app marketplace.** A capable choice for teams expecting complex, high-volume, multi-channel support as they grow.

---

### LIMITATIONS

**Expensive and complex for early startups:** per-agent seats plus separately-billed AI/per-resolution fees, annual contracts, and a configuration learning curve. Simpler tools (Intercom, Freshdesk, Help Scout) are cheaper and faster for small teams.

---

Pricing and limits read from [Zendesk's own site](#) . Verified to the edition date above.

# Build

## BUILT FOR

Zero license cost, you own the data and the server upkeep.

---

## CHOOSE THIS IF YOU ARE

Pre-launch or tiny teams handling a trickle of support with near-zero budget.

### BEST WHEN

Choose this over any paid helpdesk when ticket volume doesn't yet justify per-seat fees and you'd rather spend engineering time than SaaS dollars.

### AVOID IF

You need SLAs, collision detection, reporting, or AI deflection, or support is becoming a real function; a shared inbox loses tickets and has no metrics.

## WHAT IT DOES

**Email out of the box; add web chat,** WhatsApp, Instagram, FB, Telegram via self-hosted Chatwoot. Phone/voice not included. Each extra channel is your integration and maintenance work, not a toggle.

**None by default.** You bolt on your own. Chatwoot supports LLM integrations, or script canned replies. Building a reliable AI deflection agent is real engineering. Off-the-shelf tools ship this. DIY means you own prompts, guardrails and hallucination risk.

**Chatwoot has a built-in help-center/KB.** A plain Gmail inbox has none. You'd add a static docs site (Docusaurus, Notion). Self-service is what cuts ticket volume. Budget time to actually write and maintain articles.

**Gmail gives labels and basic assignment,** with no real SLAs or analytics. Chatwoot adds assignment, teams, labels, basic reports. Neither matches enterprise SLA automation; fine for a small team, thin for compliance-driven or high-volume support.

**Google Groups collaborative inbox or Gmail delegation lets a few people share a queue with basic assignment,** but collisions and "who's got this?" are common. Chatwoot adds private notes and agent presence, much better once you're past 2-3 people.

## COST

**No per-seat fees, which is the big win.** Costs are a VPS (~\$10-40/mo) + your setup and ops time. Beats Zendesk/Intercom (\$50-100+/agent/mo) only if you'd otherwise pay for several seats and can maintain the box yourself.

No vendor page: this route is assembled from the components named above.

---

## ECOSYSTEM & EXIT

**Chatwoot ships a full REST API and webhooks.**

Gmail has the Gmail API. You wire CRM/billing links yourself. Flexible but integration is code you write and maintain, with no marketplace of one-click apps.

**Gmail shared inbox: minutes.** Self-hosted Chatwoot: hours to a day (Docker, domain, SMTP, SSL, upgrades). No white-glove migration of history; you script imports. Ongoing patching and backups are on you, a real ops cost.

---

## WHERE IT STANDS OUT

**Ideal for an early team with a technical member.** Near-zero per-seat cost, full data ownership (good for privacy-sensitive users), and no vendor lock-in while volume is low.

---

## LIMITATIONS

**Won't scale gracefully: no mature SLAs,** workforce management, advanced routing or polished AI deflection, and you own uptime, security patches and backups. A self-host outage is your outage. Migrate to a real helpdesk before support becomes a department.

# Workflow Automation

Compared: Zapier · Make · n8n · Activepieces · Gumloop · Pipedream · Build

# At a glance

| TOOL         | CHOOSE THIS IF YOU ARE                                                                                                               | STARTING PRICE                                                                                                                         | MAIN BENEFIT                                                                                                                                                                                                                                                                                                                    |
|--------------|--------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Zapier       | Non-technical operators who want the largest app library<br>the fastest path to a working automation.                                | Free tier of 100 tasks/mo<br>limited to two-step Zaps                                                                                  | <b>Its standout is reach: 9,000+ app integrations, the most of the group, paired with a low-friction setup for non-technical users.</b><br>When the priority is connecting to a specific SaaS product with minimal effort, the odds of a ready-made integration existing are highest here of any tool compared.                 |
| Make         | Visual builders who want capable multi-step scenarios at a fraction of Zapier's cost                                                 | Free tier of 1,000 credits/mo<br>(Make renamed 'operations' to 'credits' on Aug 27 2025)                                               | <b>Its standout is visual automation at a low entry price (Core is \$12/mo for 10,000 credits, formerly 'operations'), so complex, multi-branch scenarios are reachable on a modest budget.</b><br>The per-step credit model gives useful run headroom per dollar for workflows that aren't excessively step-heavy on each run. |
| n8n          | Technical teams that want source-available automation, self-hosting, and execution-based (not per-step) pricing.                     | The self -hosted Community Edition is free with unlimited executions<br>so run volume carries no marginal cost if you host it yourself | <b>Its standout is free, unlimited self-hosted executions combined with developer-grade flexibility: custom-code nodes, branching and full debugging.</b><br>For a technical team, run volume costs nothing when self-hosted and few no-code limits apply, trading infrastructure effort for capability and control.            |
| Activepieces | Teams that want a truly open-source (MIT), self-hostable automation platform with unlimited self-hosted runs at zero recurring cost. | MIT -licensed and free to self -host with no usage cap                                                                                 | <b>Its standout is the combination of a permissive MIT open-source license, free self-hosting, and 700+ integrations, plus agent tooling via MCP.</b><br>It offers open-source data control without the steepest developer curve, positioning itself between no-code simplicity and fully developer-oriented tools.             |
| Gumloop      | Teams building AI-first, agentic workflows (scraping, LLM chaining, and data extraction) over simple app-to-app glue.                | Gumloop discontinued its permanent free plan in mid-2026<br>replacing it with a 14-day Pro trial                                       | <b>Its standout is being AI-native.</b><br>Models do the substantive work across the workflow, with OpenAI, Anthropic or Google models pluggable into any node. For automations whose core task is interpretation, generation or extraction, the model-first design fits better than bolting AI onto a rule engine.             |
| Pipedream    | Developers who want to drop real code between steps and pay for compute time not step count.                                         | Free tier: 100 credits/day<br>3 active workflows, 3 connected accounts and unlimited testing                                           | <b>Strong code-in-the-loop automation: real programming languages inside workflows, 3,000+ integrations, and a usable free compute tier.</b><br>Suited to technical founders who outgrow Zapier's rigidity but don't want to self-host like n8n.                                                                                |
| Build        | Engineering teams who want total control near-zero SaaS spend for automation.                                                        | Software is free/open - source<br>you pay for a VPS (~\$5-40/mo) and your time                                                         | <b>Best for a technical team running high automation volume or handling sensitive data.</b><br>Massive cost savings vs per-task SaaS, no vendor limits, and complete control over logic and where data lives.                                                                                                                   |

Full detail for every tool follows on its own page.

# Zapier

## BUILT FOR

Non-technical operators and small teams connecting SaaS apps through simple trigger-action automations across thousands of integrations. / Best fit: A fit when breadth of app connectors and ease for non-coders matters more than deep logic or cost control.

---

## CHOOSE THIS IF YOU ARE

Non-technical operators who want the largest app library and the fastest path to a working automation.

## BEST WHEN

Breadth of 9,000+ integrations and ease matter most. Free gives 100 tasks/mo (two-step Zaps), Professional from \$19.99/mo (750 tasks) billed annually (about \$29.99 monthly), Team from \$69/mo annually (about \$103.50 monthly). A task is one successful action, so triggers and filters are free.

## AVOID IF

Your workflows are multi-step and high-volume.

## WHAT IT DOES

**Workflows are built as multi-step 'Zaps' and billed per task**, i.e. per action step, which ties cost to run volume rather than to how many workflows you build. Cloud-only, with no self-hosting option, and aimed at non-technical users.

**Supports multi-step Zaps with conditional paths and built-in error handling.** Adequate for branching logic, but the builder is list-based rather than a visual canvas (less visual than Make). Tracing and debugging an elaborate flow is harder. Best suited to linear or moderately branched automations.

## Built-in AI actions plus an AI assistant

**available across the connector catalog**, so AI steps can be added to standard Zaps without external setup. AI is layered onto the existing automation model, not its core; it complements the 9,000+ app integrations without becoming the centerpiece.

## The most approachable of the five for non-

**technical users.** Linear Zaps and a large pre-built connector library mean common automations can be assembled without code or data-mapping expertise. The trade-off is that the same linear model limits how complex a single workflow can practically get.

## COST

**Free tier of 100 tasks/mo, limited to two-step Zaps.** Professional starts at \$19.99/mo (annual) for 750 tasks. Team is \$69/mo (annual). Billing is per task, i.e. per action step, so cost scales directly with execution volume rather than with the number of workflows you build.

---

## ECOSYSTEM & EXIT

**9,000+ apps, the most of the five tools compared here.** That breadth is Zapier's main draw for non-technical users, since most SaaS products a founder already relies on are likely to have a pre-built Zapier integration rather than requiring custom code or workarounds to connect.

**No self-hosting.** Zapier is cloud-only. All workflow data passes through the vendor's managed infrastructure. There is no option to run it inside your own environment, which matters for teams with data-residency or full-control requirements that a hosted-only platform cannot satisfy.

---

## WHERE IT STANDS OUT

**Its standout is reach: 9,000+ app integrations,** the most of the group, paired with a low-friction setup for non-technical users. When the priority is connecting to a specific SaaS product with minimal effort, the odds of a ready-made integration existing are highest here of any tool compared.

## LIMITATIONS

**Per-task, per-action-step billing means cost rises directly with execution volume,** which can get expensive for high-throughput or multi-step automations. It is also cloud-only with no self-hosting. Teams needing data-residency control or volume-independent pricing will hit its structural limits.

---

Pricing and limits read from Zapier's own site . Verified to the edition date above.

# Make

## BUILT FOR

---

### CHOOSE THIS IF YOU ARE

#### BEST WHEN

You want a rich drag-and-drop canvas with cheap runs: Free (1,000 credits/mo) then Core \$12, Pro \$21, Teams \$38/mo, all paid plans starting at 10,000 credits (Make moved from "operations" to "credits" on Aug 27 2025); routers, iterators and error handling make complex logic visual.

#### WHAT IT DOES

**Every step in a scenario (trigger, filter or action) consumes one operation.** A 10-step scenario uses 10 operations per run, tying cost to workflow complexity as much as volume. Cloud-only, with AI modules available.

**Visual scenarios with routers and filters for branching**, plus detailed run inspection that shows the data passing through each module on a given execution. The canvas makes multi-branch logic easier to follow, and per-run inspection helps trace where a scenario failed and which data caused it.

#### AVOID IF

You want per-workflow-execution billing regardless of module count or full self-hosting, unlimited self-hosted flows, or the very largest app catalog and simplest UX.

**AI is delivered as modules you drop into a scenario alongside other steps**, so a model call becomes one operation within a broader visual flow. AI is an optional component of a scenario and never the driver, fitting the operations-based model where each AI step counts as an operation like any other.

**A steeper learning curve than Zapier.** The visual canvas, routers, filters and operation model all take time to grasp. Once learned, it expresses more complex logic than a linear builder, so the added ramp-up buys more capability, better suited to builders willing to invest time in the tool.



## COST

**Free tier of 1,000 credits/mo (Make renamed 'operations' to 'credits' on Aug 27 2025).** Core is \$12/mo for 10,000 credits, Pro \$21/mo, Teams \$38/mo. Because every step (trigger, filter or action) costs one credit, a 10-step scenario consumes 10 credits per run, so the real capacity of each tier depends heavily on how complex your scenarios are.

---

## ECOSYSTEM & EXIT

**A broad app catalog paired with deep,** granular module options per connector. No specific integration count is documented, but the depth means each app typically exposes multiple modules. Scenarios can target specific API actions, not only high-level triggers and actions.

**No self-hosting.** Make is cloud-only, so scenarios and the data they process run entirely on the vendor's infrastructure. Teams needing on-premises deployment or in-house data control won't find it here. The platform is offered solely as a managed cloud service with no self-hosted path.

---

## WHERE IT STANDS OUT

**Its standout is visual automation at a low entry price (Core is \$12/mo for 10,000 credits, formerly 'operations'),** so complex, multi-branch scenarios are reachable on a modest budget. The per-step credit model gives useful run headroom per dollar for workflows that aren't excessively step-heavy on each run.

## LIMITATIONS

Because every step consumes an operation, **complex scenarios burn through an allowance quickly (a 10-step scenario uses 10 operations per run).** Operation counting adds up on elaborate or high-frequency workflows. It is also cloud-only, with no self-hosted deployment option available.

---

Pricing and limits read from **Make's own site** . Verified to the edition date above.

# n8n

## BUILT FOR

Developers and technical teams wanting an open-source, self-hostable automation platform with code nodes and full data control. Best fit: pick it when self-hosting, custom code, and unlimited executions beat the convenience of hosted no-code tools.

---

## CHOOSE THIS IF YOU ARE

Technical teams that want source-available automation, self-hosting, and execution-based (not per-step) pricing.

### BEST WHEN

Complex, high-volume workflows make per-step billing painful. One execution = one full workflow run no matter how many nodes, so a 30-step flow costs the same as a 3-step one. Cloud is \$24/mo Starter or \$60/mo Pro, and the Community edition is free to self-host with code nodes and AI/LLM building blocks.

### AVOID IF

You're non-technical and want the easiest UX and biggest app library, a purely visual cloud canvas, or a fully MIT-licensed self-host.

---

## WHAT IT DOES

**The self-hosted Community Edition is free with unlimited executions**, while managed Cloud plans are metered by execution count. Combines a large catalog of built-in nodes with custom-code nodes, and runs either self-hosted or on cloud.

**The most flexible option here for logic**, with branching, custom-code nodes for arbitrary transformation, and full execution debugging for inspecting and re-running individual steps. The code escape hatch means edge cases a pure no-code builder can't express can still be handled directly within a flow.

**Provides dedicated AI and agent nodes plus LangChain-style building blocks.** You can assemble multi-step AI agents and chains inside a workflow. Combined with code nodes, this supports custom AI logic beyond single model calls: chaining, tool use and agent patterns built from composable nodes.

**Developer-oriented, with the steepest learning curve of the five.** Concepts like nodes, expressions, code steps and self-hosting assume technical comfort, and there is more to configure than in a no-code tool. The payoff is flexibility, but non-developers will find the ramp-up meaningfully harder.

## COST

**The self-hosted Community Edition is free with unlimited executions**, so run volume carries no marginal cost if you host it yourself. Managed cloud is metered by execution: Cloud Starter is \$24/mo for 2,500 executions and Pro is \$60/mo for 10,000 executions, shifting the trade-off to infrastructure effort.

## ECOSYSTEM & EXIT

**A large catalog of built-in nodes**, backed by custom-code nodes for anything not covered natively. No exact count is documented, but the code-node escape hatch means integrations aren't limited to the pre-built list. A developer can call arbitrary APIs directly from within a workflow.

**Free self-hosting via the Community Edition**, giving full data control since workflows and data can run entirely inside your own infrastructure, with unlimited executions and no per-run cost. This is a core differentiator for teams with data-residency, privacy or cost-at-scale requirements.

## WHERE IT STANDS OUT

**Its standout is free, unlimited self-hosted executions combined with developer-grade flexibility**: custom-code nodes, branching and full debugging. For a technical team, run volume costs nothing when self-hosted and few no-code limits apply, trading infrastructure effort for capability and control.

## LIMITATIONS

**It is developer-oriented, with the steepest learning curve here**. Non-technical users may struggle. And while self-hosting is free, it requires infrastructure to deploy, secure, update and maintain; the free, unlimited executions come with real operational overhead and technical responsibility.

---

Pricing and limits read from [n8n's own site](#) . Verified to the edition date above.

# Activepieces

## BUILT FOR

Developers and small teams wanting an open-source, MIT-licensed automation tool with AI steps and a simpler self-hosted setup.

---

## CHOOSE THIS IF YOU ARE

Teams that want a truly open-source (MIT), self-hostable automation platform with unlimited self-hosted runs at zero recurring cost.

### BEST WHEN

Predictable billing and no self-hosted run caps dominate: self-host free for unlimited, production-ready flows, or managed cloud tiers (free single-user, Plus \$16/mo and Team \$166/mo, billed yearly, credit-based).

### AVOID IF

You need the broadest connector catalog and consumer-grade polish, the most mature visual scenario tooling, or the largest self-hosted community and node library today.

## WHAT IT DOES

**The Community Edition is free to self-host with unlimited runs.** The managed cloud uses a credit-based model (a free single-user tier with daily credits, Plus at \$16/mo and Team at \$166/mo, both billed yearly). Ships 700+ integrations, extensible via open-source pieces.

**Offers branching and code steps within a visual builder**, so no-code users get conditional logic while developers can drop into code where needed. It sits as a middle ground: more extensibility than a purely no-code tool through code steps, without requiring code for standard branching flows.

**Offers AI 'pieces' and agent support**, including ~400 MCP servers that let agents connect to external tools and data via the Model Context Protocol. This positions it for agent-style automations where the model calls tools instead of just returning single completions, and the MCP catalog is self-hostable.

**Approachable as a no-code builder while leaving optional code steps for more technical users.** A non-developer can build standard flows and a developer can extend them. It sits between Zapier-style simplicity and n8n's developer depth, aiming to serve both skill levels within one interface.

## COST

**MIT-licensed and free to self-host with no usage cap.** The managed cloud is credit-based: a free single-user tier with daily-refreshing credits and unlimited flows, then Plus at \$16/mo (billed yearly, up to 5 users) and Team at \$166/mo (billed yearly, 25 users). Self-hosting stays uncapped on runs.

## ECOSYSTEM & EXIT

**700+ integrations, extensible through open-source 'pieces' that anyone can contribute or self-author.** Fewer than Zapier's catalog, but the open contribution model means the list can grow via the community, and self-hosters can add private connectors themselves instead of waiting on the vendor.

**MIT-licensed and completely free to self-host,** giving full data control with workflows running in your own environment. The permissive MIT license also allows modifying and redistributing the platform, which is broader freedom than a fair-code or source-available license grants self-hosters.

## WHERE IT STANDS OUT

**Its standout is the combination of a permissive MIT open-source license,** free self-hosting, and 700+ integrations, plus agent tooling via MCP. It offers open-source data control without the steepest developer curve, positioning itself between no-code simplicity and fully developer-oriented tools.

## LIMITATIONS

**A newer entrant with a smaller integration catalog (700+) than the most-integrated tools in this group.** A specific niche connector may not exist yet. Self-hosting also carries the usual infrastructure and maintenance overhead, though the open-source pieces model adds missing connectors.

---

Pricing and limits read from [Activepieces's own site](#) . Verified to the edition date above.

# Gumloop

## BUILT FOR

Operators and growth teams building AI-heavy automations for scraping, enrichment, and content workflows on a visual canvas. Best fit: choose it when your automations center on LLM steps and data pipelines rather than plain app-to-app connectors.

---

## CHOOSE THIS IF YOU ARE

Teams building AI-first, agentic workflows (scraping, LLM chaining, and data extraction) over simple app-to-app glue.

### BEST WHEN

The job is AI automation, not integrations: a node-based canvas purpose-built for LLM steps, web scraping and document processing with Pro at \$37/mo.

### AVOID IF

You mainly need broad SaaS connectors and CRUD syncing, execution-based pricing for step-heavy classic workflows, or a free self-hosted open-source base; Gumloop's strength is AI, not connector breadth.

## WHAT IT DOES

**You plug OpenAI, Anthropic or Google models into any node**, so the model, not fixed rule logic, drives each step. Cloud-only, billed on a credit system tied to model usage.

**Logic is organized as nodes on a visual canvas**, but oriented around AI steps instead of deterministic rule branches. The model handles interpretation inside a node. Dedicated error-handling or debugging tooling is not documented in the sources used here, so its strength is AI-driven steps over rigid rule logic.

**AI-first by design.** You plug OpenAI, Anthropic or Google models into any node. The model performs the core work of each step, not an add-on. This makes AI the default unit of automation here. Workflows are built around model calls, with connectors supplying their inputs and outputs.

**A visual canvas oriented around AI steps**, which keeps building AI workflows approachable for those comfortable with prompting and model behavior. The learning curve centers on designing effective model-driven nodes rather than wiring deterministic logic. It suits AI-first builders over rule-based ones.

## COST

**Gumloop discontinued its permanent free plan in mid-2026**, replacing it with a 14-day Pro trial. Pro is \$37/mo for ~20,000 credits/mo (240,000/year) with unlimited seats and scales up a credit slider. Enterprise is custom-priced. Because usage is credit-metered and credits fund AI model calls, cost tracks how much model work each workflow does, not just run count.

## ECOSYSTEM & EXIT

**Connectors here play a supporting role.** Instead of being the automation themselves, they feed data into AI model nodes that do the processing. No integration count is documented; the design assumes connectors supply inputs and outputs while the model handles the interpretation in between.

**No self-hosting.** Gumloop is cloud-only. Both workflows and the data fed to its AI models run on the vendor's infrastructure. Because AI nodes also send data to external model providers, teams with strict data-control or residency needs have limited ability to contain where that data flows.

## WHERE IT STANDS OUT

**Its standout is being AI-native.** Models do the substantive work across the workflow, with OpenAI, Anthropic or Google models pluggable into any node. For automations whose core task is interpretation, generation or extraction, the model-first design fits better than bolting AI onto a rule engine.

## LIMITATIONS

**Costs are credit-based and tied to AI model usage.** Heavier model work consumes credits faster and makes spend less predictable than flat per-task pricing. It is also cloud-only with no self-hosting, meaning data, including inputs sent to AI models, must pass through the vendor's infrastructure.

---

Pricing and limits read from [Gumloop's own site](#) . Verified to the edition date above.

# Pipedream

## BUILT FOR

Developers wanting a code-first automation platform with generous free executions, hosted functions, and thousands of API integrations. Best fit: pick it when you are comfortable writing Node or Python glue and want developer-grade workflows over pure no-code tools.

---

## CHOOSE THIS IF YOU ARE

Developers who want to drop real code between steps and pay for compute time, not step count.

### BEST WHEN

You want code-first automation with generous limits. Free gives 100 credits/day, Basic \$29/mo, Advanced \$99/mo, and billing is by execution time (~1 credit ≈ 30s at 256MB), so a typical run is about one credit regardless of steps.

### AVOID IF

You're non-technical and want pure drag-and-drop, you need self-hosting and data ownership, or your workloads are AI-agent heavy; Pipedream rewards writing code.

---

## WHAT IT DOES

**Connects 3,000+ APIs and runs on managed serverless compute**, with no infrastructure to operate.

**Suited to complex logic: inline code steps (Node/Python/Go/Bash)**, branching/control flow, data stores, and per-step inspection with full request/response logs; event history and replay make debugging stronger than pure no-code tools.

### A core focus: prebuilt OpenAI/Anthropic

**components**, an MCP/agent connector layer, AI-agent-oriented tooling, and included AI tokens on the free tier. Increasingly positioned for LLM agents and AI workflows.

**Steeper than Zapier for non-coders.** Its power lives in code steps. Very comfortable for developers; a visual builder and prebuilt actions handle basic flows, but true depth assumes some coding fluency.



## COST

**Free tier: 100 credits/day, 3 active workflows,** 3 connected accounts and unlimited testing. Paid tiers add workflows, credits and concurrency: Basic \$29/mo, Advanced \$79/mo, Business custom. 1 credit  $\approx$  30s compute at 256MB.

---

Pricing and limits read from [Pipedream's own site](#) . Verified to the edition date above.

## ECOSYSTEM & EXIT

**3,000+ prebuilt app connectors**, plus the ability to call any API directly via HTTP. Managed OAuth for connected accounts and a large registry of first-party and community components cover almost anything unlisted.

**Cloud-only and serverless, with no official self-hosted/open-source runtime like n8n or Activepieces.** Data flows through Pipedream's managed infrastructure. Great for convenience, weaker for on-prem or strict data-residency requirements.

## WHERE IT STANDS OUT

**Strong code-in-the-loop automation: real programming languages inside workflows**, 3,000+ integrations, and a usable free compute tier. Suited to technical founders who outgrow Zapier's rigidity but don't want to self-host like n8n.

## LIMITATIONS

**Not for non-coders at depth, no self-hosting**, and credit-based compute can surprise at scale. Watch item: Acquired by Workday (announced Nov 2025; deal has since closed); long-term roadmap and independence for small external users are uncertain.

# Build

## BUILT FOR

---

### CHOOSE THIS IF YOU ARE

Engineering teams who want total control and near-zero SaaS spend for automation.

#### BEST WHEN

Choose this over every hosted tool when data must stay in-house, volumes are large enough that SaaS billing hurts, or your logic is too custom for prebuilt connectors.

#### AVOID IF

You lack the ops bandwidth to run, monitor and update it, or you need hundreds of maintained connectors out of the box; DIY trades convenience and reliability for control.

### WHAT IT DOES

**Unlimited: branching, loops, custom code,** retries, whatever you build. But error handling, retries, alerting and logging are yours to implement. A silently failed cron job at 3am is a class of bug SaaS platforms surface for you.

#### **n8n/Activepieces have native LLM/agent**

**nodes;** in scripts you call any model API directly with zero markup. Most flexible and cheapest per token, but you own prompt plumbing, rate limits and cost control instead of a managed AI step.

**n8n's visual builder is approachable.** Pure scripts require a developer. Either way you must handle hosting, updates and secrets, steeper than signing into Zapier. Realistically needs a technical founder or it becomes a liability.

### COST

**Software is free/open-source.** You pay for a VPS (~\$5-40/mo) and your time. Undercuts Zapier/Make's per-task pricing dramatically at high volume, but it only pays off once task counts are high enough to offset the hours you sink into running it.

### ECOSYSTEM & EXIT

**n8n ships 400+ prebuilt nodes and generic HTTP for anything with an API.** Raw scripts can hit any endpoint but you write each auth/pagination flow yourself. Broadest reach, most manual effort; nothing is truly "one click."

**The core advantage: data never leaves your server**, no third party sees the workflows or credentials, which is strong for privacy/compliance. Trade-off: you own uptime, security patches, backups and scaling. The automation is only as reliable as the box it runs on.

---

No vendor page: this route is assembled from the components named above.

---

## WHERE IT STANDS OUT

**Best for a technical team running high automation volume or handling sensitive data.**

Massive cost savings vs per-task SaaS, no vendor limits, and complete control over logic and where data lives.

---

## LIMITATIONS

**Setup, maintenance and on-call are on you.** No vendor support when it breaks. Fragile for business-critical flows unless you build proper monitoring and redundancy. Getting it wrong means silent failures and missed events, often worse than a SaaS bill.

# Methodology

**The primary sources are linked.** Every tool name at the top of its profile links to that vendor's own site, which is the authority for its pricing, tiers and limits. Each profile also carries a direct link to the page the pricing was read from, so any figure here can be checked at source.

**Where the facts come from.** Every fact here comes from the research workbook behind this guide, compiled in August 2026. Where a claim rests on a particular source we name it in the sentence, as with the contract figures given "per Vendr data."

**The evidence standard.** We treated vendor pricing pages and documentation as authoritative on what a vendor knows best: its own tiers, limits and feature gates. Anywhere the vendor's account could differ from what teams actually run into, such as real cost, where a tool strains and what leaving it involves, we weighted independent practitioner reports and third-party contract data higher. We used no affiliate or sponsored content.

**What this guide does not carry.** The workbook keeps its citations inside the text of each claim rather than as separate metadata. There is no per-claim tag marking a fact as officially documented or user-reported, and no per-claim verification date. We would rather say that plainly than invent the apparatus. Each profile is verified to the edition date in its footer, and where a claim's origin matters we name it in the sentence. So treat the inline attributions as the provenance record, treat every price as accurate to the edition date and no later, and check anything price-critical against the vendor's own page before you commit. Prices in this category moved twice while we were compiling it.

**What was selected.** Each category compares six tools plus the do-it-yourself route. We kept lesser-known tools alongside the obvious ones, and left out anything an early-stage company couldn't realistically use.

**Limitations.** Pricing and features move fast. Figures carry their dates where the source gives one, and anything price-critical is worth checking against the vendor's current page.

# Update log

August 2026, first edition. All 14 chapters.